

Package ‘mlr’

February 13, 2016

Title Machine Learning in R

Description Interface to a large number of classification and regression techniques, including machine-readable parameter descriptions. There is also an experimental extension for survival analysis, clustering and general, example-specific cost-sensitive learning. Generic resampling, including cross-validation, bootstrapping and subsampling. Hyperparameter tuning with modern optimization techniques, for single- and multi-objective problems. Filter and wrapper methods for feature selection. Extension of basic learners with additional operations common in machine learning, also allowing for easy nested resampling. Most operations can be parallelized.

URL <https://github.com/mlr-org/mlr>

BugReports <https://github.com/mlr-org/mlr/issues>

License BSD_2_clause + file LICENSE

Encoding UTF-8

Depends R (>= 3.0.2), BBmisc (>= 1.9), ggplot2, ParamHelpers (>= 1.7), stats

Imports checkmate (>= 1.7.1), ggvis, methods, parallelMap (>= 1.3), plyr, reshape2, shiny, survival

Suggests ada, adabag, bartMachine, brnn, bst, care, caret (>= 6.0-57), class, clue, cluster, clusterSim, cValid, cmaes, CoxBoost, crs, Cubist, deepnet, DiceKriging, DiceOptim, DiscriMiner, e1071, earth, elasticnet, elmNN, emoa, extraTrees, flare, fields, FNN, fpc, frbs, FSelector, gbm, GenSA, glmnet, Hmisc, irace (>= 1.0.7), kernlab, kknn, klaR, knitr, kohonen, laGP, LiblineaR, lqa, MASS, mboost, mco, mda, mlbench, modeltools, mRMRe, nnet, nodeHarvest (>= 0.7-3), neuralnet, numDeriv, pamr, party, penalized, pls, PMCMR, pROC (>= 1.8), randomForest, randomForestSRC (>= 2.0.5), ranger (>= 0.3.0), RCurl, rFems, rjson, rknn, rmarkdown, robustbase, ROCR, rotationForest, rpart, rrlda, rsm, RSNNs, RWeka, sda, sparsediscrim, sparseLDA, stepPlr, SwarmSVM, testthat, tgp, TH.data, xgboost, XML

LazyData yes

ByteCompile yes

Version 2.8

VignetteBuilder knitr

RoxygenNote 5.0.1

NeedsCompilation yes

Author Bernd Bischl [aut, cre],

Michel Lang [aut],

Jakob Richter [aut],

Jakob Bossek [aut],

Leonard Judt [aut],

Tobias Kuehn [aut],

Erich Studerus [aut],

Lars Kotthoff [aut],

Zachary Jones [ctb],

Schiffner Julia [aut]

Maintainer Bernd Bischl <bernd_bischl@gmx.net>

Repository CRAN

Date/Publication 2016-02-13 08:37:35

R topics documented:

Aggregation	6
aggregations	7
agri.task	9
analyzeFeatSelResult	9
asROCRPrediction	10
bc.task	10
benchmark	11
BenchmarkResult	12
bh.task	13
capLargeValues	13
configureMlr	14
convertBMRTToRankMatrix	15
convertMLBenchObjToTask	16
costiris.task	17
createDummyFeatures	17
crossover	18
downsample	19
dropFeatures	19
estimateRelativeOverfitting	20
estimateResidualVariance	21
FailureModel	22
FeatSelControl	22
FeatSelResult	25
filterFeatures	26
friedmanPostHocTestBMR	27
friedmanTestBMR	28

generateCalibrationData	29
generateCritDifferencesData	30
generateFilterValuesData	32
generateLearningCurveData	33
generatePartialPredictionData	34
generateThreshVsPerfData	37
getBMRAggrPerformances	38
getBMRFeatSelResults	39
getBMRFilteredFeatures	40
getBMRLearnerIds	41
getBMRLearners	42
getBMRLearnerShortNames	42
getBMRMeasureIds	43
getBMRMeasures	44
getBMRModels	44
getBMRPerformances	45
getBMRPredictions	46
getBMRTaskIds	47
getBMRTuneResults	47
getCaretParamSet	48
getClassWeightParam	49
getConfMatrix	50
getDefaultMeasure	51
getFailureModelMsg	51
getFeatSelResult	52
getFilteredFeatures	52
getFilterValues	53
getHomogeneousEnsembleModels	54
getHyperPars	54
getLearnerModel	55
getMlrOptions	55
getMultilabelBinaryPerformances	56
getNestedTuneResultsOptPathDf	56
getNestedTuneResultsX	57
getParamSet	58
getPredictionProbabilities	58
getPredictionResponse	59
getProbabilities	60
getRRPredictions	60
getStackedBaseLearnerPredictions	61
getTaskClassLevels	61
getTaskCosts	62
getTaskData	63
getTaskDescription	64
getTaskFeatureNames	65
getTaskFormula	65
getTaskId	66
getTaskNFeats	67

getTaskSize	67
getTaskTargetNames	68
getTaskTargets	68
getTaskType	69
getTuneResult	70
hasProperties	70
imputations	71
impute	72
iris.task	74
isFailureModel	75
joinClassLevels	75
learnerArgsToControl	76
LearnerProperties	76
learners	77
listFilterMethods	77
listLearners	78
listMeasures	79
lung.task	80
makeAggregation	81
makeBaggingWrapper	82
makeClassifTask	83
makeCostMeasure	85
makeCostSensClassifWrapper	86
makeCostSensRegrWrapper	87
makeCostSensWeightedPairsWrapper	88
makeCustomResampledMeasure	88
makeDownsampleWrapper	90
makeFeatSelWrapper	91
makeFilter	92
makeFilterWrapper	93
makeFixedHoldoutInstance	94
makeImputeMethod	95
makeImputeWrapper	96
makeLearner	97
makeMeasure	99
makeModelMultiplexer	101
makeModelMultiplexerParamSet	102
makeMulticlassWrapper	103
makeMultilabelBinaryRelevanceWrapper	104
makeOverBaggingWrapper	105
makePreprocWrapper	107
makePreprocWrapperCaret	108
makeResampleDesc	108
makeResampleInstance	111
makeSMOTERWrapper	112
makeStackedLearner	113
makeTuneWrapper	115
makeUndersampleWrapper	116

makeWeightedClassesWrapper	118
makeWrappedModel	119
measures	120
mergeBenchmarkResultLearner	124
mergeBenchmarkResultTask	125
mergeSmallFactorLevels	125
mtcars.task	126
normalizeFeatures	126
oversample	127
performance	128
pid.task	129
plotBMRBoxplots	130
plotBMRRanksAsBarChart	131
plotBMRSummary	132
plotCalibration	133
plotCritDifferences	134
plotFilterValues	135
plotFilterValuesGGVIS	136
plotLearnerPrediction	137
plotLearningCurve	138
plotLearningCurveGGVIS	139
plotPartialPrediction	140
plotPartialPredictionGGVIS	141
plotROCCurves	142
plotThreshVsPerf	143
plotThreshVsPerfGGVIS	144
plotTuneMultiCritResult	145
plotTuneMultiCritResultGGVIS	146
plotViperCharts	147
predict.WrappedModel	148
Prediction	149
predictLearner	150
regr.randomForest	151
reimpute	151
removeConstantFeatures	152
removeHyperPars	153
resample	154
ResamplePrediction	156
ResampleResult	156
RLearner	157
selectFeatures	159
setAggregation	160
setHyperPars	161
setHyperPars2	162
setId	162
setPredictThreshold	163
setPredictType	163
setThreshold	164

smote	165
sonar.task	166
subsetTask	167
summarizeColumns	167
summarizeLevels	168
TaskDesc	169
train	170
trainLearner	171
TuneControl	172
TuneMultiCritControl	174
TuneMultiCritResult	176
tuneParams	177
tuneParamsMultiCrit	179
TuneResult	180
tuneThreshold	180
wdbc.task	181
yeast.task	182

Index	183
--------------	------------

Aggregation	<i>Aggregation object.</i>
-------------	----------------------------

Description

An aggregation method reduces the performance values of the test (and possibly the training sets) to a single value. To see all possible implemented aggregations look at [aggregations](#).

The aggregation can access all relevant information of the result after resampling and combine them into a single value. Though usually something very simple like taking the mean of the test set performances is done.

Object members:

id [character(1)] Name of the aggregation method.

name [character(1)] Long name of the aggregation method.

fun [function(task, perf.test, perf.train, measure, group, pred)] Aggregation function.

See Also

[makeAggregation](#)

Description

- **test.mean**
Mean of performance values on test sets.
- **test.sd**
Standard deviation of performance values on test sets.
- **test.median**
Median of performance values on test sets.
- **test.min**
Minimum of performance values on test sets.
- **test.max**
Maximum of performance values on test sets.
- **test.sum**
Sum of performance values on test sets.
- **train.mean**
Mean of performance values on training sets.
- **train.sd**
Standard deviation of performance values on training sets.
- **train.median**
Median of performance values on training sets.
- **train.min**
Minimum of performance values on training sets.
- **train.max**
Maximum of performance values on training sets.
- **train.sum**
Sum of performance values on training sets.
- **b632**
Aggregation for B632 bootstrap.
- **b632plus**
Aggregation for B632+ bootstrap.
- **testgroup.mean**
Performance values on test sets are grouped according to resampling method. The mean for every group is calculated, then the mean of those means. Mainly used for repeated CV.
- **test.join**
Performance measure on joined test sets. This is especially useful for small sample sizes where unbalanced group sizes have a significant impact on the aggregation, especially for cross-validation test.join might make sense now. For the repeated CV, the performance is calculated on each repetition and then aggregated with the arithmetic mean.

Usage`test.mean``test.sd``test.median``test.min``test.max``test.sum``test.range``test.rmse``train.mean``train.sd``train.median``train.min``train.max``train.sum``train.range``train.rmse``b632``b632plus``testgroup.mean``test.join`**Format**`None`**See Also**[Aggregation](#)

agri.task	<i>European Union Agricultural Workforces clustering task.</i>
-----------	--

Description

Contains the task (agri.task).

References

See [agriculture](#).

analyzeFeatSelResult	<i>Show and visualize the steps of feature selection.</i>
----------------------	---

Description

This function prints the steps [selectFeatures](#) took to find its optimal set of features and the reason why it stopped. It can also print information about all calculations done in each intermediate step.
Currently only implemented for sequential feature selection.

Usage

```
analyzeFeatSelResult(res, reduce = TRUE)
```

Arguments

- | | |
|--------|---|
| res | [FeatSelResult]
The result of of selectFeatures . |
| reduce | [logical(1)]
Per iteration: Print only the selected feature (or all features that were evaluated)?
Default is TRUE. |

Value

invisible(NULL) .

See Also

Other featsel: [FeatSelControl](#), [getFeatSelResult](#), [makeFeatSelWrapper](#), [selectFeatures](#)

asROCRPrediction	<i>Converts predictions to a format package ROCR can handle.</i>
------------------	--

Description

Converts predictions to a format package ROCR can handle.

Usage

```
asROCRPrediction(pred)
```

Arguments

pred	[Prediction] Prediction object.
------	--

See Also

Other predict: [getPredictionProbabilities](#), [getPredictionResponse](#), [plotViperCharts](#), [predict.WrappedModel](#), [setPredictThreshold](#), [setPredictType](#)

Other roc: [plotViperCharts](#)

bc.task	<i>Wisconsin Breast Cancer classification task.</i>
---------	---

Description

Contains the task (bc.task).

References

See [BreastCancer](#). The column "Id" and all incomplete cases have been removed from the task.

benchmark	<i>Benchmark experiment for multiple learners and tasks.</i>
-----------	--

Description

Complete benchmark experiment to compare different learning algorithms across one or more tasks w.r.t. a given resampling strategy. Experiments are paired, meaning always the same training / test sets are used for the different learners. Furthermore, you can of course pass “enhanced” learners via wrappers, e.g., a learner can be automatically tuned using [makeTuneWrapper](#).

Usage

```
benchmark(learners, tasks, resamplings, measures, keep.pred = TRUE,
          models = TRUE, show.info = getMlrOption("show.info"))
```

Arguments

learners	[(list of) Learner] Learning algorithms which should be compared.
tasks	[(list of) Task] Tasks that learners should be run on.
resamplings	[(list of) ResampleDesc ResampleInstance] Resampling strategy for each tasks. If only one is provided, it will be replicated to match the number of tasks. If missing, a 10-fold cross validation is used.
measures	[(list of) Measure] Performance measures for all tasks. If missing, the default measure of the first task is used.
keep.pred	[logical(1)] Keep the prediction data in the pred slot of the result object. If you do many experiments (on larger data sets) these objects might unnecessarily increase object size / mem usage, if you do not really need them. In this case you can set this argument to FALSE. Default is TRUE.
models	[logical(1)] Should all fitted models be stored in the ResampleResult ? Default is TRUE.
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .

Value

[BenchmarkResult](#) .

See Also

Other benchmark: [BenchmarkResult](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFIteratedFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

Examples

```
lrns = list(makeLearner("classif.lda"), makeLearner("classif.rpart"))
tasks = list(iris.task, sonar.task)
rdesc = makeResampleDesc("CV", iters = 2L)
meas = list(acc, ber)
bmr = benchmark(lrns, tasks, rdesc, measures = meas)
rmat = convertBMRTToRankMatrix(bmr)
print(rmat)
plotBMRSummary(bmr)
plotBMRBoxplots(bmr, ber, style = "violin")
plotBMRRanksAsBarChart(bmr, pos = "stack")
friedmanTestBMR(bmr)
friedmanPostHocTestBMR(bmr, p.value = 0.05)
```

BenchmarkResult

BenchmarkResult object.

Description

Result of a benchmark experiment conducted by [benchmark](#) with the following members:

results [[list of ResampleResult](#) :] A nested list of resample results, first ordered by task id, then by learner id.

measures [[list of Measure](#) :] The performance measures used in the benchmark experiment.

learners [[list of Learner](#) :] The learning algorithms compared in the benchmark experiment.

The print method of this object shows aggregated performance values for all tasks and learners.

It is recommended to retrieve required information via the `getBMR*` getter functions. You can also convert the object using [as.data.frame](#).

See Also

Other benchmark: [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFIteratedFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

bh.task	<i>Boston Housing regression task.</i>
---------	--

Description

Contains the task (bh.task).

References

See [BostonHousing](#).

capLargeValues	<i>Convert large/infinite numeric values in a data.frame or task.</i>
----------------	---

Description

Convert numeric entries which large/infinite (absolute) values in a data.frame or task. Only numeric/integer columns are affected.

Usage

```
capLargeValues(obj, target = character(0L), cols = NULL, threshold = Inf,
  impute = threshold, what = "abs")
```

Arguments

obj	[data.frame Task] Input data.
target	[character] Name of the column(s) specifying the response. Target columns will not be capped. Default is character(0).
cols	[character] Which columns to convert. Default is all numeric columns.
threshold	[numeric(1)] Threshold for capping. Every entry whose absolute value is equal or larger is converted. Default is Inf.
impute	[numeric(1)] Replacement value for large entries. Large negative entries are converted to -impute. Default is threshold.
what	[character(1)] What kind of entries are affected? "abs" means $\text{abs}(x) > \text{threshold}$, "pos" means $\text{abs}(x) > \text{threshold} \ \&\& \ x > 0$, "neg" means $\text{abs}(x) > \text{threshold} \ \&\& \ x < 0$. Default is "abs".

Value

data.frame

See Also

Other `eda_and_preprocess`: [createDummyFeatures](#), [dropFeatures](#), [mergeSmallFactorLevels](#), [normalizeFeatures](#), [removeConstantFeatures](#), [summarizeColumns](#)

Examples

```
capLargeValues(iris, threshold = 5, impute = 5)
```

configureMlr	<i>Configures the behavior of the package.</i>
--------------	--

Description

Configuration is done by setting custom [options](#).

If you do not set an option here, its current value will be kept.

If you call this function with an empty argument list, everything is set to its defaults.

Usage

```
configureMlr(show.info, on.learner.error, on.learner.warning,
             on.par.without.desc, on.par.out.of.bounds, show.learner.output)
```

Arguments

show.info	[logical(1)] Some methods of <code>mlr</code> support a <code>show.info</code> argument to enable verbose output on the console. This option sets the default value for these arguments. Setting the argument manually in one of these functions will overwrite the default value for that specific function call. Default is TRUE.
on.learner.error	[character(1)] What should happen if an error in an underlying learning algorithm is caught: “stop”: R exception is generated. “warn”: A <code>FailureModel</code> will be created, which predicts only NAs and a warning will be generated. “quiet”: Same as “warn” but without the warning. Default is “stop”.
on.learner.warning	[character(1)] What should happen if a warning in an underlying learning algorithm is generated: “warn”: The warning is generated as usual. “quiet”: The warning is suppressed. Default is “warn”.

```
on.par.without.desc
    [character(1)]
    What should happen if a parameter of a learner is set to a value, but no parameter
    description object exists, indicating a possibly wrong name:
    "stop": R exception is generated.
    "warn": Warning, but parameter is still passed along to learner.
    "quiet": Same as "warn" but without the warning.
    Default is "stop".

on.par.out.of.bounds
    [character(1)]
    What should happen if a parameter of a learner is set to an out of bounds value.
    "stop": R exception is generated.
    "warn": Warning, but parameter is still passed along to learner.
    "quiet": Same as "warn" but without the warning.
    Default is "stop".

show.learner.output
    [logical(1)]
    Should the output of the learning algorithm during training and prediction be
    shown or captured and suppressed? Default is TRUE.
```

Value

invisible(NULL) .

See Also

Other configure: [getMlrOptions](#)

convertBMRTToRankMatrix

Convert BenchmarkResult to a rank-matrix.

Description

Computes a matrix of all the ranks of different algorithms over different datasets (tasks). Ranks are computed from aggregated measures. Smaller ranks imply better methods, so for measures that are minimized, small ranks imply small scores. for measures that are maximized, small ranks imply large scores.

Usage

```
convertBMRTToRankMatrix(bmr, measure = NULL, ties.method = "average",
    aggregation = "default")
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
measure	[Measure] Performance measure. Default is the first measure used in the benchmark experiment.
ties.method	[character(1)] See rank for details.
aggregation	[character(1)] “mean” or “default”. See getBMRAggrPerformances for details on “default”.

Value

matrix with measure ranks as entries. The matrix has one row for each learner, and one column for each task.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

Examples

```
# see benchmark
```

```
convertMLBenchObjToTask
```

*Convert a machine learning benchmark / demo object from package
mlbench to a task.*

Description

We auto-set the target column, drop any column which is called “Id” and convert logicals to factors.

Usage

```
convertMLBenchObjToTask(x, n = 100L, ...)
```


Arguments

x	[character(1)] Name of an mlbench function or dataset.
n	[integer(1)] Number of observations for data simul functions. Note that for a few mlbench function this setting is not exactly respected by mlbench. Default is 100.
...	[any] Passed on to data simul functions.

Examples

```
print(convertMLBenchObjToTask("Ionosphere"))
print(convertMLBenchObjToTask("mlbench.spirals", n = 100, sd = 0.1))
```

costiris.task	<i>Iris cost-sensitive classification task.</i>
---------------	---

Description

Contains the task (costiris.task).

References

See [iris](#). The cost matrix was generated artificially following
 Tu, H.-H. and Lin, H.-T. (2010), One-sided support vector regression for multiclass cost-sensitive classification. In ICML, J. Fürnkranz and T. Joachims, Eds., Omnipress, 1095–1102.

createDummyFeatures	<i>Generate dummy variables for factor features.</i>
---------------------	--

Description

Replace all factor features with their dummy variables. Internally `model.matrix` is used. Non factor features will be left untouched and passed to the result.

Usage

```
createDummyFeatures(obj, target = character(0L), method = "1-of-n",
  exclude = character(0L))
```

Arguments

obj	[data.frame Task] Input data.
target	[character(1) character(2) character(n.classes)] Name(s) of the target variable(s). Only used when obj is a data.frame, otherwise ignored. If survival analysis is applicable, these are the names of the survival time and event columns, so it has length 2. For multilabel classification these are the names of logical columns that indicate whether a class label is present and the number of target variables corresponds to the number of classes.
method	[character(1)] Available are: “1-of-n”: For n factor levels there will be n dummy variables. “reference”: There will be n-1 dummy variables leaving out the first factor level of each variable.
exclude	[character] Names of the columns to exclude. The target does not have to be included here. Default is none.

Value

data.frame | [Task](#) . Same type as obj.

See Also

Other eda_and_preprocess: [capLargeValues](#), [dropFeatures](#), [mergeSmallFactorLevels](#), [normalizeFeatures](#), [removeConstantFeatures](#), [summarizeColumns](#)

crossover	<i>Crossover.</i>
-----------	-------------------

Description

Takes two bit strings and creates a new one of the same size by selecting the items from the first string or the second, based on a given rate (the probability of choosing an element from the first string).

Arguments

x	[logical] First parent string.
y	[logical] Second parent string.
rate	[numeric(1)] A number representing the probability of selecting an element of the first string. Default is 0.5.

Value

[crossover](#) .

downsample	<i>Downsample (subsample) a task or a data.frame.</i>
------------	---

Description

Decrease the observations in a task or a ResampleInstance to a given percentage of observations.

Usage

```
downsample(obj, perc = 1, stratify = FALSE)
```

Arguments

obj	[Task ResampleInstance] Input data or a ResampleInstance.
perc	[numeric(1)] Percentage from [0, 1]. Default is 1.
stratify	[logical(1)] Only for classification: Should the downsampled data be stratified according to the target classes? Default is FALSE.

Value

[data.frame](#) | [Task](#) | [ResampleInstance](#) . Same type as obj.

See Also

[makeResampleInstance](#)

Other downsample: [makeDownsampleWrapper](#)

dropFeatures	<i>Drop some features of task.</i>
--------------	------------------------------------

Description

Drop some features of task.

Usage

```
dropFeatures(task, features)
```

Arguments

task	[Task] The task.
features	[character] Features to drop.

Value

[Task](#) .

See Also

Other eda_and_preprocess: [capLargeValues](#), [createDummyFeatures](#), [mergeSmallFactorLevels](#), [normalizeFeatures](#), [removeConstantFeatures](#), [summarizeColumns](#)

estimateRelativeOverfitting

Estimate relative overfitting.

Description

Estimates the relative overfitting of a model as the ratio of the difference in test and train performance to the difference of test performance in the no-information case and train performance. In the no-information case the features carry no information with respect to the prediction. This is simulated by permuting features and predictions.

Usage

```
estimateRelativeOverfitting(rdesc, measures, task, learner)
```

```
## S3 method for class 'ResampleDesc'
estimateRelativeOverfitting(rdesc, measures, task,
  learner)
```

Arguments

rdesc	[ResampleDesc] Resampling strategy.
measures	[Measure] list of Measure Performance measure(s) to evaluate. Default is the default measure for the task, see here getDefaultMeasure .
task	[Task] The task.
learner	[Learner] character(1) The learner. If you pass a string the learner will be created via makeLearner .

Details

Currently only support for classification and regression tasks is implemented.

Value

`data.frame` . Relative overfitting estimate(s), named by measure(s), for each resampling iteration.

References

Bradley Efron and Robert Tibshirani; Improvements on Cross-Validation: The .632+ Bootstrap Method, Journal of the American Statistical Association, Vol. 92, No. 438. (Jun., 1997), pp. 548-560.

See Also

Other performance: [makeCostMeasure](#), [makeCustomResampledMeasure](#), [makeMeasure](#), [measures](#), [performance](#)

Examples

```
task = makeClassifTask(data = iris, target = "Species")
rdesc = makeResampleDesc("CV", iters = 2)
estimateRelativeOverfitting(rdesc, acc, task, makeLearner("classif.knn"))
estimateRelativeOverfitting(rdesc, acc, task, makeLearner("classif.lda"))
```

estimateResidualVariance

Estimate the residual variance.

Description

Estimate the residual variance of a regression model on a given task. If a regression learner is provided instead of a model, the model is trained (see [train](#)) first.

Usage

```
estimateResidualVariance(x, task, data, target)
```

Arguments

<code>x</code>	[Learner or WrappedModel] Learner or wrapped model.
<code>task</code>	[RegrTask] Regression task. If missing, data and target must be supplied.
<code>data</code>	[<code>data.frame</code>] A data frame containing the features and target variable. If missing, task must be supplied.

target	[character(1)] Name of the target variable. If missing, task must be supplied.
--------	---

FailureModel	<i>Failure model.</i>
--------------	-----------------------

Description

A subclass of [WrappedModel](#). It is created - if you set the respective option in [configureMlr](#) - when a model internally crashed during training. The model always predicts NAs.

Its encapsulated learner `.model` is simply a string: The error message that was generated when the model crashed. The following code shows how to access the message.

Examples

```
configureMlr(on.learner.error = "warn")
data = iris
data$newfeat = 1 # will make LDA crash
task = makeClassifTask(data = data, target = "Species")
m = train("classif.lda", task) # LDA crashed, but mlr catches this
print(m)
print(m$learner.model) # the error message
p = predict(m, task) # this will predict NAs
print(p)
print(performance(p))
configureMlr(on.learner.error = "stop")
```

FeatSelControl	<i>Create control structures for feature selection.</i>
----------------	---

Description

Feature selection method used by [selectFeatures](#). The following methods are available:

FeatSelControlExhaustive Exhaustive search. All feature sets (up to a certain number of features `max.features`) are searched.

FeatSelControlRandom Random search. Features vectors are randomly drawn, up to a certain number of features `max.features`. A feature is included in the current set with probability `prob`. So we are basically drawing (0,1)-membership-vectors, where each element is Bernoulli(`prob`) distributed.

FeatSelControlSequential Deterministic forward or backward search. That means extending (forward) or shrinking (backward) a feature set. Depending on the given method different approaches are taken.

sfs Sequential Forward Search: Starting from an empty model, in each step the feature increasing the performance measure the most is added to the model.

sbs Sequential Backward Search: Starting from a model with all features, in each step the feature decreasing the performance measure the least is removed from the model.

sffs Sequential Floating Forward Search: Starting from an empty model, in each step the algorithm chooses the best model from all models with one additional feature and from all models with one feature less.

sfbs Sequential Floating Backward Search: Similar to sffs but starting with a full model.

FeatSelControlGA Search via genetic algorithm. The GA is a simple (μ , λ) or ($\mu + \lambda$) algorithm, depending on the comma setting. A comma strategy selects a new population of size μ out of the $\lambda > \mu$ offspring. A plus strategy uses the joint pool of μ parents and λ offspring for selecting μ new candidates. Out of those μ features, the new λ features are generated by randomly choosing pairs of parents. These are crossed over and `crossover.rate` represents the probability of choosing a feature from the first parent instead of the second parent. The resulting offspring is mutated, i.e., its bits are flipped with probability `mutation.rate`. If `max.features` is set, offspring are repeatedly generated until the setting is satisfied.

Usage

```
makeFeatSelControlExhaustive(same.resampling.instance = TRUE,
  maxit = NA_integer_, max.features = NA_integer_, tune.threshold = FALSE,
  tune.threshold.args = list(), log.fun = NULL)
```

```
makeFeatSelControlGA(same.resampling.instance = TRUE, impute.val = NULL,
  maxit = NA_integer_, max.features = NA_integer_, comma = FALSE,
  mu = 10L, lambda, crossover.rate = 0.5, mutation.rate = 0.05,
  tune.threshold = FALSE, tune.threshold.args = list(), log.fun = NULL)
```

```
makeFeatSelControlRandom(same.resampling.instance = TRUE, maxit = 100L,
  max.features = NA_integer_, prob = 0.5, tune.threshold = FALSE,
  tune.threshold.args = list(), log.fun = NULL)
```

```
makeFeatSelControlSequential(same.resampling.instance = TRUE,
  impute.val = NULL, method, alpha = 0.01, beta = -0.001,
  maxit = NA_integer_, max.features = NA_integer_, tune.threshold = FALSE,
  tune.threshold.args = list(), log.fun = NULL)
```

Arguments

<code>same.resampling.instance</code>	[logical(1)] Should the same resampling instance be used for all evaluations to reduce variance? Default is TRUE.
<code>maxit</code>	[integer(1)] Maximal number of iterations. Note, that this is usually not equal to the number of function evaluations.
<code>max.features</code>	[integer(1)] Maximal number of features.

<code>tune.threshold</code>	[logical(1)] Should the threshold be tuned for the measure at hand, after each feature set evaluation, via <code>tuneThreshold</code> ? Only works for classification if the predict type is “prob”. Default is FALSE.
<code>tune.threshold.args</code>	[list] Further arguments for threshold tuning that are passed down to <code>tuneThreshold</code> . Default is none.
<code>log.fun</code>	[function NULL] Function used for logging. If set to NULL, the internal default will be used. Otherwise a function with arguments <code>learner</code> , <code>resampling</code> , <code>measures</code> , <code>par.set</code> , <code>control</code> , <code>opt.path</code> , <code>dob</code> , <code>x</code> , <code>y</code> , <code>remove.nas</code> , and <code>stage</code> is expected. The default displays the performance measures, the time needed for evaluating, the currently used memory and the max memory ever used before (the latter two both taken from <code>gc</code>). See the implementation for details.
<code>impute.val</code>	[numeric] If something goes wrong during optimization (e.g. the learner crashes), this value is fed back to the tuner, so the tuning algorithm does not abort. It is not stored in the optimization path, an NA and a corresponding error message are logged instead. Note that this value is later multiplied by -1 for maximization measures internally, so you need to enter a larger positive value for maximization here as well. Default is the worst obtainable value of the performance measure you optimize for when you aggregate by mean value, or Inf instead. For multi-criteria optimization pass a vector of imputation values, one for each of your measures, in the same order as your measures.
<code>comma</code>	[logical(1)] Parameter of the GA feature selection, indicating whether to use a (μ , λ) or ($\mu + \lambda$) GA. The default is FALSE.
<code>mu</code>	[integer(1)] Parameter of the GA feature selection. Size of the parent population.
<code>lambda</code>	[integer(1)] Parameter of the GA feature selection. Size of the children population (should be smaller or equal to μ).
<code>crossover.rate</code>	[numeric(1)] Parameter of the GA feature selection. Probability of choosing a bit from the first parent within the crossover mutation.
<code>mutation.rate</code>	[numeric(1)] Parameter of the GA feature selection. Probability of flipping a feature bit, i.e. switch between selecting / deselecting a feature.
<code>prob</code>	[numeric(1)] Parameter of the random feature selection. Probability of choosing a feature.
<code>method</code>	[character(1)] Parameter of the sequential feature selection. A character representing the method. Possible values are <code>sfs</code> (forward search), <code>sbs</code> (backward search), <code>sffs</code> (floating forward search) and <code>sfbs</code> (floating backward search).

alpha	[numeric(1)] Parameter of the sequential feature selection. Minimal required value of improvement difference for a forward / adding step. Default is 0.01.
beta	[numeric(1)] Parameter of the sequential feature selection. Minimal required value of improvement difference for a backward / removing step. Negative values imply that you allow a slight decrease for the removal of a feature. Default is -0.001.

Value

[FeatSelControl](#) . The specific subclass is one of [FeatSelControlExhaustive](#), [FeatSelControlRandom](#), [FeatSelControlSequential](#), [FeatSelControlGA](#).

See Also

Other featsel: [analyzeFeatSelResult](#), [getFeatSelResult](#), [makeFeatSelWrapper](#), [selectFeatures](#)

FeatSelResult	<i>Result of feature selection.</i>
---------------	-------------------------------------

Description

Container for results of feature selection. Contains the obtained features, their performance values and the optimization path which lead there.

You can visualize it using [analyzeFeatSelResult](#).

Details

Object members:

learner [[Learner](#)] Learner that was optimized.

control [[FeatSelControl](#)] Control object from feature selection.

x [character] Vector of feature names identified as optimal.

y [numeric] Performance values for optimal x.

threshold [numeric] Vector of finally found and used thresholds if `tune.threshold` was enabled in [FeatSelControl](#), otherwise not present and hence NULL.

opt.path [[OptPath](#)] Optimization path which lead to x.

filterFeatures	<i>Filter features by thresholding filter values.</i>
----------------	---

Description

First, calls [generateFilterValuesData](#). Features are then selected via select and val.

Usage

```
filterFeatures(task, method = "rf.importance", fval = NULL, perc = NULL,
  abs = NULL, threshold = NULL, mandatory.feats = NULL, ...)
```

Arguments

task	[Task] The task.
method	[character(1)] See listFilterMethods . Default is “rf.importance”.
fval	[FilterValues] Result of generateFilterValuesData . If you pass this, the filter values in the object are used for feature filtering. method and ... are ignored then. Default is NULL and not used.
perc	[numeric(1)] If set, select perc*100 top scoring features. Mutually exclusive with arguments abs and threshold.
abs	[numeric(1)] If set, select abs top scoring features. Mutually exclusive with arguments perc and threshold.
threshold	[numeric(1)] If set, select features whose score exceeds threshold. Mutually exclusive with arguments perc and abs.
mandatory.feats	[character] Mandatory features which are always included regardless of their scores
...	[any] Passed down to selected filter method.

Value

[Task](#) .

See Also

Other filter: [generateFilterValuesData](#), [getFilterValues](#), [getFilteredFeatures](#), [makeFilterWrapper](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#)

friedmanPostHocTestBMR

Perform a posthoc Friedman-Nemenyi test.

Description

Performs a `posthoc.friedman.nemenyi.test` for a `BenchmarkResult` and a selected measure. This means all pairwise comparisons of learners are performed. The null hypothesis of the post hoc test is that each pair of learners is equal. If the null hypothesis of the included ad hoc `friedman.test` can be rejected a `pairwise.htest` is returned. If not, the function returns the corresponding `friedman.test`

Usage

```
friedmanPostHocTestBMR(bmr, measure = NULL, p.value = 0.05,
  aggregation = "default")
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
measure	[Measure] Performance measure. Default is the first measure used in the benchmark experiment.
p.value	[<code>numeric(1)</code>] p-value for the tests. Default: 0.05
aggregation	[<code>character(1)</code>] “mean” or “default”. See getBMRAggrPerformances for details on “default”.

Value

`pairwise.htest` : See [posthoc.friedman.nemenyi.test](#) for details. Additionally two components are added to the list:

f.rejnull [`logical(1)`] Whether the according `friedman.test` rejects the Null hypothesis at the selected p.value

crit.difference [`list(2)`] Minimal difference the mean ranks of two learners need to have in order to be significantly different

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

Examples

```
# see benchmark
```

friedmanTestBMR	<i>Perform overall Friedman test for a BenchmarkResult.</i>
-----------------	---

Description

Performs a `friedman.test` for a selected measure. The null hypothesis is that apart from an effect of the different [Task], the location parameter (aggregated performance-measure) is the same for each `Learner`.

Usage

```
friedmanTestBMR(bmr, measure = NULL, aggregation = "default")
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
measure	[Measure] Performance measure. Default is the first measure used in the benchmark experiment.
aggregation	[character(1)] “mean” or “default”. See getBMRAggrPerformances for details on “default”.

Value

htest : See `friedman.test` for details.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

Examples

```
# see benchmark
```

generateCalibrationData

Generate classifier calibration data.

Description

A calibrated classifier is one where the predicted probability of a class closely matches the rate at which that class occurs, e.g. for data points which are assigned a predicted probability of class A of .8, approximately 80 percent of such points should belong to class A if the classifier is well calibrated. This is estimated empirically by grouping data points with similar predicted probabilities for each class, and plotting the rate of each class within each bin against the predicted probability bins.

Usage

```
generateCalibrationData(obj, breaks = "Sturges", groups = NULL,
  task.id = NULL)
```

Arguments

obj	[(list of) Prediction (list of) ResampleResult BenchmarkResult] Single prediction object, list of them, single resample result, list of them, or a benchmark result. In case of a list probably produced by different learners you want to compare, then name the list with the names you want to see in the plots, probably learner shortnames or ids.
breaks	[character(1) numeric] If character(1), the algorithm to use in generating probability bins. See hist for details. If numeric, the cut points for the bins. Default is “Sturges”.
groups	[integer(1)] The number of bins to construct. If specified, breaks is ignored. Default is NULL.
task.id	[character(1)] Selected task in BenchmarkResult to do plots for, ignored otherwise. Default is first task.

Value

CalibrationData . A list containing:

proportion	[data.frame] with columns: <ul style="list-style-type: none"> • Learner Name of learner. • bin Bins calculated according to the breaks or groups argument. • Class Class labels (for binary classification only the positive class). • Proportion Proportion of observations from class Class among all observations with posterior probabilities of class Class within the interval given in bin.
------------	--

data	[data.frame] with columns: <ul style="list-style-type: none"> • Learner Name of learner. • truth True class label. • Class Class labels (for binary classification only the positive class). • Probability Predicted posterior probability of Class. • bin Bin corresponding to Probability.
task	[TaskDesc] Task description.

References

Vuk, Miha, and Curk, Tomaz. "ROC Curve, Lift Chart, and Calibration Plot." Metodoloski zvezki. Vol. 3. No. 1 (2006): 89-108.

See Also

Other calibration: [plotCalibration](#)

Other generate_plot_data: [generateCritDifferencesData](#), [generateFilterValuesData](#), [generateLearningCurveData](#), [generatePartialPredictionData](#), [generateThreshVsPerfData](#), [getFilterValues](#)

generateCritDifferencesData

Generate data for critical-differences plot.

Description

Generates data that can be used to plot a critical differences plot. Computes the critical differences according to either the "Bonferroni-Dunn" test or the "Nemenyi" test.

"Bonferroni-Dunn" usually yields higher power as it does not compare all algorithms to each other, but all algorithms to a baseline instead.

Learners are drawn on the y-axis according to their average rank.

For test = "nemenyi" a bar is drawn, connecting all groups of not significantly different learners.

For test = "bd" an interval is drawn around the algorithm selected as baseline. All learners within this interval are not significantly different from the baseline.

Calculation:

$$CD = q_{\alpha} \sqrt{\left(\frac{k(k+1)}{6N}\right)}$$

Where q_{α} is based on the studentized range statistic. See references for details.

Usage

```
generateCritDifferencesData(bmr, measure = NULL, p.value = 0.05,
  baseline = NULL, test = "bd")
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
measure	[Measure] Performance measure. Default is the first measure used in the benchmark experiment.
p.value	[numeric(1)] P-value for the critical difference. Default: 0.05
baseline	[character(1)]: [learner.id] Select a learner.id as baseline for the test = "bd" ("Bonferroni-Dunn") critical differences diagram. The critical difference Interval will then be positioned around this learner. Defaults to best performing algorithm. For test = "nemenyi", no baseline is needed as it performs all pairwise comparisons.
test	[character(1)] Test for which the critical differences are computed. "bd" for the Bonferroni-Dunn Test, which is comparing all classifiers to a baseline, thus performing a comparison of one classifier to all others. Algorithms not connected by a single line are statistically different. then the baseline. "nemenyi" for the posthoc.friedman.nemenyi.test which is comparing all classifiers to each other. The null hypothesis that there is a difference between the classifiers can not be rejected for all classifiers that have a single grey bar connecting them.

Value

critDifferencesData . List containing:

data	[data.frame] containing the info for the descriptive part of the plot
friedman.nemenyi.test	[list] of class pairwise.htest contains the calculated posthoc.friedman.nemenyi.test
cd.info	[list] containing info on the critical difference and its positioning
baseline	baseline chosen for plotting
p.value	p.value used for the posthoc.friedman.nemenyi.test and for computation of the critical difference

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

Other generate_plot_data: [generateCalibrationData](#), [generateFilterValuesData](#), [generateLearningCurveData](#), [generatePartialPredictionData](#), [generateThreshVsPerfData](#), [getFilterValues](#)

```
generateFilterValuesData
```

Calculates feature filter values.

Description

Calculates numerical filter values for features. For a list of features, use [listFilterMethods](#).

Usage

```
generateFilterValuesData(task, method = "rf.importance",
  nselect = getTaskNFeats(task), ...)
```

Arguments

task	[Task] The task.
method	[character] Filter method(s), see above. Default is “rf.importance”.
nselect	[integer(1)] Number of scores to request. Scores are getting calculated for all features per default.
...	[any] Passed down to selected method.

Value

FilterValues . A list containing:

task.desc	[TaskDesc] Task description.
data	[data.frame] with columns: <ul style="list-style-type: none"> • name Name of feature. • type Feature column type. • A column for each method with the feature importance values.

See Also

Other filter: [filterFeatures](#), [getFilterValues](#), [getFilteredFeatures](#), [makeFilterWrapper](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#)

Other generate_plot_data: [generateCalibrationData](#), [generateCritDifferencesData](#), [generateLearningCurveData](#), [generatePartialPredictionData](#), [generateThreshVsPerfData](#), [getFilterValues](#)

generateLearningCurveData

Generates a learning curve.

Description

Observe how the performance changes with an increasing number of observations.

Usage

```
generateLearningCurveData(learners, task, resampling = NULL,
  percs = seq(0.1, 1, by = 0.1), measures, stratify = FALSE,
  show.info = getMlrOption("show.info"))
```

Arguments

learners	[(list of) Learner] Learning algorithms which should be compared.
task	[Task] The task.
resampling	[ResampleDesc ResampleInstance] Resampling strategy to evaluate the performance measure. If no strategy is given a default "Holdout" will be performed.
percs	[numeric] Vector of percentages to be drawn from the training split. These values represent the x-axis. Internally makeDownsampleWrapper is used in combination with benchmark . Thus for each percentage a different set of observations is drawn resulting in noisy performance measures as the quality of the sample can differ.
measures	[(list of) Measure] Performance measures to generate learning curves for, representing the y-axis.
stratify	[logical(1)] Only for classification: Should the downsampled data be stratified according to the target classes?
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .

Value

LearningCurveData . A list containing:

task	[Task] The task.
measures	[(list of) Measure] Performance measures.
data	[data.frame] with columns:

- learner Names of learners.
- percentage Percentages drawn from the training split.
- One column for each [Measure](#) passed to [generateLearningCurveData](#).

See Also

Other generate_plot_data: [generateCalibrationData](#), [generateCritDifferencesData](#), [generateFilterValuesData](#), [generatePartialPredictionData](#), [generateThreshVsPerfData](#), [getFilterValues](#)

Other learning_curve: [plotLearningCurveGGVIS](#), [plotLearningCurve](#)

Examples

```
r = generateLearningCurveData(list("classif.rpart", "classif.knn"),
task = sonar.task, percs = seq(0.2, 1, by = 0.2),
measures = list(tp, fp, tn, fn), resampling = makeResampleDesc(method = "Subsample", iters = 5),
show.info = FALSE)
plotLearningCurve(r)
```

generatePartialPredictionData

Generate partial predictions.

Description

Estimate how the learned prediction function is affected by one or more features. For a learned function $f(x)$ where x is partitioned into x_s and x_c , the partial dependence of f on x_s can be summarized by averaging over x_c and setting x_s to a range of values of interest, estimating $E_{x_c}(f(x_s, x_c))$. The conditional expectation of f at observation i is estimated similarly. Additionally, partial derivatives of the marginalized function w.r.t. the features can be computed.

Usage

```
generatePartialPredictionData(obj, input, features, interaction = FALSE,
derivative = FALSE, individual = FALSE, center = NULL, fun = mean,
bounds = c(qnorm(0.025), qnorm(0.975)), resample = "none", fmin, fmax,
gridsize = 10L, ...)
```

Arguments

obj	[WrappedModel] Result of train .
input	[data.frame Task] Input data.
features	[character] A vector of feature names contained in the training data. If not specified all features in the input will be used.

interaction	<p>[logical(1)]</p> <p>Whether the features should be interacted or not. If TRUE then the Cartesian product of the prediction grid for each feature is taken, and the partial prediction at each unique combination of values of the features is estimated. Note that if the length of features is greater than two, plotPartialPrediction and plotPartialPredictionGGVIS cannot be used. If FALSE each feature is considered separately. In this case features can be much longer than two. Default is FALSE.</p>
derivative	<p>[logical(1)]</p> <p>Whether or not the partial derivative of the learned function with respect to the features should be estimated. If TRUE interaction must be FALSE. The partial derivative of individual observations may be estimated. Note that computation time increases as the learned prediction function is evaluated at <code>gridsize</code> points * the number of points required to estimate the partial derivative. Additional arguments may be passed to grad (for regression or survival tasks) or jacobian (for classification tasks). Note that functions which are not smooth may result in estimated derivatives of 0 (for points where the function does not change within +/- epsilon) or estimates trending towards +/- infinity (at discontinuities). Default is FALSE.</p>
individual	<p>[logical(1)]</p> <p>Whether to plot the individual conditional expectation curves rather than the aggregated curve, i.e., rather than aggregating (using <code>fun</code>) the partial predictions of features, plot the partial predictions of all observations in data across all values of the features. The algorithm is developed in Goldstein, Kapelner, Bleich, and Pitkin (2015). Default is FALSE.</p>
center	<p>[list]</p> <p>A named list containing the fixed values of the features used to calculate an individual partial prediction which is then subtracted from each individual partial prediction made across the prediction grid created for the features: centering the individual partial prediction lines to make them more interpretable. This argument is ignored if <code>individual != TRUE</code>. Default is NULL.</p>
fun	<p>[function]</p> <p>For regression, a function that accepts a numeric vector and returns either a single number such as a measure of location such as the mean, or three numbers, which give a lower bound, a measure of location, and an upper bound. Note if three numbers are returned they must be in this order. For classification with <code>predict.type = "prob"</code> the function must accept a numeric matrix with the number of columns equal to the number of class levels of the target. For classification with <code>predict.type = "response"</code> (the default) the function must accept a character vector and output a numeric vector with length equal to the number of classes in the target feature. The default is the mean, unless <code>obj</code> is classification with <code>predict.type = "response"</code> in which case the default is the proportion of observations predicted to be in each class.</p>
bounds	<p>[numeric(2)]</p> <p>The value (lower, upper) the estimated standard error is multiplied by to estimate the bound on a confidence region for a partial prediction. Ignored if <code>predict.type != "se"</code> for the learner. Default is the 2.5 and 97.5 quantiles (-1.96, 1.96) of the Gaussian distribution.</p>

resample	[character(1)] Defines how the prediction grid for each feature is created. If “bootstrap” then values are sampled with replacement from the training data. If “subsample” then values are sampled without replacement from the training data. If “none” an evenly spaced grid between either the empirical minimum and maximum, or the minimum and maximum defined by fmin and fmax, is created. Default is “none”.
fmin	[numeric] The minimum value that each element of features can take. This argument is only applicable if resample = NULL and when the empirical minimum is higher than the theoretical minimum for a given feature. This only applies to numeric features and a NA should be inserted into the vector if the corresponding feature is a factor. Default is the empirical minimum of each numeric feature and NA for factor features.
fmax	[numeric] The maximum value that each element of features can take. This argument is only applicable if resample = “none” and when the empirical maximum is lower than the theoretical maximum for a given feature. This only applies to numeric features and a NA should be inserted into the vector if the corresponding feature is a factor. Default is the empirical maximum of each numeric feature and NA for factor features.
gridsize	[integer(1)] The length of the prediction grid created for each feature. If resample = “bootstrap” or resample = “subsample” then this defines the number of (possibly non-unique) values resampled. If resample = NULL it defines the length of the evenly spaced grid created.
...	additional arguments to be passed to predict .

Value

PartialPredictionData . A named list, which contains the partial predictions, input data, target, features, task description, and other arguments controlling the type of partial predictions made.

Object members:

data	[data.frame] Has columns for the prediction: one column for regression and survival analysis, and a column for class and the predicted probability for classification as well as a column for each element of features. If individual = TRUE then there is an additional column idx which gives the index of the data that each prediction corresponds to.
task.desc	[TaskDesc] Task description.
target	Target feature for regression, target feature levels for classification, survival and event indicator for survival.
features	[character] Features argument input.

interaction	[logical(1)] Whether or not the features were interacted (i.e. conditioning).
derivative	[logical(1)] Whether or not the partial derivative was estimated.
individual	[logical(1)] Whether the partial predictions were aggregated or the individual curves are retained.
center	[logical(1)] If individual == TRUE whether the partial prediction at the values of the features specified was subtracted from the individual partial predictions. Only displayed if individual == TRUE.

References

Goldstein, Alex, Adam Kapelner, Justin Bleich, and Emil Pitkin. "Peeking inside the black box: Visualizing statistical learning with plots of individual conditional expectation." *Journal of Computational and Graphical Statistics*. Vol. 24, No. 1 (2015): 44-65.

Friedman, Jerome. "Greedy Function Approximation: A Gradient Boosting Machine." *The Annals of Statistics*. Vol. 29, No. 5 (2001): 1189-1232.

See Also

Other generate_plot_data: [generateCalibrationData](#), [generateCritDifferencesData](#), [generateFilterValuesData](#), [generateLearningCurveData](#), [generateThreshVsPerfData](#), [getFilterValues](#)

Other partial_prediction: [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#)

Examples

```
lrn = makeLearner("regr.rpart")
fit = train(lrn, bh.task)
pd = generatePartialPredictionData(fit, bh.task, "lstat")
plotPartialPrediction(pd)

lrn = makeLearner("classif.rpart", predict.type = "prob")
fit = train(lrn, iris.task)
pd = generatePartialPredictionData(fit, iris.task, "Petal.Width")
plotPartialPrediction(pd)
```

generateThreshVsPerfData

Generate threshold vs. performance(s) for 2-class classification.

Description

Generates data on threshold vs. performance(s) for 2-class classification that can be used for plotting.

Usage

```
generateThreshVsPerfData(obj, measures, gridsize = 100L, aggregate = TRUE,
  task.id = NULL)
```

Arguments

obj	[(list of) Prediction (list of) ResampleResult BenchmarkResult] Single prediction object, list of them, single resample result, list of them, or a benchmark result. In case of a list probably produced by different learners you want to compare, then name the list with the names you want to see in the plots, probably learner shortnames or ids.
measures	Measure list of Measure Performance measure(s) to evaluate. Default is the default measure for the task, see here getDefaultMeasure .
gridsize	[integer(1)] Grid resolution for x-axis (threshold). Default is 100.
aggregate	[logical(1)] Whether to aggregate ResamplePredictions or to plot the performance of each iteration separately. Default is TRUE.
task.id	[character(1)] Selected task in BenchmarkResult to do plots for, ignored otherwise. Default is first task.

Value

ThreshVsPerfData . A named list containing the measured performance across the threshold grid, the measures, and whether the performance estimates were aggregated (only applicable for (list of) [ResampleResults](#)).

See Also

Other generate_plot_data: [generateCalibrationData](#), [generateCritDifferencesData](#), [generateFilterValuesData](#), [generateLearningCurveData](#), [generatePartialPredictionData](#), [getFilterValues](#)

Other thresh_vs_perf: [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

```
getBMRAggrPerformances
```

Extract the aggregated performance values from a benchmark result.

Description

Either a list of lists of “aggr” numeric vectors, as returned by [resample](#), or these objects are rbind-ed with extra columns “task.id” and “learner.id”.

Usage

```
getBMRAggrPerformances(bmr, task.ids = NULL, learner.ids = NULL,
  as.df = FALSE)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
task.ids	[character(1)] Restrict result to certain tasks. Default is all.
learner.ids	[character(1)] Restrict result to certain learners. Default is all.
as.df	[character(1)] Return one data.frame as result - or a list of lists of objects?. Default is FALSE.

Value

list | data.frame . See above.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRFeatSelResults *Extract the feature selection results from a benchmark result.*

Description

Returns a nested list of [FeatSelResults](#). The first level of nesting is by data set, the second by learner, the third for the benchmark resampling iterations. If as.df is TRUE, a data frame with “task.id”, “learner.id”, the resample iteration and the selected features is returned.

Note that if more than one feature is selected and a data frame is requested, there will be multiple rows for the same dataset-learner-iteration; one for each selected feature.

Usage

```
getBMRFeatSelResults(bmr, task.ids = NULL, learner.ids = NULL,
  as.df = FALSE)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
task.ids	[character(1)] Restrict result to certain tasks. Default is all.
learner.ids	[character(1)] Restrict result to certain learners. Default is all.
as.df	[character(1)] Return one data.frame as result - or a list of lists of objects?. Default is FALSE.

Value

list | data.frame . See above.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFILTEREDFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSUMMARY](#), [plotCritDifferences](#)

getBMRFILTEREDFeatures

Extract the feature selection results from a benchmark result.

Description

Returns a nested list of characters The first level of nesting is by data set, the second by learner, the third for the benchmark resampling iterations. The list at the lowest level is the list of selected features. If as.df is TRUE, a data frame with “task.id”, “learner.id”, the resample iteration and the selected features is returned.

Note that if more than one feature is selected and a data frame is requested, there will be multiple rows for the same dataset-learner-iteration; one for each selected feature.

Usage

```
getBMRFILTEREDFeatures(bmr, task.ids = NULL, learner.ids = NULL,
  as.df = FALSE)
```


Arguments

bmr	[BenchmarkResult] Benchmark result.
task.ids	[character(1)] Restrict result to certain tasks. Default is all.
learner.ids	[character(1)] Restrict result to certain learners. Default is all.
as.df	[character(1)] Return one data.frame as result - or a list of lists of objects?. Default is FALSE.

Value

list | data.frame . See above.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRLearnerIds	<i>Return learner ids used in benchmark.</i>
------------------	--

Description

Gets the IDs of the learners used in a benchmark experiment.

Usage

```
getBMRLearnerIds(bmr)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
-----	--

Value

character .

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRLearners

Return learners used in benchmark.

Description

Gets the learners used in a benchmark experiment.

Usage

```
getBMRLearners(bmr)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
-----	--

Value

list .

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRLearnerShortNames

Return learner short.names used in benchmark.

Description

Gets the learner short.names of the learners used in a benchmark experiment.

Usage

```
getBMRLearnerShortNames(bmr)
```

Arguments

bmr [\[BenchmarkResult\]](#)
Benchmark result.

Value

character .

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRMeasureIds	<i>Return measures IDs used in benchmark.</i>
------------------	---

Description

Gets the IDs of the measures used in a benchmark experiment.

Usage

```
getBMRMeasureIds(bmr)
```

Arguments

bmr [\[BenchmarkResult\]](#)
Benchmark result.

Value

list . See above.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRMeasures	<i>Return measures used in benchmark.</i>
----------------	---

Description

Gets the measures used in a benchmark experiment.

Usage

```
getBMRMeasures(bmr)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
-----	--

Value

list . See above.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRModels	<i>Extract all models from benchmark result.</i>
--------------	--

Description

A list of lists containing all [WrappedModels](#) trained in the benchmark experiment.

If models is FALSE in the call to [benchmark](#), the function will return NULL.

Usage

```
getBMRModels(bmr, task.ids = NULL, learner.ids = NULL)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
task.ids	[character(1)] Restrict result to certain tasks. Default is all.
learner.ids	[character(1)] Restrict result to certain learners. Default is all.

Value

list .

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRPerformances	<i>Extract the test performance values from a benchmark result.</i>
--------------------	---

Description

Either a list of lists of “measure.test” data.frames, as returned by [resample](#), or these objects are rbind-ed with extra columns “task.id” and “learner.id”.

Usage

```
getBMRPerformances(bmr, task.ids = NULL, learner.ids = NULL,
  as.df = FALSE)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
task.ids	[character(1)] Restrict result to certain tasks. Default is all.
learner.ids	[character(1)] Restrict result to certain learners. Default is all.
as.df	[character(1)] Return one data.frame as result - or a list of lists of objects?. Default is FALSE.

Value

list | data.frame . See above.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRPredictions	<i>Extract the predictions from a benchmark result.</i>
-------------------	---

Description

Either a list of lists of [ResamplePrediction](#) objects, as returned by [resample](#), or these objects are rbind-ed with extra columns “task.id” and “learner.id”.

If `predict.type` is “prob”, the probabilities for each class are returned in addition to the response.

If `keep.pred` is FALSE in the call to [benchmark](#), the function will return NULL.

Usage

```
getBMRPredictions(bmr, task.ids = NULL, learner.ids = NULL, as.df = FALSE)
```

Arguments

<code>bmr</code>	[BenchmarkResult] Benchmark result.
<code>task.ids</code>	[character(1)] Restrict result to certain tasks. Default is all.
<code>learner.ids</code>	[character(1)] Restrict result to certain learners. Default is all.
<code>as.df</code>	[character(1)] Return one data.frame as result - or a list of lists of objects?. Default is FALSE.

Value

`list | data.frame` . See above.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRTaskIds	<i>Return task ids used in benchmark.</i>
---------------	---

Description

Gets the task IDs used in a benchmark experiment.

Usage

```
getBMRTaskIds(bmr)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
-----	--

Value

character .

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getBMRTuneResults	<i>Extract the tuning results from a benchmark result.</i>
-------------------	--

Description

Returns a nested list of [TuneResults](#). The first level of nesting is by data set, the second by learner, the third for the benchmark resampling iterations. If `as.df` is TRUE, a data frame with the “task.id”, “learner.id”, the resample iteration, the parameter values and the performances is returned.

Usage

```
getBMRTuneResults(bmr, task.ids = NULL, learner.ids = NULL, as.df = FALSE)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
task.ids	[character(1)] Restrict result to certain tasks. Default is all.
learner.ids	[character(1)] Restrict result to certain learners. Default is all.
as.df	[character(1)] Return one data.frame as result - or a list of lists of objects?. Default is FALSE.

Value

list | data.frame . See above.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

getCaretParamSet

Get tuning parameters from a learner of the caret R-package.

Description

Constructs a grid of tuning parameters from a learner of the caret R-package. These values are then converted into a list of non-tunable parameters (`par.vals`) and a tunable [ParamSet](#) (`par.set`), which can be used by [tuneParams](#) for tuning the learner. Numerical parameters will either be specified by their lower and upper bounds or they will be discretized into specific values.

Usage

```
getCaretParamSet(learner, length = 3L, task, discretize = TRUE)
```

Arguments

learner	[character(1)] The name of the learner from caret (cf. http://topepo.github.io/caret/modellist.html). Note that the names in caret often differ from the ones in mlr.
length	[integer(1)] A length / precision parameter which is used by caret for generating the grid of tuning parameters. caret generates either as many values per tuning parameter / dimension as defined by length or only a single value (in case of non-tunable <code>par.vals</code>).

task	[Task] Learning task, which might be requested for creating the tuning grid.
discretize	[logical(1)] Should the numerical parameters be discretized? Alternatively, they will be defined by their lower and upper bounds. The default is TRUE.

Value

list(2) . A list of parameters:

- par.vals contains a list of all constant tuning parameters
- par.set is a [ParamSet](#), containing all the configurable tuning parameters

Examples

```
library(caret)
classifTask = makeClassifTask(data = iris, target = "Species")

# (1) classification (random forest) with discretized parameters
getCaretParamSet("rf", length = 9L, task = classifTask, discretize = TRUE)

# (2) regression (gradient boosting machine) without discretized parameters
library(mlbench)
data(BostonHousing)
regrTask = makeRegrTask(data = BostonHousing, target = "medv")
getCaretParamSet("gbm", length = 9L, task = regrTask, discretize = FALSE)
```

getClassWeightParam *Get the class weight parameter of a learner.*

Description

Gets the class weight parameter of a learner.

Usage

```
getClassWeightParam(learner)
```

Arguments

learner [[Learner](#) | character(1)]
The learner. If you pass a string the learner will be created via [makeLearner](#).

Value

numeric [LearnerParam](#) : A numeric parameter object, containing the class weight parameter of the given learner.

See Also

Other learner: [LearnerProperties](#), [getHyperPars](#), [getParamSet](#), [makeLearner](#), [removeHyperPars](#), [setHyperPars](#), [setId](#), [setPredictThreshold](#), [setPredictType](#)

getConfMatrix	<i>Confusion matrix.</i>
---------------	--------------------------

Description

Calculates confusion matrix for (possibly resampled) prediction. Rows indicate true classes, columns predicted classes.

Usage

```
getConfMatrix(pred, relative = FALSE)
```

Arguments

pred	[Prediction] Prediction object.
relative	[logical(1)] If TRUE rows are normalized to show relative frequencies. Default is FALSE.

Value

matrix . A confusion matrix.

See Also

[predict.WrappedModel](#)

Examples

```
# get confusion matrix after simple manual prediction
allinds = 1:150
train = sample(allinds, 75)
test = setdiff(allinds, train)
mod = train("classif.lda", iris.task, subset = train)
pred = predict(mod, iris.task, subset = test)
print(getConfMatrix(pred))
print(getConfMatrix(pred, relative = TRUE))

# now after cross-validation
r = crossval("classif.lda", iris.task, iters = 2L)
print(getConfMatrix(r$pred))
```

getDefaultMeasure	<i>Get default measure.</i>
-------------------	-----------------------------

Description

Get the default measure for a task type, task, task description or a learner. Currently these are:

classif	mmce
regr	mse
cluster	db
surv	cindex
costsens	mcp
multilabel	hamloss

Usage

```
getDefaultMeasure(x)
```

Arguments

x	[character(1) Task TaskDesc Learner] Task type, task, task description or a learner.
---	--

Value

[Measure](#) .

getFailureModelMsg	<i>Return error message of FailureModel.</i>
--------------------	--

Description

Such a model is created when one sets the corresponding option in [configureMlr](#). If no failure occurred, NA is returned.

For complex wrappers this getter returns the first error message encountered in ANY model that failed.

Usage

```
getFailureModelMsg(model)
```

Arguments

model	[WrappedModel] The model.
-------	--

Value

character(1) .

getFeatSelResult	Returns the selected feature set and optimization path after training.
------------------	--

Description

Returns the selected feature set and optimization path after training.

Usage

```
getFeatSelResult(object)
```

Arguments

object	[WrappedModel] Trained Model created with makeFeatSelWrapper .
--------	---

Value

[FeatSelResult](#) .

See Also

Other featsel: [FeatSelControl](#), [analyzeFeatSelResult](#), [makeFeatSelWrapper](#), [selectFeatures](#)

getFilteredFeatures	Returns the filtered features.
---------------------	--------------------------------

Description

Returns the filtered features.

Usage

```
getFilteredFeatures(model)
```

Arguments

model	[WrappedModel] Trained Model created with makeFilterWrapper .
-------	--

Value

character .

See Also

Other filter: [filterFeatures](#), [generateFilterValuesData](#), [getFilterValues](#), [makeFilterWrapper](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#)

getFilterValues	<i>Calculates feature filter values.</i>
-----------------	--

Description

Calculates numerical filter values for features. For a list of features, use [listFilterMethods](#).

Usage

```
getFilterValues(task, method = "rf.importance",
  nselect = getTaskNFeats(task), ...)
```

Arguments

task	[Task] The task.
method	[character(1)] Filter method, see above. Default is “rf.importance”.
nselect	[integer(1)] Number of scores to request. Scores are getting calculated for all features per default.
...	[any] Passed down to selected method.

Value

[FilterValues](#) .

Note

getFilterValues is deprecated in favor of [generateFilterValuesData](#).

See Also

Other filter: [filterFeatures](#), [generateFilterValuesData](#), [getFilteredFeatures](#), [makeFilterWrapper](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#)

Other filter: [filterFeatures](#), [generateFilterValuesData](#), [getFilteredFeatures](#), [makeFilterWrapper](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#)

Other generate_plot_data: [generateCalibrationData](#), [generateCritDifferencesData](#), [generateFilterValuesData](#), [generateLearningCurveData](#), [generatePartialPredictionData](#), [generateThreshVsPerfData](#)

```
getHomogeneousEnsembleModels
```

Deprecated, use getLearnerModel instead.

Description

Deprecated, use getLearnerModel instead.

Usage

```
getHomogeneousEnsembleModels(model, learner.models = FALSE)
```

Arguments

`model` *Deprecated.*

`learner.models` *Deprecated.*

```
getHyperPars
```

Get current parameter settings for a learner.

Description

Retrieves the current hyperparameter settings of a learner.

Usage

```
getHyperPars(learner, for.fun = c("train", "predict", "both"))
```

Arguments

`learner` [[Learner](#) | character(1)]
The learner. If you pass a string the learner will be created via [makeLearner](#).

`for.fun` [character(1)]
Restrict the returned settings to hyperparameters corresponding to when the are used (see [LearnerParam](#)). Must be a subset of: “train”, “predict” or “both”. Default is c(“train”, “predict”, “both”).

Value

`list` . A named list of values.

See Also

Other learner: [LearnerProperties](#), [getClassWeightParam](#), [getParamSet](#), [makeLearner](#), [removeHyperPars](#), [setHyperPars](#), [setId](#), [setPredictThreshold](#), [setPredictType](#)

getLearnerModel	<i>Get underlying R model of learner integrated into mlr.</i>
-----------------	---

Description

Get underlying R model of learner integrated into mlr.

Usage

```
getLearnerModel(model, more.unwrap = FALSE)
```

Arguments

model	[WrappedModel] The model, returned by e.g., train .
more.unwrap	[logical(1)] Some learners are not basic learners from R, but implemented in mlr as meta-techniques. Examples are everything that inherits from <code>HomogeneousEnsemble</code> . In these cases, the <code>learner.model</code> is often a list of mlr WrappedModels . This option allows to strip them further to basic R models. The option is simply ignored for basic learner models. Default is FALSE.

Value

any . A fitted model, depending the learner / wrapped package. E.g., a model of class [rpart](#) for learner “`classif.rpart`”.

getMlrOptions	<i>Returns a list of mlr's options.</i>
---------------	---

Description

Gets the options for mlr.

Usage

```
getMlrOptions()
```

Value

list .

See Also

Other configure: [configureMlr](#)

```
getMultilabelBinaryPerformances
```

Retrieve binary classification measures for multilabel classification predictions.

Description

Measures the quality of each binary label prediction w.r.t. some binary classification performance measure.

Usage

```
getMultilabelBinaryPerformances(pred, measures)
```

Arguments

pred	[Prediction] Multilabel Prediction object.
measures	[Measure list of Measure] Performance measure(s) to evaluate, must be applicable to binary classification performance. Default is mmce.

Value

named matrix . Performance value(s), column names are measure(s), row names are labels.

See Also

Other multilabel: [makeMultilabelBinaryRelevanceWrapper](#)

Examples

```
# see makeMultilabelBinaryRelevanceWrapper
```

```
getNestedTuneResultsOptPathDf
```

Get the opt.paths from each tuning step from the outer resampling.

Description

After you resampled a tuning wrapper (see [makeTuneWrapper](#)) with `resample(..., extract = getTuneResult)` this helper returns a data.frame with with all opt.paths combined by `rbind`. An additional column `iter` indicates to what resampling iteration the row belongs.

Usage

```
getNestedTuneResultsOptPathDf(r)
```


Arguments

`r` [\[ResampleResult\]](#)
The result of resampling of a tuning wrapper.

Value

`data.frame` . See above.

See Also

Other tune: [TuneControl](#), [getNestedTuneResultsX](#), [getTuneResult](#), [makeModelMultiplexerParamSet](#), [makeModelMultiplexer](#), [makeTuneWrapper](#), [tuneParams](#), [tuneThreshold](#)

Examples

```
# see example of makeTuneWrapper
```

`getNestedTuneResultsX` *Get the tuned hyperparameter settings from a nested tuning.*

Description

After you resampled a tuning wrapper (see [makeTuneWrapper](#)) with `resample(..., extract = getTuneResult)` this helper returns a `data.frame` with the the best found hyperparameter settings for each resampling iteration.

Usage

```
getNestedTuneResultsX(r)
```

Arguments

`r` [\[ResampleResult\]](#)
The result of resampling of a tuning wrapper.

Value

`data.frame` . One column for each tuned hyperparameter and one row for each outer resampling iteration.

See Also

Other tune: [TuneControl](#), [getNestedTuneResultsOptPathDf](#), [getTuneResult](#), [makeModelMultiplexerParamSet](#), [makeModelMultiplexer](#), [makeTuneWrapper](#), [tuneParams](#), [tuneThreshold](#)

Examples

```
# see example of makeTuneWrapper
```

getParamSet

Get a description of all possible parameter settings for a learner.

Description

Get a description of all possible parameter settings for a learner.

Usage

```
getParamSet(learner)
```

Arguments

learner [[Learner](#) | character(1)]
The learner. If you pass a string the learner will be created via [makeLearner](#).

Value

[ParamSet](#) .

See Also

Other learner: [LearnerProperties](#), [getClassWeightParam](#), [getHyperPars](#), [makeLearner](#), [removeHyperPars](#), [setHyperPars](#), [setId](#), [setPredictThreshold](#), [setPredictType](#)

getPredictionProbabilities

Get probabilities for some classes.

Description

Get probabilities for some classes.

Usage

```
getPredictionProbabilities(pred, cl)
```

Arguments

pred [[Prediction](#)]
Prediction object.

cl [character]
Names of classes. Default is either all classes for multi-class / multilabel problems or the positive class for binary classification.

Value

data.frame with numerical columns or a numerical vector if length of `cl` is 1. Order of columns is defined by `cl`.

See Also

Other predict: [asROCRPrediction](#), [getPredictionResponse](#), [plotViperCharts](#), [predict.WrappedModel](#), [setPredictThreshold](#), [setPredictType](#)

Examples

```
task = makeClassifTask(data = iris, target = "Species")
lrn = makeLearner("classif.lda", predict.type = "prob")
mod = train(lrn, task)
# predict probabilities
pred = predict(mod, newdata = iris)

# Get probabilities for all classes
head(getPredictionProbabilities(pred))

# Get probabilities for a subset of classes
head(getPredictionProbabilities(pred, c("setosa", "virginica")))
```

`getPredictionResponse` *Get response / truth from prediction object.*

Description

The following types are returned, depending on task type:

classif	factor
regr	numeric
se	numeric
cluster	integer
surv	numeric
multilabel	logical matrix, columns named with labels

Usage

```
getPredictionResponse(pred)

getPredictionSE(pred)

getPredictionTruth(pred)
```

Arguments

pred [[Prediction](#)]
Prediction object.

Value

See above.

See Also

Other predict: [asROCRPrediction](#), [getPredictionProbabilities](#), [plotViperCharts](#), [predict.WrappedModel](#), [setPredictThreshold](#), [setPredictType](#)

getProbabilities	<i>Deprecated, use getPredictionProbabilities instead.</i>
------------------	--

Description

Deprecated, use [getPredictionProbabilities](#) instead.

Usage

```
getProbabilities(pred, cl)
```

Arguments

pred Deprecated.
cl Deprecated.

getRRPredictions	<i>Get predictions from resample results.</i>
------------------	---

Description

Very simple getter.

Usage

```
getRRPredictions(res)
```

Arguments

res [[ResampleResult](#)]
The result of [resample](#) run with `keep.pred = TRUE`.

Value

ResamplePrediction .

See Also

Other resample: [ResamplePrediction](#), [ResampleResult](#), [makeResampleDesc](#), [makeResampleInstance](#), [resample](#)

getStackedBaseLearnerPredictions

Returns the predictions for each base learner.

Description

Returns the predictions for each base learner.

Usage

```
getStackedBaseLearnerPredictions(model, newdata = NULL)
```

Arguments

model	[WrappedModel] Wrapped model, result of train.
newdata	[data.frame] New observations, for which the predictions using the specified base learners should be returned. Default is NULL and extracts the base learner predictions that were made during the training.

Details

None.

getTaskClassLevels

Get the class levels for classification and multilabel tasks.

Description

NB: For multilabel, [getTaskTargetNames](#) and [getTaskClassLevels](#) actually return the same thing.

Usage

```
getTaskClassLevels(x)
```

Arguments

x [Task | TaskDesc]
Task or its description object.

Value

character .

See Also

Other task: [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskCosts	<i>Extract costs in task.</i>
--------------	-------------------------------

Description

Retuns “NULL” if the task is not of type “costsens”.

Usage

getTaskCosts(task, subset)

Arguments

task [CostSensTask]
The task.

subset [integer]
Selected cases. Default is all cases.

Value

matrix | NULL .

See Also

Other task: [getTaskClassLevels](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskData	<i>Extract data in task.</i>
-------------	------------------------------

Description

Useful in [trainLearner](#) when you add a learning machine to the package.

Usage

```
getTaskData(task, subset, features, target.extra = FALSE,  
  recode.target = "no")
```

Arguments

task	[Task] The task.
subset	[integer] Selected cases. Default is all cases.
features	[character] Selected features. Default is all.
target.extra	[logical(1)] Should target vector be returned separately? If not, a single data.frame including the target columns is returned, otherwise a list with the input data.frame and an extra vector or data.frame for the targets. Default is FALSE.
recode.target	[character(1)] Should target classes be recoded? Supported are binary classification and survival. Possible values for binary classification are "01", "-1+1" and "drop.levels". In the two latter cases the target vector is converted into a numeric vector. The positive class is coded as "+1" and the negative class either as "0" or "-1". "drop.levels" will remove empty factor levels in the target column. For survival, you may choose to recode the survival times to "left", "right" or "interval2" censored times using "lcens", "rcens" or "icens", respectively. See Surv for the format specification. Default for both binary classification and survival is "no" (do nothing).

Value

Either a data.frame or a list with data.frame data and vector target.

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

Examples

```
library("mlbench")
data(BreastCancer)

df = BreastCancer
df$Id = NULL
task = makeClassifTask(id = "BreastCancer", data = df, target = "Class", positive = "malignant")
head(getTaskData)
head(getTaskData(task, features = c("Cell.size", "Cell.shape"), recode.target = "-1+1"))
head(getTaskData(task, subset = 1:100, recode.target = "01"))
```

getTaskDescription	<i>Get a summarizing task description.</i>
--------------------	--

Description

Get a summarizing task description.

Usage

```
getTaskDescription(x)
```

Arguments

x [\[Task | TaskDesc\]](#)
Task or its description object.

Value

[TaskDesc](#) .

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskFeatureNames	<i>Get feature names of task.</i>
---------------------	-----------------------------------

Description

Target column name is not included.

Usage

```
getTaskFeatureNames(task)
```

Arguments

task	[Task]
	The task.

Value

character .

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskFormula	<i>Get formula of a task.</i>
----------------	-------------------------------

Description

This is usually simply “<target> ~ .”. For multilabel it is “<target_1> + ... + <target_k> ~ .”.

Usage

```
getTaskFormula(x, target = getTaskTargetNames(x), explicit.features = FALSE,  
  env = parent.frame())
```

Arguments

x	[Task TaskDesc] Task or its description object.
target	[character(1)] Left hand side of the formula. Default is defined by task x.
explicit.features	[logical(1)] Should the features (right hand side of the formula) be explicitly listed? Default is FALSE, i.e., they will be represented as ". ".
env	[environment] Environment of the formula. Default is parent.frame().

Value

formula .

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskId	<i>Get the id of the task.</i>
-----------	--------------------------------

Description

Get the id of the task.

Usage

getTaskId(x)

Arguments

x	[Task TaskDesc] Task or its description object.
---	--

Value

character(1) .

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskNFeats	<i>Get number of features in task.</i>
---------------	--

Description

Get number of features in task.

Usage

```
getTaskNFeats(x)
```

Arguments

x	[Task TaskDesc] Task or its description object.
---	--

Value

integer(1) .

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskSize	<i>Get number of observations in task.</i>
-------------	--

Description

Get number of observations in task.

Usage

```
getTaskSize(x)
```

Arguments

x	[Task TaskDesc] Task or its description object.
---	--

Value

integer(1) .

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskTargetNames	<i>Get the name(s) of the target column(s).</i>
--------------------	---

Description

NB: For multilabel, [getTaskTargetNames](#) and [getTaskClassLevels](#) actually return the same thing.

Usage

```
getTaskTargetNames(x)
```

Arguments

x	[Task TaskDesc] Task or its description object.
---	--

Value

character .

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargets](#), [getTaskType](#), [subsetTask](#)

getTaskTargets	<i>Get target data of task.</i>
----------------	---------------------------------

Description

Get target data of task.

Usage

```
getTaskTargets(task, recode.target = "no")
```

Arguments

task	[Task] The task.
recode.target	[character(1)] Should target classes be recoded? Only for binary classification. Possible are “no” (do nothing), “01”, and “-1+1”. In the two latter cases the target vector is converted into a numeric vector. The positive class is coded as +1 and the negative class either as 0 or -1. Default is “no”.

Value

A factor for classification or a numeric for regression, a data.frame of logical columns for multilabel.

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskType](#), [subsetTask](#)

Examples

```
task = makeClassifTask(data = iris, target = "Species")
getTaskTargets(task)
```

getTaskType	<i>Get the type of the task.</i>
-------------	----------------------------------

Description

Get the type of the task.

Usage

```
getTaskType(x)
```

Arguments

x	[Task TaskDesc] Task or its description object.
---	--

Value

[character\(1\)](#) .

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [subsetTask](#)

getTuneResult	Returns the optimal hyperparameters and optimization path after training.
---------------	---

Description

Returns the optimal hyperparameters and optimization path after training.

Usage

getTuneResult(object)

Arguments

object [\[WrappedModel\]](#)
Trained Model created with [makeTuneWrapper](#).

Value

[TuneResult](#) .

See Also

Other tune: [TuneControl](#), [getNestedTuneResultsOptPathDf](#), [getNestedTuneResultsX](#), [makeModelMultiplexerParams](#), [makeModelMultiplexer](#), [makeTuneWrapper](#), [tuneParams](#), [tuneThreshold](#)

hasProperties	Deprecated, use hasLearnerProperties instead.
---------------	---

Description

Deprecated, use hasLearnerProperties instead.

Usage

hasProperties(learner, props)

Arguments

learner Deprecated.
props Deprecated.

Description

The built-ins are:

- `imputeConstant(const)` for imputation using a constant value,
- `imputeMedian()` for imputation using the median,
- `imputeMode()` for imputation using the mode,
- `imputeMin(multiplier)` for imputing constant values shifted below the minimum using $\min(x) - \text{multiplier} * \text{diff}(\text{range}(x))$,
- `imputeMax(multiplier)` for imputing constant values shifted above the maximum using $\max(x) + \text{multiplier} * \text{diff}(\text{range}(x))$,
- `imputeNormal(mean, sd)` for imputation using normally distributed random values. Mean and standard deviation will be calculated from the data if not provided.
- `imputeHist(breaks, use.mids)` for imputation using random values with probabilities calculated using table or hist.
- `imputeLearner(learner, preimpute)` for imputations using the response of a classification or regression learner.

Usage

```
imputeConstant(const)

imputeMedian()

imputeMean()

imputeMode()

imputeMin(multiplier = 1)

imputeMax(multiplier = 1)

imputeUniform(min = NA_real_, max = NA_real_)

imputeNormal(mu = NA_real_, sd = NA_real_)

imputeHist(breaks, use.mids = TRUE)

imputeLearner(learner, features = NULL)
```

Arguments

const	[any] Constant valued use for imputation.
multiplier	[numeric(1)] Value that stored minimum or maximum is multiplied with when imputation is done.
min	[numeric(1)] Lower bound for uniform distribution. If NA (default), it will be estimated from the data.
max	[numeric(1)] Upper bound for uniform distribution. If NA (default), it will be estimated from the data.
mu	[numeric(1)] Mean of normal distribution. If missing it will be estimated from the data.
sd	[numeric(1)] Standard deviation of normal distribution. If missing it will be estimated from the data.
breaks	[numeric(1)] Number of breaks to use in hist . If missing, defaults to auto-detection via “Sturges”.
use.mids	[logical(1)] If x is numeric and a histogram is used, impute with bin mids (default) or instead draw uniformly distributed samples within bin range.
learner	[Learner] Supervised learner. Its predictions will be used for imputations. Note that the target column is not available for this operation.
features	[character] Features to use in learner for prediction. Default is NULL which uses all available features except the target column of the original task.

See Also

Other impute: [impute](#), [makeImputeMethod](#), [makeImputeWrapper](#), [reimpute](#)

impute

Impute and re-impute data

Description

Allows imputation of missing feature values through various techniques. Note that you have the possibility to re-impute a data set in the same way as the imputation was performed during training. This especially comes in handy during resampling when one wants to perform the same imputation on the test set as on the training set.

The function `impute` performs the imputation on a data set and returns, alongside with the imputed data set, an “ImputationDesc” object which can contain “learned” coefficients and helpful data. It can then be passed together with a new data set to `reimpute`.

The imputation techniques can be specified for certain features or for feature classes, see function arguments.

You can either provide an arbitrary object, use a built-in imputation method listed under [imputations](#) or create one yourself using `makeImputeMethod`.

Usage

```
impute(data, target = character(0L), classes = list(), cols = list(),
       dummy.classes = character(0L), dummy.cols = character(0L),
       dummy.type = "factor", force.dummies = FALSE, impute.new.levels = TRUE,
       recode.factor.levels = TRUE)
```

Arguments

<code>data</code>	[data.frame] Input data.
<code>target</code>	[character] Name of the column(s) specifying the response. Default is <code>character(0)</code> .
<code>classes</code>	[named list] Named list containing imputation techniques for classes of columns. E.g. <code>list(numeric = imputeMedia</code>
<code>cols</code>	[named list] Named list containing names of imputation methods to impute missing values in the data column referenced by the list element’s name. Overrides imputation set via <code>classes</code> .
<code>dummy.classes</code>	[character] Classes of columns to create dummy columns for. Default is <code>character(0)</code> .
<code>dummy.cols</code>	[character] Column names to create dummy columns (containing binary missing indicator) for. Default is <code>character(0)</code> .
<code>dummy.type</code>	[character(1)] How dummy columns are encoded. Either as 0/1 with type “numeric” or as “factor”. Default is “factor”.
<code>force.dummies</code>	[logical(1)] Force dummy creation even if the respective data column does not contain any NAs. Note that (a) most learners will complain about constant columns created this way but (b) your feature set might be stochastic if you turn this off. Default is FALSE.
<code>impute.new.levels</code>	[logical(1)] If new, unencountered factor level occur during reimputation, should these be handled as NAs and then be imputed the same way? Default is TRUE.

```
recode.factor.levels
```

[logical(1)]
 Recode factor levels after reimputation, so they match the respective element of lvls (in the description object) and therefore match the levels of the feature factor in the training data after imputation?. Default is TRUE.

Details

The description object contains these slots

target [character] See argument.

features [character] Feature names, these are the column names of data, excluding target.

lvls [named list] Mapping of column names of factor features to their levels, including newly created ones during imputation.

impute [named list] Mapping of column names to imputation functions.

dummies [named list] Mapping of column names to imputation functions.

impute.new.levels [logical(1)] See argument.

recode.factor.levels [logical(1)] See argument.

Value

```
data [data.frame]
```

list Imputed data.

```
desc [ImputationDesc]
```

 Description object.

See Also

Other impute: [imputations](#), [makeImputeMethod](#), [makeImputeWrapper](#), [reimpute](#)

Examples

```
df = data.frame(x = c(1, 1, NA), y = factor(c("a", "a", "b")), z = 1:3)
imputed = impute(df, target = character(0), cols = list(x = 99, y = imputeMode()))
print(imputed$data)
reimpute(data.frame(x = NA), imputed$desc)
```

```
iris.task
```

```
Iris classification task.
```

Description

Contains the task (`iris.task`).

References

See [iris](#).

isFailureModel	<i>Is the model a FailureModel?</i>
----------------	-------------------------------------

Description

Such a model is created when one sets the corresponding option in [configureMlr](#).

For complex wrappers this getter returns TRUE if ANY model contained in it failed.

Usage

```
isFailureModel(model)
```

Arguments

model	[WrappedModel] The model.
-------	--

Value

logical(1) .

joinClassLevels	<i>Join some class existing levels to new, larger class levels for classification problems.</i>
-----------------	---

Description

Join some class existing levels to new, larger class levels for classification problems.

Usage

```
joinClassLevels(task, new.levels)
```

Arguments

task	[Task] The task.
new.levels	[list of character] Element names specify the new class levels to create, while the corresponding element character vector specifies the existing class levels which will be joined to the new one.

Value

[Task](#) .

Examples

```
joinClassLevels(iris.task, new.levels = list(foo = c("setosa", "virginica")))
```

learnerArgsToControl	<i>Convert arguments to control structure.</i>
----------------------	--

Description

Find all elements in ... which are not missing and call control on them.

Usage

```
learnerArgsToControl(control, ...)
```

Arguments

control	[function] Function that creates control structure.
...	[any] Arguments for control structure function.

Value

Control structure for learner.

LearnerProperties	<i>Query properties of learners.</i>
-------------------	--------------------------------------

Description

Properties can be accessed with `getLearnerProperties(learner)`, which returns a character vector.

The learner properties are defined as follows:

numerics, factors, ordered Can numeric, factor or ordered factor features be handled?

missings Can missing values in features be handled?

weights Can observations be weighted during fitting?

oneclas, twoclass, multiclass Only for classif: Can one-class, two-class or multi-class classification problems be handled?

class.weights Only for classif: Can class weights be handled?

rcens, lcens, icens Only for surv: Can right, left, or interval censored data be handled?

prob For classif, cluster, multilabel, surv: Can probabilities be predicted?

se Only for regr: Can standard errors be predicted?

Usage

```
getLearnerProperties(learner)

hasLearnerProperties(learner, props)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
props	[character] Vector of properties to query.

Value

`getLearnerProperties` returns a character vector with learner properties. `hasLearnerProperties` returns a logical vector of the same length as `props`.

See Also

Other learner: [getClassWeightParam](#), [getHyperPars](#), [getParamSet](#), [makeLearner](#), [removeHyperPars](#), [setHyperPars](#), [setId](#), [setPredictThreshold](#), [setPredictType](#)

learners	<i>List of supported learning algorithms.</i>
----------	---

Description

All supported learners can be found by [listLearners](#) or as a table in the tutorial appendix: http://mlr-org.github.io/mlr-tutorial/release/html/integrated_learners/.

listFilterMethods	<i>List filter methods.</i>
-------------------	-----------------------------

Description

Returns a subset-able dataframe with filter information.

Usage

```
listFilterMethods(desc = TRUE, tasks = FALSE, features = FALSE)
```

Arguments

desc	[logical(1)] Provide more detailed information about filters.
tasks	[logical(1)] Provide information on supported tasks.
features	[logical(1)] Provide information on supported features.

Value

data.frame .

listLearners	<i>Find matching learning algorithms.</i>
--------------	---

Description

Returns learning algorithms which have specific characteristics, e.g. whether they support missing values, case weights, etc.

Note that the packages of all learners are loaded during the search if you create them. This can be a lot. If you do not create them we only inspect properties of the S3 classes. This will be a lot faster.

Note that for general cost-sensitive learning, mlr currently supports mainly “wrapper” approaches like [CostSensWeightedPairsWrapper](#), which are not listed, as they are not basic R learning algorithms. The same applies for multilabel classification, see [makeMultilabelBinaryRelevanceWrapper](#).

Usage

```
listLearners(obj = NA_character_, properties = character(0L),
  quiet = TRUE, warn.missing.packages = TRUE, check.packages = TRUE,
  create = FALSE)

## Default S3 method:
listLearners(obj, properties = character(0L),
  quiet = TRUE, warn.missing.packages = TRUE, check.packages = TRUE,
  create = FALSE)

## S3 method for class 'character'
listLearners(obj, properties = character(0L),
  quiet = TRUE, warn.missing.packages = TRUE, check.packages = TRUE,
  create = FALSE)

## S3 method for class 'Task'
listLearners(obj, properties = character(0L), quiet = TRUE,
  warn.missing.packages = TRUE, check.packages = TRUE, create = FALSE)
```

Arguments

obj	[character(1) Task] Either a task or the type of the task, in the latter case one of: “classif”, “regr”, “surv”, “costsens”, “cluster”, “multilabel”. Default is NA, matching all types.
properties	[character] Set of required properties to filter for. Default is character(0).
quiet	[logical(1)] Construct learners quietly to check their properties, shows no package startup messages. Turn off if you suspect errors. Default is TRUE.
warn.missing.packages	[logical(1)] If some learner cannot be constructed because its package is missing, should a warning be shown? Default is TRUE.
check.packages	[logical(1)] Check if required packages are installed. Calls find.package(). If create is TRUE, this is done implicitly and the value of this parameter is ignored. Default is TRUE. If set to FALSE, learners that cannot actually be constructed because of missing packages may be returned.
create	[logical(1)] Instantiate objects (or return info table)? Packages are loaded if and only if this option is TRUE. Default is FALSE.

Value

data.frame | list of [Learner](#) . Either descriptive data.frame that allows access to all properties of learners or a list of created learner objects. The latter is named by ids of listed learners.

Examples

```
## Not run:
listLearners("classif", properties = c("multiclass", "prob"))
data = iris
task = makeClassifTask(data = data, target = "Species")
listLearners(task)

## End(Not run)
```

listMeasures

Find matching measures.

Description

Returns the matching measures which have specific characteristics, e.g. whether they supports classification or regression.

Usage

```
listMeasures(obj, properties = character(0L), create = FALSE)

## Default S3 method:
listMeasures(obj, properties = character(0L),
  create = FALSE)

## S3 method for class 'character'
listMeasures(obj, properties = character(0L),
  create = FALSE)

## S3 method for class 'Task'
listMeasures(obj, properties = character(0L), create = FALSE)
```

Arguments

obj	[character(1) Task] Either a task or the type of the task, in the latter case one of: “classif”, “reg”, “surv”, “costsens”, “cluster”, “multilabel”. Default is NA, matching all types.
properties	[character] Set of required properties to filter for. See Measure for some standardized properties. Default is character(0).
create	[logical(1)] Instantiate objects (or return strings)? Default is FALSE.

Value

character | list of [Measure](#) . Class names of matching measures or instantiated objects.

lung.task	<i>NCCTG Lung Cancer survival task.</i>
-----------	---

Description

Contains the task (lung.task).

References

See [lung](#). Incomplete cases have been removed from the task.

makeAggregation	<i>Specify your own aggregation of measures.</i>
-----------------	--

Description

This is an advanced feature of mlr. It gives access to some inner workings so the result might not be compatible with everything!

Usage

```
makeAggregation(id, name = id, fun)
```

Arguments

id	[character(1)] Name of the aggregation method (preferably the same name as the generated function).
name	[character(1)] Long name of the aggregation method. Default is id.
fun	[function(task, perf.test, perf.train, measure, group, pred)] Calculates the aggregated performance. In most cases you will only need the performances perf.test and optionally perf.train on the test and training data sets. task [Task] The task. perf.test [numeric] performance results on the test data sets. perf.train [numeric] performance results on the training data sets. measure [Measure] Performance measure. group [factor] Grouping of resampling iterations. This encodes whether specific iterations 'belong together' (e.g. repeated CV). pred [Prediction] Prediction object.

Value

[Aggregation](#) .

See Also

[aggregations](#), [setAggregation](#)

Examples

```
# computes the interquartile range on all performance values
test.iqr = makeAggregation(id = "test.iqr", name = "Test set interquartile range",
  fun = function (task, perf.test, perf.train, measure, group, pred) IQR(perf.test))
```

makeBaggingWrapper *Fuse learner with the bagging technique.*

Description

Fuses a learner with the bagging method (i.e., similar to what a `randomForest` does). Creates a learner object, which can be used like any other learner object. Models can easily be accessed via [getLearnerModel](#).

Bagging is implemented as follows: For each iteration a random data subset is sampled (with or without replacement) and potentially the number of features is also restricted to a random subset. Note that this is usually handled in a slightly different way in the random forest where features are sampled at each tree split).

Prediction works as follows: For classification we do majority voting to create a discrete label and probabilities are predicted by considering the proportions of all predicted labels. For regression the mean value and the standard deviations across predictions is computed.

Note that the passed base learner must always have `predict.type = 'response'`, while the `BaggingWrapper` can estimate probabilities and standard errors, so it can be set, e.g., to `predict.type = 'prob'`. For this reason, when you call [setPredictType](#), the type is only set for the `BaggingWrapper`, not passed down to the inner learner.

Usage

```
makeBaggingWrapper(learner, bw.iters = 10L, bw.replace = TRUE, bw.size,
  bw.feats = 1)
```

Arguments

<code>learner</code>	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
<code>bw.iters</code>	[integer(1)] Iterations = number of fitted models in bagging. Default is 10.
<code>bw.replace</code>	[logical(1)] Sample bags with replacement (bootstrapping)? Default is TRUE.
<code>bw.size</code>	[numeric(1)] Percentage size of sampled bags. Default is 1 for bootstrapping and 0.632 for subsampling.
<code>bw.feats</code>	[numeric(1)] Percentage size of randomly selected features in bags. Default is 1. At least one feature will always be selected.

Value

[Learner](#) .

See Also

Other wrapper: [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRel](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTEWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

makeClassifTask	Create a classification, regression, survival, cluster, cost-sensitive classification or multilabel task.
-----------------	---

Description

The task encapsulates the data and specifies - through its subclasses - the type of the task. It also contains a description object detailing further aspects of the data.

Useful operators are: [getTaskFormula](#), [getTaskFeatureNames](#), [getTaskData](#), [getTaskTargets](#), and [subsetTask](#).

Object members:

env [environment] Environment where data for the task are stored. Use [getTaskData](#) in order to access it.

weights [numeric] See argument. NULL if not present.

blocking [factor] See argument. NULL if not present.

task.desc [[TaskDesc](#)] Encapsulates further information about the task.

Notes: For multilabel classification we assume that the presence of labels is encoded via logical columns in data. The name of the column specifies the name of the label. `target` is then a char vector that points to these columns.

Usage

```
makeClassifTask(id = deparse(substitute(data)), data, target,
  weights = NULL, blocking = NULL, positive = NA_character_,
  fixup.data = "warn", check.data = TRUE)
```

```
makeClusterTask(id = deparse(substitute(data)), data, weights = NULL,
  blocking = NULL, fixup.data = "warn", check.data = TRUE)
```

```
makeCostSensTask(id = deparse(substitute(data)), data, costs,
  blocking = NULL, fixup.data = "warn", check.data = TRUE)
```

```
makeMultilabelTask(id = deparse(substitute(data)), data, target,
  weights = NULL, blocking = NULL, positive = NA_character_,
  fixup.data = "warn", check.data = TRUE)
```

```
makeRegrTask(id = deparse(substitute(data)), data, target, weights = NULL,
```

```

blocking = NULL, fixup.data = "warn", check.data = TRUE)

makeSurvTask(id = deparse(substitute(data)), data, target,
  censoring = "rcens", weights = NULL, blocking = NULL,
  fixup.data = "warn", check.data = TRUE)

```

Arguments

id	[character(1)] Id string for object. Default is the name of the R variable passed to data.
data	[data.frame] A data frame containing the features and target variable(s).
target	[character(1) character(2) character(n.classes)] Name(s) of the target variable(s). For survival analysis these are the names of the survival time and event columns, so it has length 2. For multilabel classification it contains the names of the logical columns that encode whether a label is present or not and its length corresponds to the number of classes.
weights	[numeric] Optional, non-negative case weight vector to be used during fitting. Cannot be set for cost-sensitive learning. Default is NULL which means no (= equal) weights.
blocking	[factor] An optional factor of the same length as the number of observations. Observations with the same blocking level “belong together”. Specifically, they are either put all in the training or the test set during a resampling iteration. Default is NULL which means no blocking.
positive	[character(1)] Positive class for binary classification (otherwise ignored and set to NA). Default is the first factor level of the target attribute.
fixup.data	[character(1)] Should some basic cleaning up of data be performed? Currently this means removing empty factor levels for the columns. Possible choices are: “no” = Don’t do it. “warn” = Do it but warn about it. “quiet” = Do it but keep silent. Default is “warn”.
check.data	[logical(1)] Should sanity of data be checked initially at task creation? You should have good reasons to turn this off (one might be speed). Default is TRUE.
costs	[data.frame] A numeric matrix or data frame containing the costs of misclassification. We assume the general case of observation specific costs. This means we have n rows, corresponding to the observations, in the same order as data. The columns correspond to classes and their names are the class labels (if unnamed we use y1 to yk as labels). Each entry (i,j) of the matrix specifies the cost of predicting class j for observation i.
censoring	[character(1)] Censoring type. Allowed choices are “rcens” for right censored data (default),

“lcens” for left censored and “icens” for interval censored data using the “interval2” format. See [Surv](#) for details.

Value

[Task](#) .

See Also

Other costsens: [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeCostSensWeightedPairsWrapper](#)

Examples

```
library(mlbench)
data(BostonHousing)
data(Ionosphere)

makeClassifTask(data = iris, target = "Species")
makeRegrTask(data = BostonHousing, target = "medv")
# an example of a classification task with more than those standard arguments:
blocking = factor(c(rep(1, 51), rep(2, 300)))
makeClassifTask(id = "myIonosphere", data = Ionosphere, target = "Class",
  positive = "good", blocking = blocking)
makeClusterTask(data = iris[, -5L])
```

makeCostMeasure	<i>Creates a measure for non-standard misclassification costs.</i>
-----------------	--

Description

Creates a cost measure for non-standard classification error costs.

Usage

```
makeCostMeasure(id = "costs", minimize = TRUE, costs, task,
  combine = mean, best = NULL, worst = NULL, name = id, note = "")
```

Arguments

id	[character(1)] Name of measure. Default is “costs”.
minimize	[logical(1)] Should the measure be minimized? Otherwise you are effectively specifying a benefits matrix. Default is TRUE.
costs	[matrix] Matrix of misclassification costs. Rows and columns have to be named with class labels, order does not matter. Rows indicate true classes, columns predicted classes.

task	[ClassifTask] Classification task. Has to be passed, so validity of matrix names can be checked.
combine	[function] How to combine costs over all cases for a SINGLE test set? Note this is not the same as the aggregate argument in makeMeasure You can set this as well via setAggregation , as for any measure. Default is mean .
best	[numeric(1)] Best obtainable value for measure. Default is -Inf or Inf, depending on minimize.
worst	[numeric(1)] Worst obtainable value for measure. Default is Inf or -Inf, depending on minimize.
name	[character] Name of the measure. Default is id.
note	[character] Description and additional notes for the measure. Default is "".

Value

[Measure](#) .

See Also

Other performance: [estimateRelativeOverfitting](#), [makeCustomResampledMeasure](#), [makeMeasure](#), [measures](#), [performance](#)

makeCostSensClassifWrapper

Wraps a classification learner for use in cost-sensitive learning.

Description

Creates a wrapper, which can be used like any other learner object. The classification model can easily be accessed via [getLearnerModel](#).

This is a very naive learner, where the costs are transformed into classification labels - the label for each case is the name of class with minimal costs. (If ties occur, the label which is better on average w.r.t. costs over all training data is preferred.) Then the classifier is fitted to that data and subsequently used for prediction.

Usage

```
makeCostSensClassifWrapper(learner)
```

Arguments

learner [[Learner](#) | character(1)]
The classification learner. If you pass a string the learner will be created via [makeLearner](#).

Value

[Learner](#) .

See Also

Other costsens: [makeClassifTask](#), [makeCostSensRegrWrapper](#), [makeCostSensWeightedPairsWrapper](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRel](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTERWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

makeCostSensRegrWrapper

Wraps a regression learner for use in cost-sensitive learning.

Description

Creates a wrapper, which can be used like any other learner object. Models can easily be accessed via [getLearnerModel](#).

For each class in the task, an individual regression model is fitted for the costs of that class. During prediction, the class with the lowest predicted costs is selected.

Usage

```
makeCostSensRegrWrapper(learner)
```

Arguments

learner [[Learner](#) | character(1)]
The regression learner. If you pass a string the learner will be created via [makeLearner](#).

Value

[Learner](#) .

See Also

Other costsens: [makeClassifTask](#), [makeCostSensClassifWrapper](#), [makeCostSensWeightedPairsWrapper](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRel](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTERWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

```
makeCostSensWeightedPairsWrapper
```

Wraps a classifier for cost-sensitive learning to produce a weighted pairs model.

Description

Creates a wrapper, which can be used like any other learner object. Models can easily be accessed via [getLearnerModel](#).

For each pair of labels, we fit a binary classifier. For each observation we define the label to be the element of the pair with minimal costs. During fitting, we also weight the observation with the absolute difference in costs. Prediction is performed by simple voting.

This approach is sometimes called cost-sensitive one-vs-one (CS-OVO), because it is obviously very similar to the one-vs-one approach where one reduces a normal multi-class problem to multiple binary ones and aggregates by voting.

Usage

```
makeCostSensWeightedPairsWrapper(learner)
```

Arguments

`learner` [[Learner](#) | character(1)]
The classification learner. If you pass a string the learner will be created via [makeLearner](#).

Value

[Learner](#) .

See Also

Other costsens: [makeClassifTask](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#)

```
makeCustomResampledMeasure
```

Construct your own resampled performance measure.

Description

Construct your own performance measure, used after resampling. Note that individual training / test set performance values will be set to NA, you only calculate an aggregated value. If you can define a function that makes sense for every single training / test set, implement your own [Measure](#).

Usage

```
makeCustomResampledMeasure(measure.id, aggregation.id, minimize = TRUE,
  properties = character(0L), fun, extra.args = list(), best = NULL,
  worst = NULL, measure.name = measure.id,
  aggregation.name = aggregation.id, note = "")
```

Arguments

measure.id	[character(1)] Short name of measure.
aggregation.id	[character(1)] Short name of aggregation.
minimize	[logical(1)] Should the measure be minimized? Default is TRUE.
properties	[character] Set of measure properties. Some standard property names include: classif Is the measure applicable for classification? classif.multi Is the measure applicable for multi-class classification? regr Is the measure applicable for regression? surv Is the measure applicable for survival? costsens Is the measure applicable for cost-sensitive learning? req.pred Is prediction object required in calculation? Usually the case. req.truth Is truth column required in calculation? Usually the case. req.task Is task object required in calculation? Usually not the case req.model Is model object required in calculation? Usually not the case. req.feats Are feature values required in calculation? Usually not the case. req.prob Are predicted probabilities required in calculation? Usually not the case, example would be AUC. Default is character(0).
fun	[function(task, group, pred, extra.args)] Calculates performance value from ResamplePrediction object. For rare cases you can also use the task, the grouping or the extra arguments extra.args. task [Task] The task. group [factor] Grouping of resampling iterations. This encodes whether specific iterations 'belong together' (e.g. repeated CV). pred [Prediction] Prediction object. extra.args [list] See below.
extra.args	[list] List of extra arguments which will always be passed to fun. Default is empty list.
best	[numeric(1)] Best obtainable value for measure. Default is -Inf or Inf, depending on minimize.

worst	[numeric(1)] Worst obtainable value for measure. Default is Inf or -Inf, depending on minimize.
measure.name	[character(1)] Long name of measure. Default is measure.id.
aggregation.name	[character(1)] Long name of the aggregation. Default is aggregation.id.
note	[character] Description and additional notes for the measure. Default is "".

Value

[Measure](#) .

See Also

Other performance: [estimateRelativeOverfitting](#), [makeCostMeasure](#), [makeMeasure](#), [measures](#), [performance](#)

makeDownsampleWrapper *Fuse learner with simple downsampling (subsampling).*

Description

Creates a learner object, which can be used like any other learner object. It will only be trained on a subset of the original data to save computational time.

Usage

```
makeDownsampleWrapper(learner, dw.perc = 1, dw.stratify = FALSE)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
dw.perc	[numeric(1)] See downsample . Default is 1.
dw.stratify	[logical(1)] See downsample . Default is FALSE.

Value

[Learner](#) .

See Also

Other downsample: [downsample](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRel](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTERWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

makeFeatSelWrapper	<i>Fuse learner with feature selection.</i>
--------------------	---

Description

Fuses a base learner with a search strategy to select variables. Creates a learner object, which can be used like any other learner object, but which internally uses [selectFeatures](#). If the train function is called on it, the search strategy and resampling are invoked to select an optimal set of variables. Finally, a model is fitted on the complete training data with these variables and returned. See [selectFeatures](#) for more details.

After training, the optimal features (and other related information) can be retrieved with [getFeatSelResult](#).

Usage

```
makeFeatSelWrapper(learner, resampling, measures, bit.names, bits.to.features,
  control, show.info = getMlrOption("show.info"))
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
resampling	[ResampleInstance ResampleDesc] Resampling strategy for feature selection. If you pass a description, it is instantiated once at the beginning by default, so all points are evaluated on the same training/test sets. If you want to change that behaviour, look at FeatSelControl .
measures	[list of Measure Measure] Performance measures to evaluate. The first measure, aggregated by the first aggregation function is optimized, others are simply evaluated. Default is the default measure for the task, see here getDefaultMeasure .
bit.names	[character] Names of bits encoding the solutions. Also defines the total number of bits in the encoding. Per default these are the feature names of the task.
bits.to.features	[function(x, task)] Function which transforms an integer-0-1 vector into a character vector of selected features. Per default a value of 1 in the ith bit selects the ith feature to be in the candidate solution.

control	[see FeatSelControl] Control object for search method. Also selects the optimization algorithm for feature selection.
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .

Value

[Learner](#) .

See Also

Other featsel: [FeatSelControl](#), [analyzeFeatSelResult](#), [getFeatSelResult](#), [selectFeatures](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTEWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

Examples

```
# nested resampling with feature selection (with a pretty stupid algorithm for selection)
outer = makeResampleDesc("CV", iters = 2L)
inner = makeResampleDesc("Holdout")
ctrl = makeFeatSelControlRandom(maxit = 1)
lrn = makeFeatSelWrapper("classif.ksvm", resampling = inner, control = ctrl)
# we also extract the selected features for all iteration here
r = resample(lrn, iris.task, outer, extract = getFeatSelResult)
```

makeFilter	<i>Create a feature filter.</i>
------------	---------------------------------

Description

Creates and registers custom feature filters. Implemented filters can be listed with [listFilterMethods](#).

Usage

```
makeFilter(name, desc, pkg, supported.tasks, supported.features, fun)
```

Arguments

name	[character(1)] Identifier for the filter.
desc	[character(1)] Short description of the filter.
pkg	[character(1)] Source package where the filter is implemented.

```

supported.tasks
    [character]
    Task types supported.
supported.features
    [character]
    Feature types supported.
fun
    [function(task, nselect, ...)]
    Function which takes a task and returns a named numeric vector of scores, one
    score for each feature of task. Higher scores mean higher importance of the
    feature. At least nselect features must be calculated, the remaining may be set
    to NA or omitted, and thus will not be selected. the original order will be restored
    if necessary.

```

Value

Object of class “Filter”.

makeFilterWrapper	<i>Fuse learner with a feature filter method.</i>
-------------------	---

Description

Fuses a base learner with a filter method. Creates a learner object, which can be used like any other learner object. Internally uses [filterFeatures](#) before every model fit.

After training, the selected features can be retrieved with [getFilteredFeatures](#).

Note that observation weights do not influence the filtering and are simply passed down to the next learner.

Usage

```
makeFilterWrapper(learner, fw.method = "rf.importance", fw.perc = NULL,
  fw.abs = NULL, fw.threshold = NULL, fw.mandatory.feats = NULL, ...)
```

Arguments

```

learner      [Learner | character(1)]
              The learner. If you pass a string the learner will be created via makeLearner.
fw.method    [character(1)]
              Filter method. See listFilterMethods. Default is “rf.importance”.
fw.perc      [numeric(1)]
              If set, select fw.perc*100 top scoring features. Mutually exclusive with argu-
              ments fw.abs and fw.threshold.
fw.abs       [numeric(1)]
              If set, select fw.abs top scoring features. Mutually exclusive with arguments
              fw.perc and fw.threshold.

```

```

fw.threshold    [numeric(1)]
                If set, select features whose score exceeds fw.threshold. Mutually exclusive
                with arguments fw.perc and fw.abs.
fw.mandatory.feats [character]
                Mandatory features which are always included regardless of their scores
...             [any]
                Additional parameters passed down to the filter.

```

Value

[Learner](#) .

See Also

Other filter: [filterFeatures](#), [generateFilterValuesData](#), [getFilterValues](#), [getFilteredFeatures](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTERWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

Examples

```

task = makeClassifTask(data = iris, target = "Species")
lrn = makeLearner("classif.lda")
inner = makeResampleDesc("Holdout")
outer = makeResampleDesc("CV", iters = 2)
lrn = makeFilterWrapper(lrn, fw.perc = 0.5)
mod = train(lrn, task)
print(getFilteredFeatures(mod))
# now nested resampling, where we extract the features that the filter method selected
r = resample(lrn, task, outer, extract = function(model) {
  getFilteredFeatures(model)
})
print(r$extract)

```

makeFixedHoldoutInstance

Generate a fixed holdout instance for resampling.

Description

Generate a fixed holdout instance for resampling.

Usage

```
makeFixedHoldoutInstance(train.inds, test.inds, size)
```

Arguments

<code>train.inds</code>	[integer] Indices for training set.
<code>test.inds</code>	[integer] Indices for test set.
<code>size</code>	[integer(1)] Size of the data set to resample. The function needs to know the largest possible index of the whole data set.

Value

[ResampleInstance](#) .

<code>makeImputeMethod</code>	<i>Create a custom imputation method.</i>
-------------------------------	---

Description

This is a constructor to create your own imputation methods.

Usage

```
makeImputeMethod(learn, impute, args = list())
```

Arguments

<code>learn</code>	[function(data, target, col, ...)] Function to learn and extract information on column <code>col</code> out of data frame <code>data</code> . Argument <code>target</code> specifies the target column of the learning task. The function has to return a named list of values.
<code>impute</code>	[function(data, target, col, ...)] Function to impute missing values in <code>col</code> using information returned by <code>learn</code> on the same column. All list elements of the return values of <code>learn</code> are passed to this function into
<code>args</code>	[list] Named list of arguments to pass to <code>learn</code> via

See Also

Other impute: [imputations](#), [impute](#), [makeImputeWrapper](#), [reimpute](#)

makeImputeWrapper	<i>Fuse learner with an imputation method.</i>
-------------------	--

Description

Fuses a base learner with an imputation method. Creates a learner object, which can be used like any other learner object. Internally uses [impute](#) before training the learner and [reimpute](#) before predicting.

Usage

```
makeImputeWrapper(learner, classes = list(), cols = list(),
  dummy.cols = character(0L), dummy.type = "factor",
  impute.new.levels = TRUE, recode.factor.levels = TRUE)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
classes	[named list] Named list containing imputation techniques for classes of columns. E.g. list(numeric = imputeMedia
cols	[named list] Named list containing names of imputation methods to impute missing values in the data column referenced by the list element's name. Overrides imputation set via classes.
dummy.cols	[character] Column names to create dummy columns (containing binary missing indicator) for. Default is character(0).
dummy.type	[character(1)] How dummy columns are encoded. Either as 0/1 with type "numeric" or as "factor". Default is "factor".
impute.new.levels	[logical(1)] If new, unencountered factor level occur during reimputation, should these be handled as NAs and then be imputed the same way? Default is TRUE.
recode.factor.levels	[logical(1)] Recode factor levels after reimputation, so they match the respective element of lvls (in the description object) and therefore match the levels of the feature factor in the training data after imputation?. Default is TRUE.

Value

[Learner](#) .

See Also

Other impute: [imputations](#), [impute](#), [makeImputeMethod](#), [reimpute](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTEWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

makeLearner

Create learner object.

Description

For a classification learner the `predict.type` can be set to “prob” to predict probabilities and the maximum value selects the label. The threshold used to assign the label can later be changed using the [setThreshold](#) function.

To see all possible properties of a learner, go to: [LearnerProperties](#).

Usage

```
makeLearner(cl, id = cl, predict.type = "response",
  predict.threshold = NULL, fix.factors.prediction = FALSE, ...,
  par.vals = list(), config = list())
```

Arguments

<code>cl</code>	[character(1)] Class of learner. By convention, all classification learners start with “classif.”, all regression learners with “regr.”, all survival learners start with “surv.”, all clustering learners with “cluster.”, and all multilabel classification learners start with “multilabel.”. A list of all integrated learners is available on the learners help page.
<code>id</code>	[character(1)] Id string for object. Used to display object. Default is <code>cl</code> .
<code>predict.type</code>	[character(1)] Classification: “response” (= labels) or “prob” (= probabilities and labels by selecting the ones with maximal probability). Regression: “response” (= mean response) or “se” (= standard errors and mean response). Survival: “response” (= some sort of orderable risk) or “prob” (= time dependent probabilities). Clustering: “response” (= cluster IDS) or “prob” (= fuzzy cluster membership probabilities), Multilabel: “response” (= logical matrix indicating the predicted class labels) or “prob” (= probabilities and corresponding logical matrix indicating class labels). Default is “response”.

<code>predict.threshold</code>	<p>[numeric]</p> <p>Threshold to produce class labels. Has to be a named vector, where names correspond to class labels. Only for binary classification it can be a single numerical threshold for the positive class. See setThreshold for details on how it is applied. Default is NULL which means 0.5 / an equal threshold for each class.</p>
<code>fix.factors.prediction</code>	<p>[logical(1)]</p> <p>In some cases, problems occur in underlying learners for factor features during prediction. If the new features have LESS factor levels than during training (a strict subset), the learner might produce an error like “type of predictors in new data do not match that of the training data”. In this case one can repair this problem by setting this option to TRUE. We will simply add the missing factor levels missing from the test feature (but present in training) to that feature. Default is FALSE.</p>
<code>...</code>	<p>[any]</p> <p>Optional named (hyper)parameters. Alternatively these can be given using the <code>par.vals</code> argument.</p>
<code>par.vals</code>	<p>[list]</p> <p>Optional list of named (hyper)parameters. The arguments in <code>...</code> take precedence over values in this list. We strongly encourage you to use one or the other to pass (hyper)parameters to the learner but not both.</p>
<code>config</code>	<p>[named list]</p> <p>Named list of config option to overwrite global settings set via configureMlr for this specific learner.</p>

Value

[Learner](#) .

See Also

[\[resample\]](#), [\[predict.WrappedModel\]](#)

Other learner: [LearnerProperties](#), [getClassWeightParam](#), [getHyperPars](#), [getParamSet](#), [removeHyperPars](#), [setHyperPars](#), [setId](#), [setPredictThreshold](#), [setPredictType](#)

Examples

```
makeLearner("classif.rpart")
makeLearner("classif.lda", predict.type = "prob")
lrn = makeLearner("classif.lda", method = "t", nu = 10)
print(lrn$par.vals)
```

makeMeasure	<i>Construct performance measure.</i>
-------------	---------------------------------------

Description

A measure object encapsulates a function to evaluate the performance of a prediction. Information about already implemented measures can be obtained here: [measures](#).

A learner is trained on a training set d1, results in a model m and predicts another set d2 (which may be a different one or the training set) resulting in the prediction. The performance measure can now be defined using all of the information of the original task, the fitted model and the prediction.

Object slots:

id [character(1)] See argument.

minimize [logical(1)] See argument.

properties [character] See argument.

fun [function] See argument.

extra.args [list] See argument.

aggr [[Aggregation](#)] See argument.

best [numeric(1)] See argument.

worst [numeric(1)] See argument.

name [character(1)] See argument.

note [character(1)] See argument.

Usage

```
makeMeasure(id, minimize, properties = character(0L), fun,
  extra.args = list(), aggr = test.mean, best = NULL, worst = NULL,
  name = id, note = "")
```

Arguments

id	[character(1)] Name of measure.
minimize	[logical(1)] Should the measure be minimized? Default is TRUE.
properties	[character] Set of measure properties. Some standard property names include: classif Is the measure applicable for classification? classif.multi Is the measure applicable for multi-class classification? multilabel Is the measure applicable for multilabel classification? regr Is the measure applicable for regression? surv Is the measure applicable for survival?

	costsens Is the measure applicable for cost-sensitive learning?
	req.pred Is prediction object required in calculation? Usually the case.
	req.truth Is truth column required in calculation? Usually the case.
	req.task Is task object required in calculation? Usually not the case
	req.model Is model object required in calculation? Usually not the case.
	req.feats Are feature values required in calculation? Usually not the case.
	req.prob Are predicted probabilities required in calculation? Usually not the case, example would be AUC.
	Default is <code>character(0)</code> .
fun	<code>[function(task, model, pred, feats, extra.args)]</code> Calculates the performance value. Usually you will only need the prediction object <code>pred</code> .
	task <code>[Task]</code> The task.
	model <code>[WrappedModel]</code> The fitted model.
	pred <code>[Prediction]</code> Prediction object.
	feats <code>[data.frame]</code> The features.
	extra.args <code>[list]</code> See below.
extra.args	<code>[list]</code> List of extra arguments which will always be passed to <code>fun</code> . Default is empty list.
aggr	<code>[Aggregation]</code> Aggregation function, which is used to aggregate the values measured on test / training sets of the measure to a single value. Default is <code>test.mean</code> .
best	<code>[numeric(1)]</code> Best obtainable value for measure. Default is <code>-Inf</code> or <code>Inf</code> , depending on minimize.
worst	<code>[numeric(1)]</code> Worst obtainable value for measure. Default is <code>Inf</code> or <code>-Inf</code> , depending on minimize.
name	<code>[character]</code> Name of the measure. Default is <code>id</code> .
note	<code>[character]</code> Description and additional notes for the measure. Default is <code>""</code> .

Value

`Measure` .

See Also

Other performance: `estimateRelativeOverfitting`, `makeCostMeasure`, `makeCustomResampledMeasure`, `measures`, `performance`

Examples

```
f = function(task, model, pred, extra.args)
  sum((pred$data$response - pred$data$truth)^2)
makeMeasure(id = "my.sse", minimize = TRUE, properties = c("regr", "response"), fun = f)
```

`makeModelMultiplexer` *Create model multiplexer for model selection to tune over multiple possible models.*

Description

Combines multiple base learners by dispatching on the hyperparameter “selected.learner” to a specific model class. This allows to tune not only the model class (SVM, random forest, etc) but also their hyperparameters in one go. Combine this with [tuneParams](#) and [makeTuneControlIrace](#) for a very powerful approach, see example below.

The parameter set is the union of all (unique) base learners. In order to avoid name clashes all parameter names are prefixed with the base learner id, i.e. “[learner.id].[parameter.name]”.

The `predict.type` of the Multiplexer is inherited from the `predict.type` of the base learners.

Usage

```
makeModelMultiplexer(base.learners)
```

Arguments

`base.learners` [list of [Learner](#)]
 List of Learners with unique IDs.

Value

`ModelMultiplexer` . A [Learner](#) specialized as `ModelMultiplexer`.

Note

Note that logging output during tuning is somewhat shortened to make it more readable. I.e., the artificial prefix before parameter names is suppressed.

See Also

Other multiplexer: [makeModelMultiplexerParamSet](#)

Other tune: [TuneControl](#), [getNestedTuneResultsOptPathDf](#), [getNestedTuneResultsX](#), [getTuneResult](#), [makeModelMultiplexerParamSet](#), [makeTuneWrapper](#), [tuneParams](#), [tuneThreshold](#)

Examples

```
bls = list(
  makeLearner("classif.ksvm"),
  makeLearner("classif.randomForest")
)
lrn = makeModelMultiplexer(bls)
# simple way to construct param set for tuning
```

```

# parameter names are prefixed automatically and the 'requires'
# element is set, too, to make all parameters subordinate to 'selected.learner'
ps = makeModelMultiplexerParamSet(lrn,
  makeNumericParam("sigma", lower = -10, upper = 10, trafo = function(x) 2^x),
  makeIntegerParam("ntree", lower = 1L, upper = 500L)
)
print(ps)
rdesc = makeResampleDesc("CV", iters = 2L)
# to save some time we use random search. but you probably want something like this:
# ctrl = makeTuneControlIrace(maxExperiments = 500L)
ctrl = makeTuneControlRandom(maxit = 10L)
res = tuneParams(lrn, iris.task, rdesc, par.set = ps, control = ctrl)
print(res)
print(head(as.data.frame(res$opt.path)))

# more unique and reliable way to construct the param set
ps = makeModelMultiplexerParamSet(lrn,
  classif.ksvm = makeParamSet(
    makeNumericParam("sigma", lower = -10, upper = 10, trafo = function(x) 2^x)
  ),
  classif.randomForest = makeParamSet(
    makeIntegerParam("ntree", lower = 1L, upper = 500L)
  )
)

# this is how you would construct the param set manually, works too
ps = makeParamSet(
  makeDiscreteParam("selected.learner", values = extractSubList(bls, "id")),
  makeNumericParam("classif.ksvm.sigma", lower = -10, upper = 10, trafo = function(x) 2^x,
    requires = quote(selected.learner == "classif.ksvm")),
  makeIntegerParam("classif.randomForest.ntree", lower = 1L, upper = 500L,
    requires = quote(selected.learner == "classif.randomForst"))
)

# all three ps-objects are exactly the same internally.

```

```
makeModelMultiplexerParamSet
```

Creates a parameter set for model multiplexer tuning.

Description

Handy way to create the param set with less typing.

The following is done automatically:

- The `selected.learner` param is created
- Parameter names are prefixed.
- The `requires` field of each param is set. This makes all parameters subordinate to `selected.learner`

Usage

```
makeModelMultiplexerParamSet(multiplexer, ..., .check = TRUE)
```

Arguments

multiplexer	[ModelMultiplexer] The multiplexer learner.
...	[ParamSet Param] (a) First option: Named param sets. Names must correspond to base learners. You only need to enter the parameters you want to tune without reference to the selected.learner field in any way. (b) Second option. Just the params you would enter in the param sets. Even shorter to create. Only works when it can be uniquely identified to which learner each of your passed parameters belongs.
.check	[logical] Check that for each param in ... one param is found in the base learners. Default is TRUE

Value

[ParamSet](#) .

See Also

Other multiplexer: [makeModelMultiplexer](#)

Other tune: [TuneControl](#), [getNestedTuneResultsOptPathDf](#), [getNestedTuneResultsX](#), [getTuneResult](#), [makeModelMultiplexer](#), [makeTuneWrapper](#), [tuneParams](#), [tuneThreshold](#)

Examples

```
# See makeModelMultiplexer
```

`makeMulticlassWrapper` *Fuse learner with multiclass method.*

Description

Fuses a base learner with a multi-class method. Creates a learner object, which can be used like any other learner object. This way learners which can only handle binary classification will be able to handle multi-class problems, too.

We use a multiclass-to-binary reduction principle, where multiple binary problems are created from the multiclass task. How these binary problems are generated is defined by an error-correcting-output-code (ECOC) code book. This also allows the simple and well-known one-vs-one and one-vs-rest approaches. Decoding is currently done via Hamming decoding, see e.g. here <http://jmlr.org/papers/volume11/escalera10a/escalera10a.pdf>.

Currently, the approach always operates on the discrete predicted labels of the binary base models (instead of their probabilities) and the created wrapper cannot predict posterior probabilities.

Usage

```
makeMulticlassWrapper(learner, mcw.method = "onevsrest")
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
mcw.method	[character(1) function] “onevsone” or “onevsrest”. You can also pass a function, with signature <code>function(task)</code> and which returns a ECOC codematrix with entries +1,-1,0. Columns define new binary problems, rows correspond to classes (rows must be named). 0 means class is not included in binary problem. Default is “onevsrest”.

Value

[Learner](#) .

See Also

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTEWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

```
makeMultilabelBinaryRelevanceWrapper
```

Use binary relevance method to create a multilabel learner.

Description

Every learner which is implemented in mlr and which supports binary classification can be converted to a wrapped binary relevance multilabel learner. The multilabel classification problem is converted into simple binary classifications for each label/target on which the binary learner is applied.

Models can easily be accessed via [getLearnerModel](#).

Note that it does not make sense to set a threshold in the used base learner when you predict probabilities. On the other hand, it can make a lot of sense, to call [setThreshold](#) on the `MultilabelBinaryRelevanceWrapper` for each label individually; Or to tune these thresholds with [tuneThreshold](#); especially when you face very unbalanced class distributions for each binary label.

Usage

```
makeMultilabelBinaryRelevanceWrapper(learner)
```


Arguments

learner [\[Learner\]](#) | character(1)
 The learner. If you pass a string the learner will be created via [makeLearner](#).

Value

[Learner](#) .

References

Tsoumakas, G., & Katakis, I. (2006) *Multi-label classification: An overview*. Dept. of Informatics, Aristotle University of Thessaloniki, Greece.

See Also

Other multilabel: [getMultilabelBinaryPerformances](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTERWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

Examples

```
d = getTaskData(yeast.task)
# drop some labels so example runs faster
d = d[, c(1:3, 15:117)]
task = makeMultilabelTask(data = d, target = c("label1", "label2", "label3"))
lrn = makeMultilabelBinaryRelevanceWrapper("classif.rpart")
lrn = setPredictType(lrn, "prob")
# train, predict and evaluate
mod = train(lrn, yeast.task)
pred = predict(mod, yeast.task)
p = performance(pred)
performance(pred, measure = hamloss)
getMultilabelBinaryPerformances(pred, measures = list(mmce, auc))
# above works also with predictions from resample!
```

makeOverBaggingWrapper

Fuse learner with the bagging technique and oversampling for imbalance correction.

Description

Fuses a classification learner for binary classification with an over-bagging method for imbalance correction when we have strongly unequal class sizes. Creates a learner object, which can be used like any other learner object. Models can easily be accessed via [getLearnerModel](#).

OverBagging is implemented as follows: For each iteration a random data subset is sampled. Class examples are oversampled with replacement with a given rate. Members of the other class are either simply copied into each bag, or bootstrapped with replacement until we have as many majority class examples as in the original training data. Features are currently not changed or sampled.

Prediction works as follows: For classification we do majority voting to create a discrete label and probabilities are predicted by considering the proportions of all predicted labels.

Usage

```
makeOverBaggingWrapper(learner, obw.iters = 10L, obw.rate = 1,
  obw.maxcl = "boot", obw.cl = NULL)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
obw.iters	[integer(1)] Number of fitted models in bagging. Default is 10.
obw.rate	[numeric(1)] Factor to upsample a class in each bag. Must be between 1 and Inf, where 1 means no oversampling and 2 would mean doubling the class size. Default is 1.
obw.maxcl	[character(1)] How should other class (usually larger class) be handled? "all" means every instance of the class gets in each bag, "boot" means the class instances are bootstrapped in each iteration. Default is "boot".
obw.cl	[character(1)] Which class should be over- or undersampled. If NULL, makeOverBaggingWrapper will take the smaller class.

Value

[Learner](#) .

See Also

Other imbalance: [makeUndersampleWrapper](#), [oversample](#), [smote](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTEWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

makePreprocWrapper	<i>Fuse learner with preprocessing.</i>
--------------------	---

Description

Fuses a base learner with a preprocessing method. Creates a learner object, which can be used like any other learner object, but which internally preprocesses the data as requested. If the train or predict function is called on data / a task, the preprocessing is always performed automatically.

Usage

```
makePreprocWrapper(learner, train, predict, par.set = makeParamSet(),
  par.vals = list())
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
train	[function(data, target, args)] Function to preprocess the data before training. target is a string and denotes the target variable in data. args is a list of further arguments and parameters to influence the preprocessing. Must return a list(data, control), where data is the preprocessed data and control stores all information necessary to do the preprocessing before predictions.
predict	[function(data, target, args, control)] Function to preprocess the data before prediction. target is a string and denotes the target variable in data. args are the args that were passed to train. control is the object you returned in train.
par.set	[ParamSet] Parameter set of LearnerParam objects to describe the parameters in args. Default is empty set.
par.vals	[list] Named list of default values for params in args respectively par.set. Default is empty list.

Value

[Learner](#) .

See Also

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makeSMOTERWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

```
makePreprocWrapperCaret
```

Fuse learner with preprocessing.

Description

Fuses a learner with preprocessing methods provided by [preProcess](#). Before training the preprocessing will be performed and the preprocessing model will be stored. Before prediction the preprocessing model will transform the test data according to the trained model.

After being wrapped the learner will support missing values although this will only be the case if `ppc.knnImpute`, `ppc.bagImpute` or `ppc.medianImpute` is set to TRUE.

Usage

```
makePreprocWrapperCaret(learner, ...)
```

Arguments

<code>learner</code>	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
<code>...</code>	[any] See preProcess for parameters not listed above. If you use them you might want to define them in the <code>add.par.set</code> so that they can be tuned.

Value

[Learner](#) .

See Also

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapper](#), [makeSMOTEWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

```
makeResampleDesc
```

Create a description object for a resampling strategy.

Description

A description of a resampling algorithm contains all necessary information to create a [ResampleInstance](#), when given the size of the data set.

Usage

```
makeResampleDesc(method, predict = "test", ..., stratify = FALSE,
  stratify.cols = NULL)
```

Arguments

method	[character(1)] “CV” for cross-validation, “LOO” for leave-one-out, “RepCV” for repeated cross-validation, “Bootstrap” for out-of-bag bootstrap, “Subsample” for subsampling, “Holdout” for holdout.
predict	[character(1)] What to predict during resampling: “train”, “test” or “both” sets. Default is “test”.
...	[any] Further parameters for strategies.
	iters [integer(1)] Number of iterations, for “CV”, “Subsample” and “Bootstrap”.
	split [numeric(1)] Proportion of training cases for “Holdout” and “Subsample” between 0 and 1. Default is 2/3.
	reps [integer(1)] Repeats for “RepCV”. Here <code>iters = folds * reps</code> . Default is 10.
	folds [integer(1)] Folds in the repeated CV for RepCV. Here <code>iters = folds * reps</code> . Default is 10.
stratify	[logical(1)] Should stratification be done for the target variable? For classification tasks, this means that the resampling strategy is applied to all classes individually and the resulting index sets are joined to make sure that the proportion of observations in each training set is as in the original data set. Useful for imbalanced class sizes. For survival tasks stratification is done on the events, resulting in training sets with comparable censoring rates.
stratify.cols	[character] Stratify on specific columns referenced by name. All columns have to be factors. Note that you have to ensure yourself that stratification is possible, i.e. that each strata contains enough observations. This argument and <code>stratify</code> are mutually exclusive.

Details

Some notes on some special strategies:

Repeated cross-validation Use “RepCV”. Then you have to set the aggregation function for your preferred performance measure to “testgroup.mean” via [setAggregation](#).

B632 bootstrap Use “Bootstrap” for bootstrap and set predict to “both”. Then you have to set the aggregation function for your preferred performance measure to “b632” via [setAggregation](#).

B632+ bootstrap Use “Bootstrap” for bootstrap and set predict to “both”. Then you have to set the aggregation function for your preferred performance measure to “b632plus” via [setAggregation](#).

Fixed Holdout set Use [makeFixedHoldoutInstance](#).

Object slots:

id [character(1)] Name of resampling strategy.

iters [integer(1)] Number of iterations. Note that this is always the complete number of generated train/test sets, so for a 10-times repeated 5fold cross-validation it would be 50.

predict [character(1)] See argument.

stratify [logical(1)] See argument.

All parameters passed in ... under the respective argument name See arguments.

Value

[ResampleDesc](#) .

Standard ResampleDesc objects

For common resampling strategies you can save some typing by using the following description objects:

hout holdout a.k.a. test sample estimation (two-thirds training set, one-third testing set)

cv2 2-fold cross-validation

cv3 3-fold cross-validation

cv5 5-fold cross-validation

cv10 10-fold cross-validation

See Also

Other resample: [ResamplePrediction](#), [ResampleResult](#), [getRRPredictions](#), [makeResampleInstance](#), [resample](#)

Examples

```
# Bootstrapping
makeResampleDesc("Bootstrap", iters = 10)
makeResampleDesc("Bootstrap", iters = 10, predict = "both")

# Subsampling
makeResampleDesc("Subsample", iters = 10, split = 3/4)
makeResampleDesc("Subsample", iters = 10)

# Holdout a.k.a. test sample estimation
makeResampleDesc("Holdout")
```

makeResampleInstance *Instantiates a resampling strategy object.*

Description

This class encapsulates training and test sets generated from the data set for a number of iterations. It mainly stores a set of integer vectors indicating the training and test examples for each iteration.

Usage

```
makeResampleInstance(desc, task, size, ...)
```

Arguments

desc	[ResampleDesc character(1)] Resampling description object or name of resampling strategy. In the latter case makeResampleDesc will be called internally on the string.
task	[Task] Data of task to resample from. Prefer to pass this instead of size.
size	[integer] Size of the data set to resample. Can be used instead of task.
...	[any] Passed down to makeResampleDesc in case you passed a string in desc. Otherwise ignored.

Details

Object slots:

desc [[ResampleDesc](#)] See argument.

size [[integer\(1\)](#)] See argument.

train.ind [list of integer] List of training indices for all iterations.

test.ind [list of integer] List of test indices for all iterations.

group [factor] Optional grouping of resampling iterations. This encodes whether specific iterations 'belong together' (e.g. repeated CV), and it can later be used to aggregate performance values accordingly. Default is 'factor()'.

Value

[ResampleInstance](#) .

See Also

Other resample: [ResamplePrediction](#), [ResampleResult](#), [getRRPredictions](#), [makeResampleDesc](#), [resample](#)

Examples

```
rdesc = makeResampleDesc("Bootstrap", iters = 10)
rin = makeResampleInstance(rdesc, task = iris.task)

rdesc = makeResampleDesc("CV", iters = 50)
rin = makeResampleInstance(rdesc, size = nrow(iris))

rin = makeResampleInstance("CV", iters = 10, task = iris.task)
```

makeSMOTEWrapper	<i>Fuse learner with SMOTE oversampling for imbalance correction in binary classification.</i>
------------------	--

Description

Creates a learner object, which can be used like any other learner object. Internally uses [smote](#) before every model fit.

Note that observation weights do not influence the sampling and are simply passed down to the next learner.

Usage

```
makeSMOTEWrapper(learner, sw.rate = 1, sw.nn = 5L, sw.standardize = TRUE,
  sw.alt.logic = FALSE)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
sw.rate	[numeric(1)] Factor to oversample the smaller class. Must be between 1 and Inf, where 1 means no oversampling and 2 would mean doubling the class size. Default is 1.
sw.nn	[integer(1)] Number of nearest neighbors to consider. Default is 5.
sw.standardize	[logical(1)] Standardize input variables before calculating the nearest neighbors for data sets with numeric input variables only. For mixed variables (numeric and factor) the gower distance is used and variables are standardized anyway. Default is TRUE.
sw.alt.logic	[logical(1)] Use an alternative logic for selection of minority class observations. Instead of sampling a minority class element AND one of its nearest neighbors, each minority class element is taken multiple times (depending on rate) for the interpolation and only the corresponding nearest neighbor is sampled. Default is FALSE.

Value

[Learner](#) .

See Also

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

makeStackedLearner	Create a stacked learner object.
--------------------	----------------------------------

Description

A stacked learner uses predictions of several base learners and fits a super learner using these predictions as features in order to predict the outcome. The following stacking methods are available:

`average` Averaging of base learner predictions without weights.

`stack.nocv` Fits the super learner, where in-sample predictions of the base learners are used.

`stack.cv` Fits the super learner, where the base learner predictions are computed by crossvalidated predictions (the resampling strategy can be set via the `resampling` argument).

`hill.climb` Select a subset of base learner predictions by hill climbing algorithm.

`compress` Train a neural network to compress the model from a collection of base learners.

Usage

```
makeStackedLearner(base.learners, super.learner = NULL, predict.type = NULL,
  method = "stack.nocv", use.feat = FALSE, resampling = NULL,
  parset = list())
```

Arguments

<code>base.learners</code>	[(list of) Learner] A list of learners created with <code>makeLearner</code> .
<code>super.learner</code>	[Learner <code>character(1)</code>] The super learner that makes the final prediction based on the base learners. If you pass a string, the super learner will be created via <code>makeLearner</code> . Not used for <code>method = 'average'</code> . Default is <code>NULL</code> .
<code>predict.type</code>	[<code>character(1)</code>] Sets the type of the final prediction for <code>method = 'average'</code> . For other methods, the predict type should be set within <code>super.learner</code> . If the type of the base learner prediction, which is set up within <code>base.learners</code> , is

	<p>"prob" then <code>predict.type = 'prob'</code> will use the average of all base learner predictions and <code>predict.type = 'response'</code> will use the class with highest probability as final prediction.</p> <p>"response" then, for classification tasks with <code>predict.type = 'prob'</code>, the final prediction will be the relative frequency based on the predicted base learner classes and classification tasks with <code>predict.type = 'response'</code> will use majority vote of the base learner predictions to determine the final prediction. For regression tasks, the final prediction will be the average of the base learner predictions.</p>
method	<p>[character(1)]</p> <p>"average" for averaging the predictions of the base learners, "stack.nocv" for building a super learner using the predictions of the base learners, "stack.cv" for building a super learner using crossvalidated predictions of the base learners. "hill.climb" for averaging the predictions of the base learners, with the weights learned from hill climbing algorithm and "compress" for compressing the model to mimic the predictions of a collection of base learners while speeding up the predictions and reducing the size of the model. Default is "stack.nocv",</p>
use.feats	<p>[logical(1)]</p> <p>Whether the original features should also be passed to the super learner. Not used for method = 'average'. Default is FALSE.</p>
resampling	<p>[ResampleDesc]</p> <p>Resampling strategy for method = 'stack.cv'. Currently only CV is allowed for resampling. The default NULL uses 5-fold CV.</p>
parset	<p>the parameters for hill.climb method, including</p> <ul style="list-style-type: none"> replace Whether a base learner can be selected more than once. init Number of best models being included before the selection algorithm. bagprob The proportion of models being considered in one round of selection. bagtime The number of rounds of the bagging selection. metric The result evaluation metric function taking two parameters pred and true, the smaller the score the better. <p>the parameters for compress method, including</p> <ul style="list-style-type: none"> k the size multiplier of the generated data prob the probability to exchange values s the standard deviation of each numerical feature

Examples

```
# Classification
data(iris)
tsk = makeClassifTask(data = iris, target = "Species")
base = c("classif.rpart", "classif.lda", "classif.svm")
lrns = lapply(base, makeLearner)
lrns = lapply(lrns, setPredictType, "prob")
m = makeStackedLearner(base.learners = lrns,
  predict.type = "prob", method = "hill.climb")
tmp = train(m, tsk)
```

```

res = predict(tmp, tsk)

# Regression
data(BostonHousing, package = "mlbench")
tsk = makeRegrTask(data = BostonHousing, target = "medv")
base = c("regr.rpart", "regr.svm")
lrns = lapply(base, makeLearner)
m = makeStackedLearner(base.learners = lrns,
  predict.type = "response", method = "compress")
tmp = train(m, tsk)
res = predict(tmp, tsk)

```

makeTuneWrapper

Fuse learner with tuning.

Description

Fuses a base learner with a search strategy to select its hyperparameters. Creates a learner object, which can be used like any other learner object, but which internally uses [tuneParams](#). If the train function is called on it, the search strategy and resampling are invoked to select an optimal set of hyperparameter values. Finally, a model is fitted on the complete training data with these optimal hyperparameters and returned. See [tuneParams](#) for more details.

After training, the optimal hyperparameters (and other related information) can be retrieved with [getTuneResult](#).

Usage

```

makeTuneWrapper(learner, resampling, measures, par.set, control,
  show.info = getMlrOption("show.info"))

```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
resampling	[ResampleInstance ResampleDesc] Resampling strategy to evaluate points in hyperparameter space. If you pass a description, it is instantiated once at the beginning by default, so all points are evaluated on the same training/test sets. If you want to change that behavior, look at TuneControl .
measures	[list of Measure Measure] Performance measures to evaluate. The first measure, aggregated by the first aggregation function is optimized, others are simply evaluated. Default is the default measure for the task, see here getDefaultMeasure .
par.set	[ParamSet] Collection of parameters and their constraints for optimization. Dependent parameters with a requires field must use quote and not expression to define it.

control	[TuneControl] Control object for search method. Also selects the optimization algorithm for tuning.
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .

Value

[Learner](#) .

See Also

Other tune: [TuneControl](#), [getNestedTuneResultsOptPathDf](#), [getNestedTuneResultsX](#), [getTuneResult](#), [makeModelMultiplexerParamSet](#), [makeModelMultiplexer](#), [tuneParams](#), [tuneThreshold](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTEWrapper](#), [makeUndersampleWrapper](#), [makeWeightedClassesWrapper](#)

Examples

```
task = makeClassifTask(data = iris, target = "Species")
lrn = makeLearner("classif.rpart")
# stupid mini grid
ps = makeParamSet(
  makeDiscreteParam("cp", values = c(0.05, 0.1)),
  makeDiscreteParam("minsplit", values = c(10, 20))
)
ctrl = makeTuneControlGrid()
inner = makeResampleDesc("Holdout")
outer = makeResampleDesc("CV", iters = 2)
lrn = makeTuneWrapper(lrn, resampling = inner, par.set = ps, control = ctrl)
mod = train(lrn, task)
print(getTuneResult(mod))
# nested resampling for evaluation
# we also extract tuned hyper pars in each iteration
r = resample(lrn, task, outer, extract = getTuneResult)
print(r$extract)
getNestedTuneResultsOptPathDf(r)
getNestedTuneResultsX(r)
```

makeUndersampleWrapper

Fuse learner with simple ove/underrsampling for imbalancey correction in binary classification.

Description

Creates a learner object, which can be used like any other learner object. Internally uses [oversample](#) or [undersample](#) before every model fit.

Note that observation weights do not influence the sampling and are simply passed down to the next learner.

Usage

```
makeUndersampleWrapper(learner, usw.rate = 1, usw.cl = NULL)
```

```
makeOversampleWrapper(learner, osw.rate = 1, osw.cl = NULL)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
usw.rate	[numeric(1)] Factor to downsample a class. Must be between 0 and 1, where 1 means no downsampling, 0.5 implies reduction to 50 percent and 0 would imply reduction to 0 observations. Default is 1.
usw.cl	[character(1)] Class that should be undersampled. Default is NULL, which means the larger one.
osw.rate	[numeric(1)] Factor to oversample a class. Must be between 1 and Inf, where 1 means no oversampling and 2 would mean doubling the class size. Default is 1.
osw.cl	[character(1)] Class that should be oversampled. Default is NULL, which means the smaller one.

Value

[Learner](#) .

See Also

Other imbalancecy: [makeOverBaggingWrapper](#), [oversample](#), [smote](#)

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTERWrapper](#), [makeTuneWrapper](#), [makeWeightedClassesWrapper](#)

```
makeWeightedClassesWrapper
```

Wraps a classifier for weighted fitting where each class receives a weight.

Description

Creates a wrapper, which can be used like any other learner object.

Fitting is performed in a weighted fashion where each observation receives a weight, depending on the class it belongs to, see `wcw.weight`. This might help to mitigate problems caused by imbalanced class distributions.

This weighted fitting can be achieved in two ways:

a) The learner already has a parameter for class weighting, so one weight can directly be defined per class. Example: “`classif.ksvm`” and parameter `class.weights`. In this case we don’t really do anything fancy. We convert `wcw.weight` a bit, but basically simply bind its value to the class weighting param. The wrapper in this case simply offers a convenient, consistent fashion for class weighting - and tuning! See example below.

b) The learner does not have a direct parameter to support class weighting, but supports observation weights, so `hasLearnerProperties(learner, 'weights')` is TRUE. This means that an individual, arbitrary weight can be set per observation during training. We set this weight depending on the class internally in the wrapper. Basically we introduce something like a new “`class.weights`” parameter for the learner via observation weights.

Usage

```
makeWeightedClassesWrapper(learner, wcw.param = NULL, wcw.weight = 1)
```

Arguments

<code>learner</code>	[Learner <code>character(1)</code>] The classification learner. If you pass a string the learner will be created via makeLearner .
<code>wcw.param</code>	[<code>character(1)</code>] Name of already existing learner parameter, which allows class weighting. The default (<code>wcw.param = NULL</code>) will use the parameter defined in the learner (<code>class.weights.param</code>). During training, the parameter must accept a named vector of class weights, where length equals the number of classes.
<code>wcw.weight</code>	[<code>numeric</code>] Weight for each class. Must be a vector of the same number of elements as classes are in task, and must also be in the same order as the class levels are in <code>getTaskDescription(task)\$class.levels</code> . For convenience, one must pass a single number in case of binary classification, which is then taken as the weight of the positive class, while the negative class receives a weight of 1. Default is 1.

Value

[Learner](#) .

See Also

Other wrapper: [makeBaggingWrapper](#), [makeCostSensClassifWrapper](#), [makeCostSensRegrWrapper](#), [makeDownsampleWrapper](#), [makeFeatSelWrapper](#), [makeFilterWrapper](#), [makeImputeWrapper](#), [makeMulticlassWrapper](#), [makeMultilabelBinaryRelevanceWrapper](#), [makeOverBaggingWrapper](#), [makePreprocWrapperCaret](#), [makePreprocWrapper](#), [makeSMOTEWrapper](#), [makeTuneWrapper](#), [makeUndersampleWrapper](#)

Examples

```
# using the direct parameter of the SVM (which is already defined in the learner)
lrn = makeWeightedClassesWrapper("classif.ksvm", wcw.weight = 0.01)
res = holdout(lrn, sonar.task)
print(getConfMatrix(res$pred))

# using the observation weights of logreg
lrn = makeWeightedClassesWrapper("classif.logreg", wcw.weight = 0.01)
res = holdout(lrn, sonar.task)
print(getConfMatrix(res$pred))

# tuning the imbalance param and the SVM param in one go
lrn = makeWeightedClassesWrapper("classif.ksvm", wcw.param = "class.weights")
ps = makeParamSet(
  makeNumericParam("wcw.weight", lower = 1, upper = 10),
  makeNumericParam("C", lower = -12, upper = 12, trafo = function(x) 2^x),
  makeNumericParam("sigma", lower = -12, upper = 12, trafo = function(x) 2^x)
)
ctrl = makeTuneControlRandom(maxit = 3L)
rdesc = makeResampleDesc("CV", iters = 2L, stratify = TRUE)
res = tuneParams(lrn, sonar.task, rdesc, par.set = ps, control = ctrl)
print(res)
print(res$opt.path)
```

makeWrappedModel

Induced model of learner.

Description

Result from [train](#).

It internally stores the underlying fitted model, the subset used for training, features used for training, levels of factors in the data set and computation time that was spent for training.

Object members: See arguments.

The constructor `makeWrappedModel` is mainly for internal use.

Usage

```
makeWrappedModel(learner, learner.model, task.desc, subset, features,
  factor.levels, time)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
learner.model	[any] Underlying model.
task.desc	[TaskDesc] Task description object.
subset	[integer] Subset used for training.
features	[character] Features used for training.
factor.levels	[named list of character] Levels of factor variables (features and potentially target) in training data. Named by variable name, non-factors do not occur in the list.
time	[numeric(1)] Computation time for model fit in seconds.

Value

[WrappedModel](#) .

measures	<i>Performance measures.</i>
----------	------------------------------

Description

A performance measure is evaluated after a single train/predict step and returns a single number to assess the quality of the prediction (or maybe only the model, think AIC). The measure itself knows whether it wants to be minimized or maximized and for what tasks it is applicable.

All supported measures can be found by [listMeasures](#) or as a table in the tutorial appendix: <http://mlr-org.github.io/mlr-tutorial/release/html/measures/>.

If you want a measure for a misclassification cost matrix, look at [makeCostMeasure](#). If you want to implement your own measure, look at [makeMeasure](#).

Most measures can directly be accessed via the function named after the scheme measureX (e.g. `measureSSE`).

For clustering measures, we compact the predicted cluster IDs such that they form a continuous series starting with 1. If this is not the case, some of the measures will generate warnings.

Usage

featperc

timetrain

timepredict

timeboth

sse

measureSSE(truth, response)

mse

measureMSE(truth, response)

rmse

measureRMSE(truth, response)

medse

measureMEDSE(truth, response)

sae

measureSAE(truth, response)

mae

measureMAE(truth, response)

medae

measureMEDAE(truth, response)

rsq

measureRSQ(truth, response)

expvar

measureEXPVAR(truth, response)

arsq

mmce

```
measureMMCE(truth, response)

acc

measureACC(truth, response)

ber

multiclass.auc

auc

measureAUC(probabilites, truth, negative, positive)

brier

measureBrier(probabilites, truth, negative, positive)

bac

measureBAC(truth, response, negative, positive)

tp

measureTP(truth, response, positive)

tn

measureTN(truth, response, negative)

fp

measureFP(truth, response, positive)

fn

measureFN(truth, response, negative)

tpr

measureTPR(truth, response, positive)

tnr

measureTNR(truth, response, negative)

fpr
```

measureFPR(truth, response, negative, positive)

fnr

measureFNR(truth, response, negative, positive)

ppv

measurePPV(truth, response, positive)

npv

measureNPV(truth, response, negative)

fdr

measureFDR(truth, response, positive)

mcc

measureMCC(truth, response, negative, positive)

f1

gmean

measureGMEAN(truth, response, negative, positive)

gpr

measureGPR(truth, response, positive)

hamloss

measureHAMLOSS(truth, response)

cindex

meancosts

mcp

db

dunn

G1

G2

silhouette

Arguments

truth	[factor] Vector of the true class.
response	[factor] Vector of the predicted class.
probabilites	[numeric] The probabilites for the positive class.
negative	[character(1)] The name of the negative class.
positive	[character(1)] The name of the positive class.

Format

none

See Also

Other performance: [estimateRelativeOverfitting](#), [makeCostMeasure](#), [makeCustomResampledMeasure](#), [makeMeasure](#), [performance](#)

mergeBenchmarkResultLearner

Merge different learners of BenchmarkResult objects.

Description

Combines the [BenchmarkResult](#) objects that were performed with different learners on the same set of Task(s). This can be helpful if you, e.g. forgot to run one learner on the set of tasks you used.

Usage

```
mergeBenchmarkResultLearner(...)
```

Arguments

... [\[BenchmarkResult\]](#)
BenchmarkResult objects that should be merged.

mergeBenchmarkResultTask

Merge different tasks of BenchmarkResult objects.

Description

Combines the [BenchmarkResult](#) objects that were performed on different tasks with the same set of learner(s). This can be helpful if you, e.g. forgot to run the set of learners on a new task

Usage

```
mergeBenchmarkResultTask(...)
```

Arguments

... [\[BenchmarkResult\]](#)
BenchmarkResult objects that should be merged.

mergeSmallFactorLevels

Merges small levels of factors into new level.

Description

Merges factor levels that occur only infrequently into combined levels with a higher frequency.

Usage

```
mergeSmallFactorLevels(task, cols = NULL, min.perc = 0.01,  
  new.level = ".merged")
```

Arguments

task	[Task] The task.
cols	[character] Which columns to convert. Default is all factor and character columns.
min.perc	[numeric(1)] The smallest levels of a factor are merged until their combined proportion w.r.t. the length of the factor exceeds min.perc. Must be between 0 and 1. Default is 0.01.
new.level	[character(1)] New name of merged level. Default is “.merged”

Value

Task, where merged levels are combined into a new level of name new.level.

See Also

Other eda_and_preprocess: [capLargeValues](#), [createDummyFeatures](#), [dropFeatures](#), [normalizeFeatures](#), [removeConstantFeatures](#), [summarizeColumns](#)

mtcars.task	<i>Motor Trend Car Road Tests clustering task.</i>
-------------	--

Description

Contains the task (mtcars.task).

References

See [mtcars](#).

normalizeFeatures	<i>Normalize features.</i>
-------------------	----------------------------

Description

Normalize features by different methods. Internally [normalize](#) is used for every feature column. Non numerical features will be left untouched and passed to the result. For constant features most methods fail, special behaviour for this case is implemented.

Usage

```
normalizeFeatures(task, method = "standardize", exclude = character(0L),
  range = c(0, 1), on.constant = "quiet")
```

Arguments

- | | |
|--------|--|
| task | [Task]
The task. |
| method | [character(1)]
Normalizing method. Available are:
“center”: Subtract mean.
“scale”: Divide by standard deviation.
“standardize”: Center and scale.
“range”: Scale to a given range. |

exclude	[character] Names of the columns to exclude. The target does not have to be included here. Default is none.
range	[numeric(2)] Range for method “range”. Default is c(0,1).
on.constant	[character(1)] How should constant vectors be treated? Only used, of “method != center”, since this methods does not fail for constant vectors. Possible actions are: “quiet”: Depending on the method, treat them quietly: “scale”: No division by standard deviation is done, input values. will be returned untouched. “standardize”: Only the mean is subtracted, no division is done. “range”: All values are mapped to the mean of the given range. “warn”: Same behaviour as “quiet”, but print a warning message. “stop”: Stop with an error.

Value

Task .

See Also

[normalize](#)

Other eda_and_preprocess: [capLargeValues](#), [createDummyFeatures](#), [dropFeatures](#), [mergeSmallFactorLevels](#), [removeConstantFeatures](#), [summarizeColumns](#)

oversample	<i>Over- or undersample binary classification task to handle class imbalance.</i>
------------	---

Description

Oversampling: For a given class (usually the smaller one) all existing observations are taken and copied and extra observations are added by randomly sampling with replacement from this class.

Undersampling: For a given class (usually the larger one) the number of observations is reduced (downsampled) by randomly sampling without replacement from this class.

Usage

```
oversample(task, rate, cl = NULL)

undersample(task, rate, cl = NULL)
```

Arguments

task	[Task] The task.
rate	[numeric(1)] Factor to upsample or downsample a class. For undersampling: Must be between 0 and 1, where 1 means no downsampling, 0.5 implies reduction to 50 percent and 0 would imply reduction to 0 observations. For oversampling: Must be between 1 and Inf, where 1 means no oversampling and 2 would mean doubling the class size.
cl	[character(1)] Which class should be over- or undersampled. If NULL, oversample will select the smaller and undersample the larger class.

Value[Task](#) .**See Also**Other imbalancy: [makeOverBaggingWrapper](#), [makeUndersampleWrapper](#), [smote](#)

performance

*Measure performance of prediction.***Description**

Measures the quality of a prediction w.r.t. some performance measure.

Usage

performance(pred, measures, task = NULL, model = NULL, feats = NULL)

Arguments

pred	[Prediction] Prediction object.
measures	[Measure list of Measure] Performance measure(s) to evaluate. Default is the default measure for the task, see here getDefaultMeasure .
task	[Task] Learning task, might be requested by performance measure, usually not needed except for clustering.
model	[WrappedModel] Model built on training data, might be requested by performance measure, usually not needed.

feats [data.frame]
 Features of predicted data, usually not needed except for clustering. If the prediction was generated from a task, you can also pass this instead and the features are extracted from it.

Value

named numeric . Performance value(s), named by measure(s).

See Also

Other performance: [estimateRelativeOverfitting](#), [makeCostMeasure](#), [makeCustomResampledMeasure](#), [makeMeasure](#), [measures](#)

Examples

```
training.set = seq(1, nrow(iris), by = 2)
test.set = seq(2, nrow(iris), by = 2)

task = makeClassifTask(data = iris, target = "Species")
lrn = makeLearner("classif.lda")
mod = train(lrn, task, subset = training.set)
pred = predict(mod, newdata = iris[test.set, ])
performance(pred, measures = mmce)

# Compute multiple performance measures at once
ms = list("mmce" = mmce, "acc" = acc, "timetrain" = timetrain)
performance(pred, measures = ms, task, mod)
```

pid.task	<i>PimaIndiansDiabetes classification task.</i>
----------	---

Description

Contains the task (pid.task).

References

See [PimaIndiansDiabetes](#). Note that this is the uncorrected version from mlbench.

plotBMRBoxplots	Create box or violin plots for a <i>BenchmarkResult</i> .
-----------------	---

Description

Plots box or violin plots for a selected measure across all iterations of the resampling strategy, faceted by the `task.id`.

Usage

```
plotBMRBoxplots(bmr, measure = NULL, style = "box", order.lrns = NULL,
  order.tsks = NULL, pretty.names = TRUE)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
measure	[Measure] Performance measure. Default is the first measure used in the benchmark experiment.
style	[character(1)] Type of plot, can be “box” for a boxplot or “violin” for a violin plot. Default is “box”.
order.lrns	[character(n.learners)] Character vector with <code>learner.ids</code> in new order.
order.tsks	[character(n.tasks)] Character vector with <code>task.ids</code> in new order.
pretty.names	[logical(1)] Whether to use the Measure name instead of the id in the plot. Default is TRUE.

Value

ggplot2 plot object.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#)

Other plot: [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Examples

```
# see benchmark
```

```
plotBMRRanksAsBarChart
```

Create a bar chart for ranks in a `BenchmarkResult`.

Description

Plots a bar chart from the ranks of algorithms. Alternatively, tiles can be plotted for every rank-task combination, see `pos` for details. In all plot variants the ranks of the learning algorithms are displayed on the x-axis. Areas are always colored according to the `learner.id`.

Usage

```
plotBMRRanksAsBarChart(bmr, measure = NULL, ties.method = "average",
  aggregation = "default", pos = "stack", order.lrns = NULL,
  order.tsks = NULL)
```

Arguments

<code>bmr</code>	[BenchmarkResult] Benchmark result.
<code>measure</code>	[Measure] Performance measure. Default is the first measure used in the benchmark experiment.
<code>ties.method</code>	[<code>character(1)</code>] See rank for details.
<code>aggregation</code>	[<code>character(1)</code>] “mean” or “default”. See getBMRAggrPerformances for details on “default”.
<code>pos</code>	[<code>character(1)</code>] Optionally set how the bars are positioned in <code>ggplot2</code> . Ranks are plotted on the x-axis. “tile” plots a heat map with task as the y-axis. Allows identification of the performance in a special task. “stack” plots a stacked bar plot. Allows for comparison of learners within and across ranks. “dodge” plots a bar plot with bars next to each other instead of stacked bars.
<code>order.lrns</code>	[<code>character(n.learners)</code>] Character vector with <code>learner.ids</code> in new order.
<code>order.tsks</code>	[<code>character(n.tasks)</code>] Character vector with <code>task.ids</code> in new order.

Value

`ggplot2` plot object.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRSummary](#), [plotCritDifferences](#)

Other plot: [plotBMRBoxplots](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Examples

```
# see benchmark
```

plotBMRSummary	<i>Plot a benchmark summary.</i>
----------------	----------------------------------

Description

Creates a scatter plot, where each line refers to a task. On that line the aggregated scores for all learners are plotted, for that task. Optionally, you can apply a rank transformation or just use one of ggplot2's transformations like [scale_x_log10](#).

Usage

```
plotBMRSummary(bmr, measure = NULL, trafo = "none", order.tsks = NULL,
  pointsize = 4L, jitter = 0.05)
```

Arguments

bmr	[BenchmarkResult] Benchmark result.
measure	[Measure] Performance measure. Default is the first measure used in the benchmark experiment.
trafo	[character(1)] Currently either “none” or “rank”, the latter performing a rank transformation (with average handling of ties) of the scores per task. NB: You can add always add scale_x_log10 to the result to put scores on a log scale. Default is “none”.
order.tsks	[character(n.tasks)] Character vector with task.ids in new order.
pointsize	[numeric(1)] Point size for ggplot2 geom_point for data points. Default is 4.
jitter	[numeric(1)] Small vertical jitter to deal with overplotting in case of equal scores. Default is 0.05.

Value

ggplot2 plot object.

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotCritDifferences](#)

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Examples

```
# see benchmark
```

plotCalibration	<i>Plot calibration data using ggplot2.</i>
-----------------	---

Description

Plots calibration data from [generateCalibrationData](#).

Usage

```
plotCalibration(obj, smooth = FALSE, reference = TRUE, rag = TRUE)
```

Arguments

obj	[CalibrationData] Result of generateCalibrationData .
smooth	[logical(1)] Whether to use a loess smoother. Default is FALSE.
reference	[logical(1)] Whether to plot a reference line showing perfect calibration. Default is TRUE.
rag	[logical(1)] Whether to include a rag plot which shows a rug plot on the top which pertains to positive cases and on the bottom which pertains to negative cases. Default is TRUE.

Value

ggplot2 plot object.

See Also

Other calibration: [generateCalibrationData](#)

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Examples

```
## Not run:
lrns = list(makeLearner("classif.rpart", predict.type = "prob"),
            makeLearner("classif.nnet", predict.type = "prob"))
fit = lapply(lrns, train, task = iris.task)
pred = lapply(fit, predict, task = iris.task)
names(pred) = c("rpart", "nnet")
out = generateCalibrationData(pred, groups = 3)
plotCalibration(out)

fit = lapply(lrns, train, task = sonar.task)
pred = lapply(fit, predict, task = sonar.task)
names(pred) = c("rpart", "lda")
out = generateCalibrationData(pred)
plotCalibration(out)

## End(Not run)
```

plotCritDifferences *Plot critical differences for a selected measure.*

Description

Plots a critical-differences diagram for all classifiers and a selected measure. If a baseline is selected for the Bonferroni-Dunn test, the critical difference interval will be positioned around the baseline. If not, the best performing algorithm will be chosen as baseline. The positioning of some descriptive elements can be moved by modifying the generated data.

Usage

```
plotCritDifferences(obj, baseline = NULL, pretty.names = TRUE)
```

Arguments

obj	[critDifferencesData] Result of generateCritDifferencesData function.
baseline	[character(1)]: [learner.id] Overwrites baseline from generateCritDifferencesData ! Select a [learner.id] as baseline for the critical difference diagram, the critical difference will be positioned around this learner. Defaults to best performing algorithm.

```
pretty.names    [logical(1)]:
                Should learner short names be used instead of learner.id?
```

Value

ggplot2 plot object.

References

Janez Demsar, Statistical Comparisons of Classifiers over Multiple Data Sets, JMLR, 2006

See Also

Other benchmark: [BenchmarkResult](#), [benchmark](#), [convertBMRTToRankMatrix](#), [friedmanPostHocTestBMR](#), [friedmanTestBMR](#), [generateCritDifferencesData](#), [getBMRAggrPerformances](#), [getBMRFeatSelResults](#), [getBMRFilteredFeatures](#), [getBMRLearnerIds](#), [getBMRLearnerShortNames](#), [getBMRLearners](#), [getBMRMeasureIds](#), [getBMRMeasures](#), [getBMRModels](#), [getBMRPerformances](#), [getBMRPredictions](#), [getBMRTaskIds](#), [getBMRTuneResults](#), [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#)

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Examples

```
# see benchmark
```

plotFilterValues	<i>Plot filter values using ggplot2.</i>
------------------	--

Description

Plot filter values using ggplot2.

Usage

```
plotFilterValues(fvalues, sort = "dec", n.show = 20L,
  feat.type.cols = FALSE)
```

Arguments

fvalues	[FilterValues] Filter values.
sort	[character(1)] Sort features like this. “dec” = decreasing, “inc” = increasing, “none” = no sorting. Default is decreasing.

n.show	[integer(1)] Number of features (maximal) to show. Default is 20.
feat.type.cols	[logical(1)] Colors for factor and numeric features. FALSE means no colors. Default is FALSE.

Value

ggplot2 plot object.

See Also

Other filter: [filterFeatures](#), [generateFilterValuesData](#), [getFilterValues](#), [getFilteredFeatures](#), [makeFilterWrapper](#), [plotFilterValuesGGVIS](#)

Other plot: [plotBMBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Examples

```
fv = generateFilterValuesData(iris.task, method = "chi.squared")
plotFilterValues(fv)
```

plotFilterValuesGGVIS *Plot filter values using ggvis.*

Description

Plot filter values using ggvis.

Usage

```
plotFilterValuesGGVIS(fvalues, feat.type.cols = FALSE)
```

Arguments

fvalues	[FilterValues] Filter values.
feat.type.cols	[logical(1)] Colors for factor and numeric features. FALSE means no colors. Default is FALSE.

Value

a ggvis plot object.

See Also

Other filter: [filterFeatures](#), [generateFilterValuesData](#), [getFilterValues](#), [getFilteredFeatures](#), [makeFilterWrapper](#), [plotFilterValues](#)

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPrediction](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Examples

```
## Not run:
fv = generateFilterValuesData(iris.task, method = "chi.squared")
plotFilterValuesGGVIS(fv)

## End(Not run)
```

`plotLearnerPrediction` Visualizes a learning algorithm on a 1D or 2D data set.

Description

Trains the model for 1 or 2 selected features, then displays it via [ggplot](#). Good for teaching or exploring models.

For classification and clustering, only 2D plots are supported. The data points, the classification and potentially through color alpha blending the posterior probabilities are shown.

For regression, 1D and 2D plots are supported. 1D shows the data, the estimated mean and potentially the estimated standard error. 2D does not show estimated standard error, but only the estimated mean via background color.

The plot title displays the model id, its parameters, the training performance and the cross-validation performance.

Usage

```
plotLearnerPrediction(learner, task, features = NULL, measures, cv = 10L,
  ..., gridsize, pointsize = 2, prob.alpha = TRUE, se.band = TRUE,
  err.mark = "train", bg.cols = c("darkblue", "green", "darkred"),
  err.col = "white", err.size = pointsize, greyscale = FALSE)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
task	[Task] The task.
features	[character] Selected features for model. By default the first 2 features are used.

measures	[Measure list of Measure] Performance measure(s) to evaluate. Default is the default measure for the task, see here getDefaultMeasure .
cv	[integer(1)] Do cross-validation and display in plot title? Number of folds. 0 means no CV. Default is 10.
...	[any] Parameters for learner.
gridsize	[integer(1)] Grid resolution per axis for background predictions. Default is 500 for 1D and 100 for 2D.
pointsize	[numeric(1)] Pointsize for ggplot2 geom_point for data points. Default is 2.
prob.alpha	[logical(1)] For classification: Set alpha value of background to probability for predicted class? Allows visualization of “confidence” for prediction. If not, only a constant color is displayed in the background for the predicted label. Default is TRUE.
se.band	[logical(1)] For regression in 1D: Show band for standard error estimation? Default is TRUE.
err.mark	[character(1)]: For classification: Either mark error of the model on the training data (“train”) or during cross-validation (“cv”) or not at all with “none”. Default is “train”.
bg.cols	[character(3)] Background colors for classification and regression. Sorted from low, medium to high. Default is TRUE.
err.col	[character(1)] For classification: Color of misclassified data points. Default is “white”
err.size	[integer(1)] For classification: Size of misclassified data points. Default is pointsize.
greyscale	[logical(1)] Should the plot be greyscale completely? Default is FALSE.

Value

The ggplot2 object.

plotLearningCurve	<i>Plot learning curve data using ggplot2.</i>
-------------------	--

Description

Visualizes data size (percentage used for model) vs. performance measure(s).

Usage

```
plotLearningCurve(obj, facet = "measure", pretty.names = TRUE)
```

Arguments

obj	[LearningCurveData] Result of generateLearningCurveData , with class LearningCurveData.
facet	[character(1)] Selects “measure” or “learner” to be the facetting variable. The variable mapped to facet must have more than one unique value, otherwise it will be ignored. The variable not chosen is mapped to color if it has more than one unique value. The default is “measure”.
pretty.names	[logical(1)] Whether to use the Measure name instead of the id in the plot. Default is TRUE.

Value

ggplot2 plot object.

See Also

Other learning_curve: [generateLearningCurveData](#), [plotLearningCurveGGVIS](#)

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

plotLearningCurveGGVIS

Plot learning curve data using ggvis.

Description

Visualizes data size (percentage used for model) vs. performance measure(s).

Usage

```
plotLearningCurveGGVIS(obj, interaction = "measure", pretty.names = TRUE)
```

Arguments

obj	[LearningCurveData] Result of generateLearningCurveData .
-----	--

interaction	[character(1)] Selects “measure” or “learner” to be used in a Shiny application making the interaction variable selectable via a drop-down menu. This variable must have more than one unique value, otherwise it will be ignored. The variable not chosen is mapped to color if it has more than one unique value. Note that if there are multiple learners and multiple measures interactivity is necessary as ggvis does not currently support facetting or subplots. The default is “measure”.
pretty.names	[logical(1)] Whether to use the Measure name instead of the id in the plot. Default is TRUE.

Value

a ggvis plot object.

See Also

Other learning_curve: [generateLearningCurveData](#), [plotLearningCurve](#)

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

plotPartialPrediction *Plot a partial prediction with ggplot2.*

Description

Plot a partial prediction from [generatePartialPredictionData](#) using ggplot2.

Usage

```
plotPartialPrediction(obj, geom = "line", facet = NULL, p = 1)
```

Arguments

obj	[PartialPredictionData] Generated by generatePartialPredictionData .
geom	[character(1)] The type of geom to use to display the data. Can be “line” or “tile”. For tiling at least two features must be used with <code>interaction = TRUE</code> in the call to generatePartialPredictionData . This may be used in conjunction with the <code>facet</code> argument if three features are specified in the call to generatePartialPredictionData . Default is “line”.
facet	[character(1)] The name of a feature to be used for facetting. This feature must have been an element of the <code>features</code> argument to generatePartialPredictionData and is only applicable when said argument had length greater than 1. If generatePartialPredictionData

is called with the interaction argument FALSE (the default) with argument features of length greater than one, then facet is ignored and each feature is plotted in its own facet. Note that if any of the elements of the features argument of [generatePartialPredictionData](#) are factors, they will be coerced to numerics. Default is NULL.

p [numeric(1)]
If individual = TRUE then sample allows the user to sample without replacement from the output to make the display more readable. Each row is sampled with probability p. Default is 1.

Value

ggplot2 plot object.

See Also

Other partial_prediction: [generatePartialPredictionData](#), [plotPartialPredictionGGVIS](#)

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

plotPartialPredictionGGVIS

Plot a partial prediction using ggvis.

Description

Plot a partial prediction from [generatePartialPredictionData](#) using ggvis.

Usage

```
plotPartialPredictionGGVIS(obj, interact = NULL, p = 1)
```

Arguments

obj [PartialPredictionData]
Generated by [generatePartialPredictionData](#).

interact [character(1)]
The name of a feature to be mapped to an interactive sidebar using Shiny. This feature must have been an element of the features argument to [generatePartialPredictionData](#) and is only applicable when said argument had length greater than 1. If [generatePartialPredictionData](#) is called with the interaction argument FALSE (the default) with argument features of length greater than one, then interact is ignored and the feature displayed is controlled by an interactive side panel. Default is NULL.

p [numeric(1)]
 If `individual = TRUE` then `sample` allows the user to sample without replacement from the output to make the display more readable. Each row is sampled with probability `p`. Default is 1.

Value

a ggvis plot object.

See Also

Other partial_prediction: [generatePartialPredictionData](#), [plotPartialPrediction](#)

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

plotROCCurves	<i>Plots a ROC curve using ggplot2.</i>
---------------	---

Description

Plots a ROC curve from predictions.

Usage

```
plotROCCurves(obj, measures, diagonal = TRUE, pretty.names = TRUE)
```

Arguments

obj [ThreshVsPerfData]
 Result of [generateThreshVsPerfData](#).

measures [list(2) of [Measure](#)]
 Default is the first 2 measures passed to [generateThreshVsPerfData](#).

diagonal [logical(1)]
 Whether to plot a dashed diagonal line. Default is TRUE.

pretty.names [logical(1)]
 Whether to use the [Measure](#) name instead of the id in the plot. Default is TRUE.

Value

a ggvis plot object.

See Also

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Other thresh_vs_perf: [generateThreshVsPerfData](#), [plotThreshVsPerfGGVIS](#), [plotThreshVsPerf](#)

Examples

```
lrn = makeLearner("classif.rpart", predict.type = "prob")
fit = train(lrn, sonar.task)
pred = predict(fit, task = sonar.task)
roc = generateThreshVsPerfData(pred, list(fpr, tpr))
plotROCCurves(roc)

r = bootstrapB632plus(lrn, sonar.task, iters = 3)
roc_r = generateThreshVsPerfData(r, list(fpr, tpr), aggregate = FALSE)
plotROCCurves(roc_r)

r2 = crossval(lrn, sonar.task, iters = 3)
roc_l = generateThreshVsPerfData(list(boot = r, cv = r2), list(fpr, tpr), aggregate = FALSE)
plotROCCurves(roc_l)
```

plotThreshVsPerf	<i>Plot threshold vs. performance(s) for 2-class classification using ggplot2.</i>
------------------	--

Description

Plots threshold vs. performance(s) data that has been generated with [generateThreshVsPerfData](#).

Usage

```
plotThreshVsPerf(obj, facet = "measure", mark.th = NA_real_,
  pretty.names = TRUE)
```

Arguments

obj	[ThreshVsPerfData] Result of generateThreshVsPerfData .
facet	[character(1)] Selects “measure” or “learner” to be the facetting variable. The variable mapped to facet must have more than one unique value, otherwise it will be ignored. The variable not chosen is mapped to color if it has more than one unique value. The default is “measure”.

mark.th	[numeric(1)] Mark given threshold with vertical line? Default is NA which means not to do it.
pretty.names	[logical(1)] Whether to use the Measure name instead of the id in the plot. Default is TRUE.

Value

ggplot2 plot object.

See Also

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#)

Other thresh_vs_perf: [generateThreshVsPerfData](#), [plotROCCurves](#), [plotThreshVsPerfGGVIS](#)

Examples

```
lrn = makeLearner("classif.rpart", predict.type = "prob")
mod = train(lrn, sonar.task)
pred = predict(mod, sonar.task)
pvs = generateThreshVsPerfData(pred, list(acc, setAggregation(acc, train.mean)))
plotThreshVsPerf(pvs)
```

plotThreshVsPerfGGVIS *Plot threshold vs. performance(s) for 2-class classification using ggvis.*

Description

Plots threshold vs. performance(s) data that has been generated with [generateThreshVsPerfData](#).

Usage

```
plotThreshVsPerfGGVIS(obj, interaction = "measure", mark.th = NA_real_,
  pretty.names = TRUE)
```

Arguments

obj	[ThreshVsPerfData] Result of generateThreshVsPerfData .
interaction	[character(1)] Selects “measure” or “learner” to be used in a Shiny application making the interaction variable selectable via a drop-down menu. This variable must have more than one unique value, otherwise it will be ignored. The variable not chosen is mapped to color if it has more than one unique value. Note that if there are multiple learners and multiple measures interactivity is necessary as ggvis does not currently support faceting or subplots. The default is “measure”.

mark.th	[numeric(1)] Mark given threshold with vertical line? Default is NA which means not to do it.
pretty.names	[logical(1)] Whether to use the Measure name instead of the id in the plot. Default is TRUE.

Value

a ggvis plot object.

See Also

Other plot: [plotBMRBoxplots](#), [plotBMRRanksAsBarChart](#), [plotBMRSummary](#), [plotCalibration](#), [plotCritDifferences](#), [plotFilterValuesGGVIS](#), [plotFilterValues](#), [plotLearningCurveGGVIS](#), [plotLearningCurve](#), [plotPartialPredictionGGVIS](#), [plotPartialPrediction](#), [plotROCCurves](#), [plotThreshVsPerf](#)

Other thresh_vs_perf: [generateThreshVsPerfData](#), [plotROCCurves](#), [plotThreshVsPerf](#)

Examples

```
## Not run:
lrn = makeLearner("classif.rpart", predict.type = "prob")
mod = train(lrn, sonar.task)
pred = predict(mod, sonar.task)
pvs = generateThreshVsPerfData(pred, list(tp, fpr))
plotThreshVsPerfGGVIS(pvs)

## End(Not run)
```

plotTuneMultiCritResult

Plots multi-criteria results after tuning using ggplot2.

Description

Visualizes the pareto front and possibly the dominated points.

Usage

```
plotTuneMultiCritResult(res, path = TRUE, col = NULL, shape = NULL,
  pointsize = 2)
```

Arguments

res	[TuneMultiCritResult] Result of tuneParamsMultiCrit .
path	[logical(1)] Visualize all evaluated points (or only the non-dominated pareto front)? For the full path, the size of the points on the front is slightly increased. Default is TRUE.

col	[character(1)] Which column of res\$opt.path should be mapped to ggplot2 color? Default is NULL, which means none.
shape	[character(1)] Which column of res\$opt.path should be mapped to ggplot2 shape? Default is NULL, which means none.
pointsize	[numeric(1)] Point size for ggplot2 geom_point for data points. Default is 2.

Value

ggplot2 plot object.

See Also

Other tune_multicrit: [TuneMultiCritControl](#), [plotTuneMultiCritResultGGVIS](#), [tuneParamsMultiCrit](#)

Examples

```
# see tuneParamsMultiCrit
```

```
plotTuneMultiCritResultGGVIS
```

Plots multi-criteria results after tuning using ggvis.

Description

Visualizes the pareto front and possibly the dominated points.

Usage

```
plotTuneMultiCritResultGGVIS(res, path = TRUE)
```

Arguments

res	[TuneMultiCritResult] Result of tuneParamsMultiCrit .
path	[logical(1)] Visualize all evaluated points (or only the non-dominated pareto front)? Points are colored according to their location. Default is TRUE.

Value

a ggvis plot object.

See Also

Other tune_multicrit: [TuneMultiCritControl](#), [plotTuneMultiCritResult](#), [tuneParamsMultiCrit](#)

Examples

```
# see tuneParamsMultiCrit
```

plotViperCharts	<i>Visualize binary classification predictions via ViperCharts system.</i>
-----------------	--

Description

This includes ROC, lift charts, cost curves, and so on. Please got to <http://viper.ijs.si> for further info.

For resampled learners, the predictions from different iterations are combined into one. That is, for example for cross-validation, the predictions appear on a single line even though they were made by different models. There is currently no facility to separate the predictions for different resampling iterations.

Usage

```
plotViperCharts(obj, chart = "rocc", browse = TRUE, auth.key = NULL,
  task.id = NULL)
```

Arguments

obj	[(list of) Prediction (list of) ResampleResult BenchmarkResult] Single prediction object, list of them, single resample result, list of them, or a benchmark result. In case of a list probably produced by different learners you want to compare, then name the list with the names you want to see in the plots, probably learner shortnames or ids.
chart	[character(1)] First chart to display in focus in browser. All other charts can be displayed by clicking on the browser page menu. Default is "rocc".
browse	[logical(1)] Open ViperCharts plot in web browser? If not you simple get the URL returned. Calls browseURL . Default is TRUE.
auth.key	[character(1)] API key to use for call to Viper charts website. Only required if you want the chart to be private. Default is NULL.
task.id	[character(1)] Selected task in BenchmarkResult to do plots for, ignored otherwise. Default is first task.

Value

character(1) . Invisibly returns the ViperCharts URL.

References

Sluban and Lavrač - ViperCharts: Visual Performance Evaluation Platform, ECML PKDD 2013, pp. 650-653, LNCS 8190, Springer, 2013.

See Also

Other predict: [asROCRPrediction](#), [getPredictionProbabilities](#), [getPredictionResponse](#), [predict.WrappedModel](#), [setPredictThreshold](#), [setPredictType](#)

Other roc: [asROCRPrediction](#)

Examples

```
## Not run:
lrn1 = makeLearner("classif.logreg", predict.type = "prob")
lrn2 = makeLearner("classif.rpart", predict.type = "prob")
b = benchmark(list(lrn1, lrn2), pid.task)
z = plotViperCharts(b, chart = "lift", browse = TRUE)

## End(Not run)
```

`predict.WrappedModel` *Predict new data.*

Description

Predict the target variable of new data using a fitted model. What is stored exactly in the [\[Prediction\]](#) object depends on the `predict.type` setting of the [Learner](#). If `predict.type` was set to “prob” probability thresholding can be done calling the [setThreshold](#) function on the prediction object.

The row names of the input task or newdata are preserved in the output.

Usage

```
## S3 method for class 'WrappedModel'
predict(object, task, newdata, subset, ...)
```

Arguments

<code>object</code>	[WrappedModel] Wrapped model, result of train .
<code>task</code>	[Task] The task. If this is passed, data from this task is predicted.
<code>newdata</code>	[data.frame] New observations which should be predicted. Pass this alternatively instead of <code>task</code> .

subset	[integer integer] An index vector specifying the training cases to be used for fitting. By default the complete data set is used. Logical vectors will be transformed to integer with which .
...	[any] Currently ignored.

Value

[Prediction](#) .

See Also

Other predict: [asROCRPrediction](#), [getPredictionProbabilities](#), [getPredictionResponse](#), [plotViperCharts](#), [setPredictThreshold](#), [setPredictType](#)

Examples

```
# train and predict
train.set = seq(1, 150, 2)
test.set = seq(2, 150, 2)
model = train("classif.lda", iris.task, subset = train.set)
p = predict(model, newdata = iris, subset = test.set)
print(p)
predict(model, task = iris.task, subset = test.set)

# predict now probabilities instead of class labels
lrn = makeLearner("classif.lda", predict.type = "prob")
model = train(lrn, iris.task, subset = train.set)
p = predict(model, task = iris.task, subset = test.set)
print(p)
getPredictionProbabilities(p)
```

Prediction

Prediction object.

Description

Result from [predict.WrappedModel](#). Use `as.data.frame` to access all information in a convenient format. The function [getPredictionProbabilities](#) is useful to access predicted probabilities.

The data member of the object contains always the following columns: `id`, index numbers of predicted cases from the task, `response` either a numeric or a factor, the predicted response values, `truth`, either a numeric or a factor, the true target values. If probabilities were predicted, as many numeric columns as there were classes named `prob.classname`. If standard errors were predicted, a numeric column named `se`.

Object members:

predict.type [character(1)] Type set in [setPredictType](#).
data [data.frame] See details.
threshold [numeric(1)] Threshold set in predict function.
task.desc [[TaskDesc](#)] Task description object.
time [numeric(1)] Time learner needed to generate predictions.

predictLearner	<i>Predict new data with an R learner.</i>
----------------	--

Description

Mainly for internal use. Predict new data with a fitted model. You have to implement this method if you want to add another learner to this package.

Usage

```
predictLearner(.learner, .model, .newdata, ...)
```

Arguments

.learner	[RLearner] Wrapped learner.
.model	[WrappedModel] Model produced by training.
.newdata	[data.frame] New data to predict. Does not include target column.
...	[any] Additional parameters, which need to be passed to the underlying predict function.

Details

Your implementation must adhere to the following: Predictions for the observations in .newdata must be made based on the fitted model (.model\$learner.model). All parameters in ... must be passed to the underlying predict function.

Value

- For classification: Either a factor with class labels for type “response” or, if the learner supports this, a matrix of class probabilities for type “prob”. In the latter case the columns must be named with the class labels.
- For regression: Either a numeric vector for type “response” or, if the learner supports this, a matrix with two columns for type “se”. In the latter case the first column contains the estimated response (mean value) and the second column the estimated standard errors.

- For survival: Either a numeric vector with some sort of orderable risk for type “response” or, if supported, a numeric vector with time dependent probabilities for type “prob”.
- For clustering: Either an integer with cluster IDs for type “response” or, if supported, a matrix of membership probabilities for type “prob”.
- For multilabel: A logical matrix that indicates predicted class labels for type “response” or, if supported, a matrix of class probabilities for type “prob”. The columns must be named with the class labels.

regr.randomForest	<i>regression using randomForest.</i>
-------------------	---------------------------------------

Description

a mlr learner for regrssion tasks using [randomForest](#).

Details

if `predict.type = "se"` the `se.method` (by default “jackknife”) is estimated, using the methods described in Sexton and Laake (2009).

If `se.method = "bootstrap"` the standard error of a prediction is estimated by bootstrapping the random forest, where the number of bootstrap replicates and the number of trees in the ensemble are controlled by `se.boot` and `ntree.for.se` respectively, and then taking the standard deviation of the predictions.

If `se.method = "jackknife"`, the default, the standard error of a prediction is estimated by computing the jackknife-after-bootstrap, the mean-squared difference between the prediction made by only using trees which did not contain said observation and the ensemble prediction.

For both “jackknife” and “bootstrap”, a Monte-Carlo bias correction is applied and, in the case that this results in a negative variance estimate, the values are truncated at 0.

References

[Joseph Sexton] and [Petter Laake],; [Standard errors for bagged and random forest estimators], Computational Statistics and Data Analysis Volume 53, 2009, [801-811].

reimpute	<i>Re-impute a data set</i>
----------	-----------------------------

Description

This function accepts a data frame and a imputation description as returned by [impute](#) to perform the following actions:

1. Restore dropped columns, setting them to NA
2. Add dummy variables for columns as specified in `impute`
3. Optionally check factors for new levels to treat them as NAs
4. Reorder factor levels to ensure identical integer representation as before
5. Impute missing values using previously collected data

Usage

```
reimpute(x, desc)
```

Arguments

x	[data.frame] Object to reimpute. Currently only data frames are supported.
desc	[ImputationDesc] Imputation description as returned by impute .

Value

Imputed x.

See Also

Other impute: [imputations](#), [impute](#), [makeImputeMethod](#), [makeImputeWrapper](#)

```
removeConstantFeatures
```

Remove constant features from a data set.

Description

Constant features can lead to errors in some models and obviously provide no information in the training set that can be learned from. With the argument “perc”, there is a possibility to also remove features for which less than “perc” percent of the observations differ from the mode value.

Usage

```
removeConstantFeatures(task, perc = 0, dont.rm = character(0L),
  na.ignore = FALSE, tol = .Machine$double.eps^0.5,
  show.info = getMlrOption("show.info"))
```

Arguments

task	[Task] The task.
perc	[numeric(1)] The percentage of a feature values in [0, 1) that must differ from the mode value. Default is 0, which means only constant features with exactly one observed level are removed.
dont.rm	[character] Names of the columns which must not be deleted. Default is no columns.
na.ignore	[logical(1)] Should NAs be ignored in the percentage calculation? (Or should they be treated as a single, extra level in the percentage calculation?) Default is FALSE.

tol	[numeric(1)] Numerical tolerance to treat two numbers as equal. Variables stored as double will get rounded accordingly before computing the mode. Default is <code>sqrt(.Machine\$double.eps)</code> .
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .

Value

[Task](#) .

See Also

Other `eda_and_preprocess`: [capLargeValues](#), [createDummyFeatures](#), [dropFeatures](#), [mergeSmallFactorLevels](#), [normalizeFeatures](#), [summarizeColumns](#)

removeHyperPars	<i>Remove hyperparameters settings of a learner.</i>
-----------------	--

Description

Remove settings (previously set through `mlr`) for some parameters. Which means that the default behavior for that param will now be used.

Usage

```
removeHyperPars(learner, ids = character(0L))
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
ids	[character] Parameter names to remove settings for. Default is <code>character(0L)</code> .

Value

[Learner](#) .

See Also

Other learner: [LearnerProperties](#), [getClassWeightParam](#), [getHyperPars](#), [getParamSet](#), [makeLearner](#), [setHyperPars](#), [setId](#), [setPredictThreshold](#), [setPredictType](#)

resample

*Fit models according to a resampling strategy.***Description**

The function `resample` fits a model specified by [Learner](#) on a [Task](#) and calculates predictions and performance [measures](#) for all training and all test sets specified by either a resampling description ([ResampleDesc](#)) or resampling instance ([ResampleInstance](#)).

You are able to return all fitted models (parameter models) or extract specific parts of the models (parameter extract) as returning all of them completely might be memory intensive.

The remaining functions on this page are convenience wrappers for the various existing resampling strategies. Note that if you need to work with precomputed training and test splits (i.e., resampling instances), you have to stick with `resample`.

Usage

```
resample(learner, task, resampling, measures, weights = NULL,
  models = FALSE, extract, keep.pred = TRUE, ...,
  show.info = getMlrOption("show.info"))

crossval(learner, task, iters = 10L, stratify = FALSE, measures,
  models = FALSE, keep.pred = TRUE, ...,
  show.info = getMlrOption("show.info"))

repcv(learner, task, folds = 10L, reps = 10L, stratify = FALSE, measures,
  models = FALSE, keep.pred = TRUE, ...,
  show.info = getMlrOption("show.info"))

holdout(learner, task, split = 2/3, stratify = FALSE, measures,
  models = FALSE, keep.pred = TRUE, ...,
  show.info = getMlrOption("show.info"))

subsample(learner, task, iters = 30, split = 2/3, stratify = FALSE,
  measures, models = FALSE, keep.pred = TRUE, ...,
  show.info = getMlrOption("show.info"))

bootstrap00B(learner, task, iters = 30, stratify = FALSE, measures,
  models = FALSE, keep.pred = TRUE, ...,
  show.info = getMlrOption("show.info"))

bootstrapB632(learner, task, iters = 30, stratify = FALSE, measures,
  models = FALSE, keep.pred = TRUE, ...,
  show.info = getMlrOption("show.info"))

bootstrapB632plus(learner, task, iters = 30, stratify = FALSE, measures,
  models = FALSE, keep.pred = TRUE, ...,
  show.info = getMlrOption("show.info"))
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
task	[Task] The task.
resampling	[ResampleDesc or ResampleInstance] Resampling strategy. If a description is passed, it is instantiated automatically.
measures	[Measure list of Measure] Performance measure(s) to evaluate. Default is the default measure for the task, see here getDefaultMeasure .
weights	[numeric] Optional, non-negative case weight vector to be used during fitting. If given, must be of same length as observations in task and in corresponding order. Overwrites weights specified in the task. By default NULL which means no weights are used unless specified in the task.
models	[logical(1)] Should all fitted models be returned? Default is FALSE.
extract	[function] Function used to extract information from a fitted model during resampling. Is applied to every WrappedModel resulting from calls to train during resampling. Default is to extract nothing.
keep.pred	[logical(1)] Keep the prediction data in the pred slot of the result object. If you do many experiments (on larger data sets) these objects might unnecessarily increase object size / mem usage, if you do not really need them. In this case you can set this argument to FALSE. Default is TRUE.
...	[any] Further hyperparameters passed to learner.
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .
iters	[integer(1)] See ResampleDesc .
stratify	[logical(1)] See ResampleDesc .
folds	[integer(1)] See ResampleDesc .
reps	[integer(1)] See ResampleDesc .
split	[numeric(1)] See ResampleDesc .

Value[ResampleResult](#) .

See Also

Other resample: [ResamplePrediction](#), [ResampleResult](#), [getRRPredictions](#), [makeResampleDesc](#), [makeResampleInstance](#)

Examples

```
task = makeClassifTask(data = iris, target = "Species")
rdesc = makeResampleDesc("CV", iters = 2)
r = resample(makeLearner("classif.qda"), task, rdesc)
print(r$aggr)
print(r$measures.test)
print(r$pred)
```

ResamplePrediction	<i>Prediction from resampling.</i>
--------------------	------------------------------------

Description

Contains predictions from resampling, returned (among other stuff) by function [resample](#). Can basically be used in the same way as [Prediction](#), its super class. The main differences are: (a) The internal data.frame (member data) contains an additional column `iter`, specifying the iteration of the resampling strategy, and and additional columns `set`, specifying whether the prediction was from an observation in the “train” or “test” set. (b) The prediction time is a numeric vector, its length equals the number of iterations.

See Also

Other resample: [ResampleResult](#), [getRRPredictions](#), [makeResampleDesc](#), [makeResampleInstance](#), [resample](#)

ResampleResult	<i>ResampleResult object.</i>
----------------	-------------------------------

Description

A resample result is created by [resample](#) and contains the following object members:

task.id [character(1) :] Name of the Task.

learner.id [character(1) :] Name of the Learner.

measures.test [data.frame :] Gives you access to performance measurements on the individual test sets. Rows correspond to sets in resampling iterations, columns to performance measures.

measures.train [data.frame :] Gives you access to performance measurements on the individual training sets. Rows correspond to sets in resampling iterations, columns to performance measures. Usually not available, only if specifically requested, see general description above.

aggr [numeric :] Named vector of aggregated performance values. Names are coded like this <measure>.<aggregation>.

err.msgs [data.frame :] Number of rows equals resampling iterations and columns are: “iter”, “train”, “predict”. Stores error messages generated during train or predict, if these were caught via `configureMlr`.

pred [`ResamplePrediction` :] Container for all predictions during resampling.

models [list of `WrappedModel` :] List of fitted models or NULL.

extract [list :] List of extracted parts from fitted models or NULL.

runtime [numeric(1) :] Time in seconds it took to execute the resampling.

The print method of this object gives a short overview, including task and learner ids, aggregated measures as well as mean and standard deviation of the measures.

See Also

Other resample: [ResamplePrediction](#), [getRRRPredictions](#), [makeResampleDesc](#), [makeResampleInstance](#), [resample](#)

RLearner

Internal construction / wrapping of learner object.

Description

Wraps an already implemented learning method from R to make it accessible to mlr. Call this method in your constructor. You have to pass an id (name), the required package(s), a description object for all changeable parameters (you do not have to do this for the learner to work, but it is strongly recommended), and use property tags to define features of the learner.

For a general overview on how to integrate a learning algorithm into mlr’s system, please read the section in the online tutorial: http://mlr-org.github.io/mlr-tutorial/release/html/create_learner/index.html

To see all possible properties of a learner, go to: [LearnerProperties](#).

Usage

```
makeRLearner()
```

```
makeRLearnerClassif(cl, package, par.set, par.vals = list(),
  properties = character(0L), name = cl, short.name = cl, note = "",
  class.weights.param = NULL)
```

```
makeRLearnerMultilabel(cl, package, par.set, par.vals = list(),
  properties = character(0L), name = cl, short.name = cl, note = "")
```

```
makeRLearnerRegr(cl, package, par.set, par.vals = list(),
  properties = character(0L), name = cl, short.name = cl, note = "")
```

```
makeRLearnerSurv(cl, package, par.set, par.vals = list(),
  properties = character(0L), name = cl, short.name = cl, note = "")

makeRLearnerCluster(cl, package, par.set, par.vals = list(),
  properties = character(0L), name = cl, short.name = cl, note = "")
```

Arguments

<code>cl</code>	[character(1)] Class of learner. By convention, all classification learners start with “ <code>classif.</code> ”, all regression learners with “ <code>regr.</code> ”, all survival learners start with “ <code>surv.</code> ”, all clustering learners with “ <code>cluster.</code> ”, and all multilabel classification learners start with “ <code>multilabel.</code> ”. A list of all integrated learners is available on the learners help page.
<code>package</code>	[character] Package(s) to load for the implementation of the learner.
<code>par.set</code>	[ParamSet] Parameter set of (hyper)parameters and their constraints. Dependent parameters with a <code>requires</code> field must use quote and not expression to define it.
<code>par.vals</code>	[list] Always set hyperparameters to these values when the object is constructed. Useful when default values are missing in the underlying function. The values can later be overwritten when the user sets hyperparameters. Default is empty list.
<code>properties</code>	[character] Set of learner properties. See above. Default is <code>character(0)</code> .
<code>name</code>	[character(1)] Meaningful name for learner. Default is <code>id</code> .
<code>short.name</code>	[character(1)] Short name for learner. Should only be a few characters so it can be used in plots and tables. Default is <code>id</code> .
<code>note</code>	[character(1)] Additional notes regarding the learner and its integration in <code>mlr</code> . Default is “”.
<code>class.weights.param</code>	[character(1)] Name of the parameter, which can be used for providing class weights.

Value

RLearner . The specific subclass is one of [RLearnerClassif](#), [RLearnerCluster](#), [RLearnerMultilabel](#), [RLearnerRegr](#), [RLearnerSurv](#).

selectFeatures	<i>Feature selection by wrapper approach.</i>
----------------	---

Description

Optimizes the features for a classification or regression problem by choosing a variable selection wrapper approach. Allows for different optimization methods, such as forward search or a genetic algorithm. You can select such an algorithm (and its settings) by passing a corresponding control object. For a complete list of implemented algorithms look at the subclasses of [FeatSelControl](#).

All algorithms operate on a 0-1-bit encoding of candidate solutions. Per default a single bit corresponds to a single feature, but you are able to change this by using the arguments `bit.names` and `bits.to.features`. Thus allowing you to switch on whole groups of features with a single bit.

Usage

```
selectFeatures(learner, task, resampling, measures, bit.names, bits.to.features,
               control, show.info = getMlrOption("show.info"))
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
task	[Task] The task.
resampling	[ResampleInstance ResampleDesc] Resampling strategy for feature selection. If you pass a description, it is instantiated once at the beginning by default, so all points are evaluated on the same training/test sets. If you want to change that behaviour, look at FeatSelControl .
measures	[list of Measure Measure] Performance measures to evaluate. The first measure, aggregated by the first aggregation function is optimized, others are simply evaluated. Default is the default measure for the task, see here getDefaultMeasure .
bit.names	[character] Names of bits encoding the solutions. Also defines the total number of bits in the encoding. Per default these are the feature names of the task.
bits.to.features	[function(x, task)] Function which transforms an integer-0-1 vector into a character vector of selected features. Per default a value of 1 in the <i>i</i> th bit selects the <i>i</i> th feature to be in the candidate solution.
control	[see FeatSelControl] Control object for search method. Also selects the optimization algorithm for feature selection.
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .

Value

[FeatSelResult](#) .

See Also

Other featsel: [FeatSelControl](#), [analyzeFeatSelResult](#), [getFeatSelResult](#), [makeFeatSelWrapper](#)

Examples

```
rdesc = makeResampleDesc("Holdout")
ctrl = makeFeatSelControlSequential(method = "sfs", maxit = NA)
res = selectFeatures("classif.rpart", iris.task, rdesc, control = ctrl)
analyzeFeatSelResult(res)
```

setAggregation	<i>Set aggregation function of measure.</i>
----------------	---

Description

Set how this measure will be aggregated after resampling. To see possible aggregation functions: [aggregations](#).

Usage

```
setAggregation(measure, aggr)
```

Arguments

measure	[Measure] Performance measure.
aggr	[Aggregation] Aggregation function.

Value

[Measure](#) with changed aggregation behaviour.

setHyperPars	<i>Set the hyperparameters of a learner object.</i>
--------------	---

Description

Set the hyperparameters of a learner object.

Usage

```
setHyperPars(learner, ..., par.vals = list())
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
...	[any] Named (hyper)parameters with new setting. Alternatively these can be passed using the <code>par.vals</code> argument.
par.vals	[list] Optional list of named (hyper)parameter settings. The arguments in <code>...</code> take precedence over values in this list.

Value

[Learner](#) .

See Also

Other learner: [LearnerProperties](#), [getClassWeightParam](#), [getHyperPars](#), [getParamSet](#), [makeLearner](#), [removeHyperPars](#), [setId](#), [setPredictThreshold](#), [setPredictType](#)

Examples

```
cl1 = makeLearner("classif.ksvm", sigma = 1)
cl2 = setHyperPars(cl1, sigma = 10, par.vals = list(C = 2))
print(cl1)
# note the now set and altered hyperparameters:
print(cl2)
```

setHyperPars2	<i>Only exported for internal use.</i>
---------------	--

Description

Only exported for internal use.

Usage

```
setHyperPars2(learner, par.vals)
```

Arguments

learner	[Learner] The learner.
par.vals	[list] List of named (hyper)parameter settings.

setId	<i>Set the id of a learner object.</i>
-------	--

Description

Set the id of a learner object.

Usage

```
setId(learner, id)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
id	[character(1)] New id for learner.

Value

[Learner](#) .

See Also

Other learner: [LearnerProperties](#), [getClassWeightParam](#), [getHyperPars](#), [getParamSet](#), [makeLearner](#), [removeHyperPars](#), [setHyperPars](#), [setPredictThreshold](#), [setPredictType](#)

setPredictThreshold	<i>Set the probability threshold the learner should use.</i>
---------------------	--

Description

See `predict.threshold` in [makeLearner](#) and [setThreshold](#).

For complex wrappers only the top-level `predict.type` is currently set.

Usage

```
setPredictThreshold(learner, predict.threshold)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
predict.threshold	[numeric] Threshold to produce class labels. Has to be a named vector, where names correspond to class labels. Only for binary classification it can be a single numerical threshold for the positive class. See setThreshold for details on how it is applied. Default is NULL which means 0.5 / an equal threshold for each class.

Value

[Learner](#) .

See Also

Other learner: [LearnerProperties](#), [getClassWeightParam](#), [getHyperPars](#), [getParamSet](#), [makeLearner](#), [removeHyperPars](#), [setHyperPars](#), [setId](#), [setPredictType](#)

Other predict: [asROCRPrediction](#), [getPredictionProbabilities](#), [getPredictionResponse](#), [plotViperCharts](#), [predict.WrappedModel](#), [setPredictType](#)

setPredictType	<i>Set the type of predictions the learner should return.</i>
----------------	---

Description

Possible prediction types are: Classification: Labels or class probabilities (including labels). Regression: Numeric or response or standard errors (including numeric response). Survival: Linear predictor or survival probability.

For complex wrappers the predict type is usually also passed down the encapsulated learner in a recursive fashion.

Usage

```
setPredictType(learner, predict.type)
```

Arguments

learner [[Learner](#) | character(1)]
The learner. If you pass a string the learner will be created via [makeLearner](#).

predict.type [character(1)]
Classification: “response” or “prob”. Regression: “response” or “se”. Survival: “response” (linear predictor) or “prob”. Clustering: “response” or “prob”. Default is “response”.

Value

[Learner](#) .

See Also

Other learner: [LearnerProperties](#), [getClassWeightParam](#), [getHyperPars](#), [getParamSet](#), [makeLearner](#), [removeHyperPars](#), [setHyperPars](#), [setId](#), [setPredictThreshold](#)

Other predict: [asROCRPrediction](#), [getPredictionProbabilities](#), [getPredictionResponse](#), [plotViperCharts](#), [predict.WrappedModel](#), [setPredictThreshold](#)

setThreshold	<i>Set threshold of prediction object.</i>
--------------	--

Description

Set threshold of prediction object for classification or multilabel classification. Creates corresponding discrete class response for the newly set threshold. For binary classification: The positive class is predicted if the probability value exceeds the threshold. For multiclass: Probabilities are divided by corresponding thresholds and the class with maximum resulting value is selected. The result of both are equivalent if in the multi-threshold case the values are greater than 0 and sum to 1. For multilabel classification: A label is predicted (with entry TRUE) if a probability matrix entry exceeds the threshold of the corresponding label.

Usage

```
setThreshold(pred, threshold)
```

Arguments

pred [[Prediction](#)]
Prediction object.

threshold [numeric]
Threshold to produce class labels. Has to be a named vector, where names correspond to class labels. Only for binary classification it can be a single numerical threshold for the positive class.

Value

[Prediction](#) with changed threshold and corresponding response.

See Also

[predict.WrappedModel](#)

Examples

```
# create task and train learner (LDA)
task = makeClassifTask(data = iris, target = "Species")
lrn = makeLearner("classif.lda", predict.type = "prob")
mod = train(lrn, task)

# predict probabilities and compute performance
pred = predict(mod, newdata = iris)
performance(pred, measures = mmce)
head(as.data.frame(pred))

# adjust threshold and predict probabilities again
threshold = c(setosa = 0.4, versicolor = 0.3, virginica = 0.3)
pred = setThreshold(pred, threshold = threshold)
performance(pred, measures = mmce)
head(as.data.frame(pred))
```

smote

Synthetic Minority Oversampling Technique to handle class imbalance in binary classification.

Description

In each iteration, samples one minority class element x_1 , then one of x_1 's nearest neighbors: x_2 . Both points are now interpolated / convex-combined, resulting in a new virtual data point x_3 for the minority class.

The method handles factor features, too. The gower distance is used for nearest neighbor calculation, see [daisy](#). For interpolation, the new factor level for x_3 is sampled from the two given levels of x_1 and x_2 per feature.

Usage

```
smote(task, rate, nn = 5L, standardize = TRUE, alt.logic = FALSE)
```

Arguments

task [\[Task\]](#)
The task.

rate	[numeric(1)] Factor to upsample the smaller class. Must be between 1 and Inf, where 1 means no oversampling and 2 would mean doubling the class size.
nn	[integer(1)] Number of nearest neighbors to consider. Default is 5.
standardize	[integer(1)] Standardize input variables before calculating the nearest neighbors for data sets with numeric input variables only. For mixed variables (numeric and factor) the gower distance is used and variables are standardized anyway. Default is TRUE.
alt.logic	[integer(1)] Use an alternative logic for selection of minority class observations. Instead of sampling a minority class element AND one of its nearest neighbors, each minority class element is taken multiple times (depending on rate) for the interpolation and only the corresponding nearest neighbor is sampled. Default is FALSE.

Value

[Task](#) .

References

Chawla, N., Bowyer, K., Hall, L., & Kegelmeyer, P. (2000) *SMOTE: Synthetic Minority Over-sampling TEchnique*. In International Conference of Knowledge Based Computer Systems, pp. 46-57. National Center for Software Technology, Mumbai, India, Allied Press.

See Also

Other imbalancecy: [makeOverBaggingWrapper](#), [makeUndersampleWrapper](#), [oversample](#)

sonar.task

Sonar classification task.

Description

Contains the task (sonar.task).

References

See [Sonar](#).

subsetTask	<i>Subset data in task.</i>
------------	-----------------------------

Description

Subset data in task.

Usage

```
subsetTask(task, subset, features)
```

Arguments

task	[Task] The task.
subset	[integer logical(n)] Selected cases. Default is all cases.
features	[character] Selected inputs. Note that target feature is always included in the resulting task, you should not pass it here. Default is all features.

Value

[Task](#) . Task with subsetted data.

See Also

Other task: [getTaskClassLevels](#), [getTaskCosts](#), [getTaskData](#), [getTaskDescription](#), [getTaskFeatureNames](#), [getTaskFormula](#), [getTaskId](#), [getTaskNFeats](#), [getTaskSize](#), [getTaskTargetNames](#), [getTaskTargets](#), [getTaskType](#)

Examples

```
task = makeClassifTask(data = iris, target = "Species")
subsetTask(task, subset = 1:100)
```

summarizeColumns	<i>Summarize columns of data.frame or task.</i>
------------------	---

Description

Summarizes a data.frame, somewhat differently than the normal [summary](#) function of R. The function is mainly useful as a basic EDA tool on data.frames before they are converted to tasks, but can be used on tasks as well.

Columns can be of type numeric, integer, logical, factor, or character. Characters and logicals will be treated as factors.

Usage

```
summarizeColumns(obj)
```

Arguments

obj [data.frame | [Task](#)]
Input data.

Value

data.frame . With columns:

name	Name of column.
type	Data type of column.
na	Number of NAs in column.
disp	Measure of dispersion, for numerics and integers sd is used, for categorical columns the qualitative variation.
mean	Mean value of column, NA for categorical columns.
median	Median value of column, NA for categorical columns.
mad	MAD of column, NA for categorical columns.
min	Minimal value of column, for categorical columns the size of the smallest category.
max	Maximal value of column, for categorical columns the size of the largest category.
nlevs	For categorical columns, the number of factor levels, NA else.

See Also

Other eda_and_preprocess: [capLargeValues](#), [createDummyFeatures](#), [dropFeatures](#), [mergeSmallFactorLevels](#), [normalizeFeatures](#), [removeConstantFeatures](#)

Examples

```
summarizeColumns(iris)
```

summarizeLevels	<i>Summarizes factors of a data.frame by tabling them.</i>
-----------------	--

Description

Characters and logicals will be treated as factors.

Usage

```
summarizeLevels(obj, cols = NULL)
```


Arguments

obj	[data.frame Task] Input data.
cols	[character] Restrict result to columns in cols. Default is all factor, character and logical columns of obj.

Value

`list` . Named list of tables.

TaskDesc	<i>Description object for task.</i>
----------	-------------------------------------

Description

Description object for task, encapsulates basic properties of the task without having to store the complete data set.

Details

Object members:

id [character(1)] Id string of task.

type [character(1)] Type of task, “classif” for classification, “regr” for regression, “surv” for survival and “cluster” for cluster analysis, “costsens” for cost-sensitive classification, and “multilabel” for multilabel classification.

target [character(0) | character(1) | character(2) | character(n.classes)] Name(s) of the target variable(s). For “surv” these are the names of the survival time and event columns, so it has length 2. For “costsens” it has length 0, as there is no target column, but a cost matrix instead. For “multilabel” these are the names of logical columns that indicate whether a class label is present and the number of target variables corresponds to the number of classes.

size [integer(1)] Number of cases in data set.

n.feats [integer(2)] Number of features, named vector with entries: “numerics”, “factors”, “ordered”.

has.missings [logical(1)] Are missing values present?

has.weights [logical(1)] Are weights specified for each observation?

has.blocking [logical(1)] Is a blocking factor for cases available in the task?

class.levels [character] All possible classes. Only present for “classif”, “costsens”, and “multilabel”.

positive [character(1)] Positive class label for binary classification. Only present for “classif”, NA for multiclass.

negative [character(1)] Negative class label for binary classification. Only present for “classif”, NA for multiclass.

censoring [character(1)] Censoring type for survival analysis. Only present for “surv”, one of “rcens” for right censored data, “lcens” for left censored data, and “icens” for interval censored data.

train	<i>Train a learning algorithm.</i>
-------	------------------------------------

Description

Given a [Task](#), creates a model for the learning machine which can be used for predictions on new data.

Usage

```
train(learner, task, subset, weights = NULL)
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
task	[Task] The task.
subset	[integer logical] An index vector specifying the training cases to be used for fitting. By default the complete data set is used. Logical vectors will be transformed to integer with which .
weights	[numeric] Optional, non-negative case weight vector to be used during fitting. If given, must be of same length as subset and in corresponding order. By default NULL which means no weights are used unless specified in the task (Task). Weights from the task will be overwritten.

Value

[WrappedModel](#) .

See Also

[predict.WrappedModel](#)

Examples

```
training.set = sample(1:nrow(iris), nrow(iris) / 2)

## use linear discriminant analysis to classify iris data
task = makeClassifTask(data = iris, target = "Species")
learner = makeLearner("classif.lda", method = "mle")
```

```
mod = train(learner, task, subset = training.set)
print(mod)

## use random forest to classify iris data
task = makeClassifTask(data = iris, target = "Species")
learner = makeLearner("classif.rpart", minsplit = 7, predict.type = "prob")
mod = train(learner, task, subset = training.set)
print(mod)
```

trainLearner*Train an R learner.*

Description

Mainly for internal use. Trains a wrapped learner on a given training set. You have to implement this method if you want to add another learner to this package.

Usage

```
trainLearner(.learner, .task, .subset, .weights = NULL, ...)
```

Arguments

<code>.learner</code>	[RLearner] Wrapped learner.
<code>.task</code>	[Task] Task to train learner on.
<code>.subset</code>	[integer] Subset of cases for training set, index the task with this. You probably want to use getTaskData for this purpose.
<code>.weights</code>	[numeric] Weights for each observation.
<code>...</code>	[any] Additional (hyper)parameters, which need to be passed to the underlying train function.

Details

Your implementation must adhere to the following: The model must be fitted on the subset of `.task` given by `.subset`. All parameters in `...` must be passed to the underlying training function.

Value

any . Model of the underlying learner.

TuneControl

Create control structures for tuning.

Description

The following tuners are available:

makeTuneControlGrid Grid search. All kinds of parameter types can be handled. You can either use their correct param type and resolution, or discretize them yourself by always using [makeDiscreteParam](#) in the `par.set` passed to [tuneParams](#).

makeTuneControlRandom Random search. All kinds of parameter types can be handled.

makeTuneControlDesign Completely pre-specify a `data.frame` of design points to be evaluated during tuning. All kinds of parameter types can be handled.

makeTuneControlCMAES CMA Evolution Strategy with method [cma_es](#). Can handle `numeric(vector)` and `integer(vector)` hyperparameters, but no dependencies. For integers the internally proposed numeric values are automatically rounded. The sigma variance parameter is initialized to 1/4 of the span of box-constraints per parameter dimension.

makeTuneControlGenSA Generalized simulated annealing with method [GenSA](#). Can handle `numeric(vector)` and `integer(vector)` hyperparameters, but no dependencies. For integers the internally proposed numeric values are automatically rounded.

makeTuneControlIrace Tuning with iterated F-Racing with method [irace](#). All kinds of parameter types can be handled. We return the best of the final elite candidates found by irace in the last race. Its estimated performance is the mean of all evaluations ever done for that candidate.

Some notes on irace: For resampling you have to pass a [ResampleDesc](#), not a [ResampleInstance](#). The resampling strategy is randomly instantiated `n.instances` times and these are the instances in the sense of irace (instances element of `tunerConfig` in [irace](#)). Also note that irace will always store its tuning results in a file on disk, see the package documentation for details on this and how to change the file path.

Usage

```
makeTuneControlCMAES(same.resampling.instance = TRUE, impute.val = NULL,
  start = NULL, tune.threshold = FALSE, tune.threshold.args = list(),
  log.fun = NULL, final.dw.perc = NULL, budget = NULL, ...)
```

```
makeTuneControlDesign(same.resampling.instance = TRUE, impute.val = NULL,
  design = NULL, tune.threshold = FALSE, tune.threshold.args = list(),
  log.fun = NULL, budget = NULL)
```

```
makeTuneControlGenSA(same.resampling.instance = TRUE, impute.val = NULL,
  start = NULL, tune.threshold = FALSE, tune.threshold.args = list(),
  log.fun = NULL, final.dw.perc = NULL, budget = NULL, ...)
```

```
makeTuneControlGrid(same.resampling.instance = TRUE, impute.val = NULL,
  resolution = 10L, tune.threshold = FALSE, tune.threshold.args = list(),
```

```

log.fun = NULL, final.dw.perc = NULL, budget = NULL)

makeTuneControlIrace(impute.val = NULL, n.instances = 100L,
  show.irace.output = FALSE, tune.threshold = FALSE,
  tune.threshold.args = list(), log.fun = NULL, final.dw.perc = NULL,
  budget = NULL, ...)

makeTuneControlRandom(same.resampling.instance = TRUE, maxit = 100L,
  tune.threshold = FALSE, tune.threshold.args = list(), log.fun = NULL,
  final.dw.perc = NULL, budget = NULL)

```

Arguments

same.resampling.instance	<p>[logical(1)]</p> <p>Should the same resampling instance be used for all evaluations to reduce variance? Default is TRUE.</p>
impute.val	<p>[numeric]</p> <p>If something goes wrong during optimization (e.g. the learner crashes), this value is fed back to the tuner, so the tuning algorithm does not abort. It is not stored in the optimization path, an NA and a corresponding error message are logged instead. Note that this value is later multiplied by -1 for maximization measures internally, so you need to enter a larger positive value for maximization here as well. Default is the worst obtainable value of the performance measure you optimize for when you aggregate by mean value, or Inf instead. For multi-criteria optimization pass a vector of imputation values, one for each of your measures, in the same order as your measures.</p>
start	<p>[list]</p> <p>Named list of initial parameter values.</p>
tune.threshold	<p>[logical(1)]</p> <p>Should the threshold be tuned for the measure at hand, after each hyperparameter evaluation, via tuneThreshold? Only works for classification if the predict type is “prob”. Default is FALSE.</p>
tune.threshold.args	<p>[list]</p> <p>Further arguments for threshold tuning that are passed down to tuneThreshold. Default is none.</p>
log.fun	<p>[function NULL]</p> <p>Function used for logging. If set to NULL, the internal default will be used. Otherwise a function with arguments learner, resampling, measures, par.set, control, opt.path, dob, x, y, remove.nas, and stage is expected. The default displays the performance measures, the time needed for evaluating, the currently used memory and the max memory ever used before (the latter two both taken from gc). See the implementation for details.</p>
final.dw.perc	<p>[boolean]</p> <p>If a Learner wrapped by a makeDownsampleWrapper is used, you can define the value of dw.perc which is used to train the Learner with the final parameter setting found by the tuning. Default is NULL which will not change anything.</p>

budget	[integer(1)] Maximum budget for tuning. This value restricts the number of function evaluations. In case of <code>makeTuneControlGrid</code> this number must be identical to the size of the grid. For <code>makeTuneControlRandom</code> the budget equals the number of iterations (<code>maxit</code>) performed by the random search algorithm. Within the <code>cma_es</code> the budget corresponds to the product of the number of generations (<code>maxit</code>) and the number of offsprings per generation (<code>lambda</code>). <code>GenSA</code> defines the budget via the argument <code>max.call</code> . However, one should note that this algorithm does not stop its local search before its end. This behaviour might lead to an extension of the defined budget and will result in a warning. In <code>irace</code> , budget is passed to <code>maxExperiments</code> .
...	[any] Further control parameters passed to the control arguments of <code>cma_es</code> or <code>GenSA</code> , as well as towards the <code>tunerConfig</code> argument of <code>irace</code> .
design	[data.frame] data.frame containing the different parameter settings to be evaluated. The columns have to be named according to the <code>ParamSet</code> which will be used in <code>tune()</code> . Proper designs can be created with <code>generateDesign</code> for instance.
resolution	[integer] Resolution of the grid for each numeric/integer parameter in <code>par.set</code> . For vector parameters, it is the resolution per dimension. Either pass one resolution for all parameters, or a named vector. See <code>generateGridDesign</code> . Default is 10.
n.instances	[integer(1)] Number of random resampling instances for <code>irace</code> , see details. Default is 100.
show.irace.output	[logical(1)] Show console output of <code>irace</code> while tuning? Default is FALSE.
maxit	[integer(1) NULL] Number of iterations for random search. Default is 100.

Value

`TuneControl` . The specific subclass is one of `TuneControlGrid`, `TuneControlRandom`, `TuneControlCMAES`, `TuneControlGenSA`, `TuneControlIrace`.

See Also

Other tune: `getNestedTuneResultsOptPathDf`, `getNestedTuneResultsX`, `getTuneResult`, `makeModelMultiplexerParam`, `makeModelMultiplexer`, `makeTuneWrapper`, `tuneParams`, `tuneThreshold`

`TuneMultiCritControl` *Create control structures for multi-criteria tuning.*

Description

The following tuners are available:

makeTuneMultiCritControlGrid Grid search. All kinds of parameter types can be handled. You can either use their correct param type and resolution, or discretize them yourself by always using [makeDiscreteParam](#) in the `par.set` passed to [tuneParams](#).

makeTuneMultiCritControlRandom Random search. All kinds of parameter types can be handled.

makeTuneMultiCritControlNSGA2 Evolutionary method [nsga2](#). Can handle numeric(vector) and integer(vector) hyperparameters, but no dependencies. For integers the internally proposed numeric values are automatically rounded.

Usage

```
makeTuneMultiCritControlGrid(same.resampling.instance = TRUE,
  resolution = 10L, log.fun = NULL, final.dw.perc = NULL, budget = NULL)
```

```
makeTuneMultiCritControlNSGA2(same.resampling.instance = TRUE,
  impute.val = NULL, log.fun = NULL, final.dw.perc = NULL,
  budget = NULL, ...)
```

```
makeTuneMultiCritControlRandom(same.resampling.instance = TRUE,
  maxit = 100L, log.fun = NULL, final.dw.perc = NULL, budget = NULL)
```

Arguments

<code>same.resampling.instance</code>	[logical(1)] Should the same resampling instance be used for all evaluations to reduce variance? Default is TRUE.
<code>resolution</code>	[integer] Resolution of the grid for each numeric/integer parameter in <code>par.set</code> . For vector parameters, it is the resolution per dimension. Either pass one resolution for all parameters, or a named vector. See generateGridDesign . Default is 10.
<code>log.fun</code>	[function NULL] Function used for logging. If set to NULL, the internal default will be used. Otherwise a function with arguments <code>learner</code> , <code>resampling</code> , <code>measures</code> , <code>par.set</code> , <code>control</code> , <code>opt.path</code> , <code>dob</code> , <code>x</code> , <code>y</code> , <code>remove.nas</code> , and <code>stage</code> is expected. The default displays the performance measures, the time needed for evaluating, the currently used memory and the max memory ever used before (the latter two both taken from gc). See the implementation for details.
<code>final.dw.perc</code>	[boolean] If a Learner wrapped by a makeDownsampleWrapper is used, you can define the value of <code>dw.perc</code> which is used to train the Learner with the final parameter setting found by the tuning. Default is NULL which will not change anything.

<code>budget</code>	<code>[integer(1)]</code> Maximum budget for tuning. This value restricts the number of function evaluations. In case of <code>makeTuneMultiCritControlGrid</code> this number must be identical to the size of the grid. For <code>makeTuneMultiCritControlRandom</code> the budget equals the number of iterations (<code>maxit</code>) performed by the random search algorithm. And in case of <code>makeTuneMultiCritControlNSGA2</code> the budget corresponds to the product of the maximum number of generations (<code>max(generations)</code>) + 1 (for the initial population) and the size of the population (<code>popsize</code>).
<code>impute.val</code>	<code>[numeric]</code> If something goes wrong during optimization (e.g. the learner crashes), this value is fed back to the tuner, so the tuning algorithm does not abort. It is not stored in the optimization path, an NA and a corresponding error message are logged instead. Note that this value is later multiplied by -1 for maximization measures internally, so you need to enter a larger positive value for maximization here as well. Default is the worst obtainable value of the performance measure you optimize for when you aggregate by mean value, or Inf instead. For multi-criteria optimization pass a vector of imputation values, one for each of your measures, in the same order as your measures.
<code>...</code>	<code>[any]</code> Further control parameters passed to the control arguments of <code>cma_es</code> or <code>GenSA</code> , as well as towards the <code>tunerConfig</code> argument of <code>irace</code> .
<code>maxit</code>	<code>[integer(1)]</code> Number of iterations for random search. Default is 100.

Value

`TuneMultiCritControl` . The specific subclass is one of `TuneMultiCritControlGrid`, `TuneMultiCritControlRandom`, `TuneMultiCritControlNSGA2`.

See Also

Other `tune_multicrit`: `plotTuneMultiCritResultGGVIS`, `plotTuneMultiCritResult`, `tuneParamsMultiCrit`

`TuneMultiCritResult` *Result of multi-criteria tuning.*

Description

Container for results of hyperparameter tuning. Contains the obtained pareto set and front and the optimization path which lead there.

Object members:

learner [`Learner`] Learner that was optimized.

control [`TuneControl`] Control object from tuning.

x [list] List of lists of non-dominated hyperparameter settings in pareto set. Note that when you have trafos on some of your params, x will always be on the TRANSFORMED scale so you directly use it.

y [matrix] Pareto front for x.

opt.path [OptPath] Optimization path which lead to x. Note that when you have trafos on some of your params, the opt.path always contains the UNTRANSFORMED values on the original scale. You can simply call `trafoOptPath(opt.path)` to transform them, or, as `data.frame(trafoOptPath(opt.pat`

tuneParams

Hyperparameter tuning.

Description

Optimizes the hyperparameters of a learner. Allows for different optimization methods, such as grid search, evolutionary strategies, iterated F-race, etc. You can select such an algorithm (and its settings) by passing a corresponding control object. For a complete list of implemented algorithms look at [TuneControl](#).

Multi-criteria tuning can be done with [tuneParamsMultiCrit](#).

Usage

```
tuneParams(learner, task, resampling, measures, par.set, control,
  show.info = getMlrOption("show.info"))
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
task	[Task] The task.
resampling	[ResampleInstance ResampleDesc] Resampling strategy to evaluate points in hyperparameter space. If you pass a description, it is instantiated once at the beginning by default, so all points are evaluated on the same training/test sets. If you want to change that behavior, look at TuneControl .
measures	[list of Measure Measure] Performance measures to evaluate. The first measure, aggregated by the first aggregation function is optimized, others are simply evaluated. Default is the default measure for the task, see here getDefaultMeasure .
par.set	[ParamSet] Collection of parameters and their constraints for optimization. Dependent parameters with a requires field must use quote and not expression to define it.

control	[TuneControl] Control object for search method. Also selects the optimization algorithm for tuning.
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .

Value

[TuneResult](#) .

See Also

Other tune: [TuneControl](#), [getNestedTuneResultsOptPathDf](#), [getNestedTuneResultsX](#), [getTuneResult](#), [makeModelMultiplexerParamSet](#), [makeModelMultiplexer](#), [makeTuneWrapper](#), [tuneThreshold](#)

Examples

```
# a grid search for an SVM (with a tiny number of points...)
# note how easily we can optimize on a log-scale
ps = makeParamSet(
  makeNumericParam("C", lower = -12, upper = 12, trafo = function(x) 2^x),
  makeNumericParam("sigma", lower = -12, upper = 12, trafo = function(x) 2^x)
)
ctrl = makeTuneControlGrid(resolution = 2L)
rdesc = makeResampleDesc("CV", iters = 2L)
res = tuneParams("classif.ksvm", iris.task, rdesc, par.set = ps, control = ctrl)
print(res)
print(as.data.frame(res$opt.path))
print(as.data.frame(trafoOptPath(res$opt.path)))

## Not run:
# we optimize the SVM over 3 kernels simultaneously
# note how we use dependent params (requires = ...) and iterated F-racing here
ps = makeParamSet(
  makeNumericParam("C", lower = -12, upper = 12, trafo = function(x) 2^x),
  makeDiscreteParam("kernel", values = c("vanilladot", "polydot", "rbfdot")),
  makeNumericParam("sigma", lower = -12, upper = 12, trafo = function(x) 2^x,
    requires = quote(kernel == "rbfdot")),
  makeIntegerParam("degree", lower = 2L, upper = 5L,
    requires = quote(kernel == "polydot"))
)
print(ps)
ctrl = makeTuneControlIrace(maxExperiments = 200L)
rdesc = makeResampleDesc("Holdout")
res = tuneParams("classif.ksvm", iris.task, rdesc, par.set = ps, control = ctrl)
print(res)
print(head(as.data.frame(res$opt.path)))

## End(Not run)
```

tuneParamsMultiCrit *Hyperparameter tuning for multiple measures at once.*

Description

Optimizes the hyperparameters of a learner in a multi-criteria fashion. Allows for different optimization methods, such as grid search, evolutionary strategies, etc. You can select such an algorithm (and its settings) by passing a corresponding control object. For a complete list of implemented algorithms look at [TuneMultiCritControl](#).

Usage

```
tuneParamsMultiCrit(learner, task, resampling, measures, par.set, control,
  show.info = getMlrOption("show.info"))
```

Arguments

learner	[Learner character(1)] The learner. If you pass a string the learner will be created via makeLearner .
task	[Task] The task.
resampling	[ResampleInstance ResampleDesc] Resampling strategy to evaluate points in hyperparameter space. If you pass a description, it is instantiated once at the beginning by default, so all points are evaluated on the same training/test sets. If you want to change that behavior, look at TuneMultiCritControl .
measures	[list of Measure] Performance measures to optimize simultaneously.
par.set	[ParamSet] Collection of parameters and their constraints for optimization. Dependent parameters with a requires field must use quote and not expression to define it.
control	[TuneMultiCritControl] Control object for search method. Also selects the optimization algorithm for tuning.
show.info	[logical(1)] Print verbose output on console? Default is set via configureMlr .

Value

[TuneMultiCritResult](#) .

See Also

Other tune_multicrit: [TuneMultiCritControl](#), [plotTuneMultiCritResultGGVIS](#), [plotTuneMultiCritResult](#)

Examples

```
# multi-criteria optimization of (tpr, fpr) with NGS-II
lrn = makeLearner("classif.ksvm")
rdesc = makeResampleDesc("Holdout")
ps = makeParamSet(
  makeNumericParam("C", lower = -12, upper = 12, trafo = function(x) 2^x),
  makeNumericParam("sigma", lower = -12, upper = 12, trafo = function(x) 2^x)
)
ctrl = makeTuneMultiCritControlNSGA2(popsi = 4L, generations = 1L)
res = tuneParamsMultiCrit(lrn, sonar.task, rdesc, par.set = ps,
  measures = list(tpr, fpr), control = ctrl)
plotTuneMultiCritResult(res, path = TRUE)
```

TuneResult	Result of tuning.
------------	-------------------

Description

Container for results of hyperparameter tuning. Contains the obtained point in search space, its performance values and the optimization path which lead there.

Object members:

learner [[Learner](#)] Learner that was optimized.

control [[TuneControl](#)] Control object from tuning.

x [[list](#)] Named list of hyperparameter values identified as optimal. Note that when you have trafos on some of your params, x will always be on the TRANSFORMED scale so you directly use it.

y [[numeric](#)] Performance values for optimal x.

threshold [[numeric](#)] Vector of finally found and used thresholds if tune.threshold was enabled in [TuneControl](#), otherwise not present and hence NULL.

opt.path [[OptPath](#)] Optimization path which lead to x. Note that when you have trafos on some of your params, the opt.path always contains the UNTRANSFORMED values on the original scale. You can simply call `trafoOptPath(opt.path)` to transform them, or, as `as.data.frame(trafoOptPath(opt.pat`

tuneThreshold	Tune prediction threshold.
---------------	----------------------------

Description

Optimizes the threshold of predictions based on probabilities. Works for classification and multi-label tasks. Uses [optimizeSubInts](#) for normal binary class problems and [cma_es](#) for multiclass and multilabel problems.

Usage

```
tuneThreshold(pred, measure, task, model, nsub = 20L, control = list())
```

Arguments

pred	[Prediction] Prediction object.
measure	[Measure] Performance measure to optimize. Default is the default measure for the task.
task	[Task] Learning task. Rarely needed, only when required for the performance measure.
model	[WrappedModel] Fitted model. Rarely needed, only when required for the performance measure.
nsub	[integer(1)] Passed to optimizeSubInts for 2class problems. Default is 20.
control	[list] Control object for cma_es when used. Default is empty list.

Value

list . A named list with with the following components: th is the optimal threshold, perf the performance value.

See Also

Other tune: [TuneControl](#), [getNestedTuneResultsOptPathDf](#), [getNestedTuneResultsX](#), [getTuneResult](#), [makeModelMultiplexerParamSet](#), [makeModelMultiplexer](#), [makeTuneWrapper](#), [tuneParams](#)

wpbc.task	<i>Wisconsin Prognostic Breast Cancer (WPBC) survival task.</i>
-----------	---

Description

Contains the task (wpbc.task).

References

See [wpbc](#). Incomplete cases have been removed from the task.

`yeast.task`*Yeast multilabel classification task.*

Description

Contains the task (`yeast.task`).

Source

<http://sourceforge.net/projects/mulan/files/datasets/yeast.rar>

References

Elisseeff, A., & Weston, J. (2001): A kernel method for multi-labelled classification. In Advances in neural information processing systems (pp. 681-687).

Index

- *Topic **datasets**
 - aggregations, [7](#)
 - measures, [120](#)
- *Topic **data**
 - agri.task, [9](#)
 - bc.task, [10](#)
 - bh.task, [13](#)
 - costiris.task, [17](#)
 - iris.task, [74](#)
 - lung.task, [80](#)
 - mtcars.task, [126](#)
 - pid.task, [129](#)
 - sonar.task, [166](#)
 - wdbc.task, [181](#)
 - yeast.task, [182](#)
- acc (measures), [120](#)
- Aggregation, [6](#), [8](#), [81](#), [99](#), [100](#), [160](#)
- aggregations, [6](#), [7](#), [81](#), [160](#)
- agri.task, [9](#)
- agriculture, [9](#)
- analyzeFeatSelResult, [9](#), [25](#), [52](#), [92](#), [160](#)
- arsq (measures), [120](#)
- as.data.frame, [12](#)
- asROCRPrediction, [10](#), [59](#), [60](#), [148](#), [149](#), [163](#), [164](#)
- auc (measures), [120](#)
- b632 (aggregations), [7](#)
- b632plus (aggregations), [7](#)
- bac (measures), [120](#)
- bc.task, [10](#)
- benchmark, [11](#), [12](#), [16](#), [27](#), [28](#), [31](#), [33](#), [39–48](#), [130](#), [132](#), [133](#), [135](#)
- BenchmarkResult, [11](#), [12](#), [12](#), [16](#), [27–29](#), [31](#), [38–48](#), [124](#), [125](#), [130–133](#), [135](#), [147](#)
- ber (measures), [120](#)
- bh.task, [13](#)
- bootstrapB632 (resample), [154](#)
- bootstrapB632plus (resample), [154](#)
- bootstrapO0B (resample), [154](#)
- BostonHousing, [13](#)
- BreastCancer, [10](#)
- brier (measures), [120](#)
- browseURL, [147](#)
- CalibrationData
 - (generateCalibrationData), [29](#)
- capLargeValues, [13](#), [18](#), [20](#), [126](#), [127](#), [153](#), [168](#)
- cindex (measures), [120](#)
- ClassifTask, [86](#)
- ClassifTask (makeClassifTask), [83](#)
- ClusterTask (makeClassifTask), [83](#)
- cma_es, [172](#), [174](#), [176](#), [180](#), [181](#)
- configureMlr, [11](#), [14](#), [22](#), [33](#), [51](#), [55](#), [75](#), [92](#), [98](#), [116](#), [153](#), [155](#), [157](#), [159](#), [178](#), [179](#)
- convertBMRTToRankMatrix, [12](#), [15](#), [27](#), [28](#), [31](#), [39–48](#), [130](#), [132](#), [133](#), [135](#)
- convertMLBenchObjToTask, [16](#)
- costiris.task, [17](#)
- CostSensClassifModel
 - (makeCostSensClassifWrapper), [86](#)
- CostSensClassifWrapper
 - (makeCostSensClassifWrapper), [86](#)
- CostSensRegrModel
 - (makeCostSensRegrWrapper), [87](#)
- CostSensRegrWrapper
 - (makeCostSensRegrWrapper), [87](#)
- CostSensTask, [62](#)
- CostSensTask (makeClassifTask), [83](#)
- CostSensWeightedPairsModel
 - (makeCostSensWeightedPairsWrapper), [88](#)
- CostSensWeightedPairsWrapper, [78](#)
- CostSensWeightedPairsWrapper
 - (makeCostSensWeightedPairsWrapper), [88](#)

- createDummyFeatures, [14](#), [17](#), [20](#), [126](#), [127](#), [153](#), [168](#)
- crossover, [18](#), [19](#)
- crossval (resample), [154](#)
- cv10 (makeResampleDesc), [108](#)
- cv2 (makeResampleDesc), [108](#)
- cv3 (makeResampleDesc), [108](#)
- cv5 (makeResampleDesc), [108](#)
- daisy, [165](#)
- db (measures), [120](#)
- downsample, [19](#), [90](#), [91](#)
- dropFeatures, [14](#), [18](#), [19](#), [126](#), [127](#), [153](#), [168](#)
- dunn (measures), [120](#)
- estimateRelativeOverfitting, [20](#), [86](#), [90](#), [100](#), [124](#), [129](#)
- estimateResidualVariance, [21](#)
- expvar (measures), [120](#)
- f1 (measures), [120](#)
- FailureModel, [22](#)
- fdr (measures), [120](#)
- featperc (measures), [120](#)
- FeatSelControl, [9](#), [22](#), [25](#), [52](#), [91](#), [92](#), [159](#), [160](#)
- FeatSelControlExhaustive, [25](#)
- FeatSelControlExhaustive (FeatSelControl), [22](#)
- FeatSelControlGA, [25](#)
- FeatSelControlGA (FeatSelControl), [22](#)
- FeatSelControlRandom, [25](#)
- FeatSelControlRandom (FeatSelControl), [22](#)
- FeatSelControlSequential, [25](#)
- FeatSelControlSequential (FeatSelControl), [22](#)
- FeatSelResult, [9](#), [25](#), [39](#), [52](#), [160](#)
- filterFeatures, [26](#), [32](#), [53](#), [93](#), [94](#), [136](#), [137](#)
- FilterValues, [26](#), [53](#), [135](#), [136](#)
- FilterValues (generateFilterValuesData), [32](#)
- fn (measures), [120](#)
- fnr (measures), [120](#)
- fp (measures), [120](#)
- fpr (measures), [120](#)
- friedman.test, [27](#), [28](#)
- friedmanPostHocTestBMR, [12](#), [16](#), [27](#), [28](#), [31](#), [39–48](#), [130](#), [132](#), [133](#), [135](#)
- friedmanTestBMR, [12](#), [16](#), [27](#), [28](#), [31](#), [39–48](#), [130](#), [132](#), [133](#), [135](#)
- G1 (measures), [120](#)
- G2 (measures), [120](#)
- gc, [24](#), [173](#), [175](#)
- generateCalibrationData, [29](#), [31](#), [32](#), [34](#), [37](#), [38](#), [53](#), [133](#), [134](#)
- generateCritDifferencesData, [12](#), [16](#), [27](#), [28](#), [30](#), [30](#), [32](#), [34](#), [37–48](#), [53](#), [130](#), [132–135](#)
- generateDesign, [174](#)
- generateFilterValuesData, [26](#), [30](#), [31](#), [32](#), [34](#), [37](#), [38](#), [53](#), [94](#), [136](#), [137](#)
- generateGridDesign, [174](#), [175](#)
- generateLearningCurveData, [30–32](#), [33](#), [34](#), [37](#), [38](#), [53](#), [139](#), [140](#)
- generatePartialPredictionData, [30–32](#), [34](#), [34](#), [38](#), [53](#), [140–142](#)
- generateThreshVsPerfData, [30–32](#), [34](#), [37](#), [37](#), [53](#), [142–145](#)
- GenSA, [172](#), [174](#), [176](#)
- geom_point, [132](#), [138](#), [146](#)
- getBMRAggrPerformances, [12](#), [16](#), [27](#), [28](#), [31](#), [38](#), [40–48](#), [130–133](#), [135](#)
- getBMRFeatSelResults, [12](#), [16](#), [27](#), [28](#), [31](#), [39](#), [39](#), [41–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRFilteredFeatures, [12](#), [16](#), [27](#), [28](#), [31](#), [39](#), [40](#), [40](#), [42–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRLearnerIds, [12](#), [16](#), [27](#), [28](#), [31](#), [39–41](#), [41](#), [42–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRLearners, [12](#), [16](#), [27](#), [28](#), [31](#), [39–42](#), [42](#), [43–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRLearnerShortNames, [12](#), [16](#), [27](#), [28](#), [31](#), [39–42](#), [42](#), [43–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRMeasureIds, [12](#), [16](#), [27](#), [28](#), [31](#), [39–43](#), [43](#), [44–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRMeasures, [12](#), [16](#), [27](#), [28](#), [31](#), [39–43](#), [44](#), [45–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRModels, [12](#), [16](#), [27](#), [28](#), [31](#), [39–44](#), [44](#), [45–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRPerformances, [12](#), [16](#), [27](#), [28](#), [31](#), [39–45](#), [45](#), [46–48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRPredictions, [12](#), [16](#), [27](#), [28](#), [31](#), [39–45](#), [46](#), [47](#), [48](#), [130](#), [132](#), [133](#), [135](#)
- getBMRTaskIds, [12](#), [16](#), [27](#), [28](#), [31](#), [39–46](#), [47](#), [48](#), [130](#), [132](#), [133](#), [135](#)

- getBMRTuneResults, [12](#), [16](#), [27](#), [28](#), [31](#),
[39–47](#), [47](#), [130](#), [132](#), [133](#), [135](#)
- getCaretParamSet, [48](#)
- getClassWeightParam, [49](#), [54](#), [58](#), [77](#), [98](#),
[153](#), [161–164](#)
- getConfMatrix, [50](#)
- getDefaultMeasure, [20](#), [38](#), [51](#), [91](#), [115](#), [128](#),
[138](#), [155](#), [159](#), [177](#)
- getFailureModelMsg, [51](#)
- getFeatSelResult, [9](#), [25](#), [52](#), [91](#), [92](#), [160](#)
- getFilteredFeatures, [26](#), [32](#), [52](#), [53](#), [93](#), [94](#),
[136](#), [137](#)
- getFilterValues, [26](#), [30–32](#), [34](#), [37](#), [38](#), [53](#),
[53](#), [94](#), [136](#), [137](#)
- getHomogeneousEnsembleModels, [54](#)
- getHyperPars, [50](#), [54](#), [58](#), [77](#), [98](#), [153](#),
[161–164](#)
- getLearnerModel, [55](#), [82](#), [86–88](#), [104](#), [106](#)
- getLearnerProperties
(LearnerProperties), [76](#)
- getMlrOptions, [15](#), [55](#)
- getMultilabelBinaryPerformances, [56](#),
[105](#)
- getNestedTuneResultsOptPathDf, [56](#), [57](#),
[70](#), [101](#), [103](#), [116](#), [174](#), [178](#), [181](#)
- getNestedTuneResultsX, [57](#), [57](#), [70](#), [101](#),
[103](#), [116](#), [174](#), [178](#), [181](#)
- getParamSet, [50](#), [54](#), [58](#), [77](#), [98](#), [153](#), [161–164](#)
- getPredictionProbabilities, [10](#), [58](#), [60](#),
[148](#), [149](#), [163](#), [164](#)
- getPredictionResponse, [10](#), [59](#), [59](#), [148](#),
[149](#), [163](#), [164](#)
- getPredictionSE
(getPredictionResponse), [59](#)
- getPredictionTruth
(getPredictionResponse), [59](#)
- getProbabilities, [60](#)
- getRRPredictions, [60](#), [110](#), [111](#), [156](#), [157](#)
- getStackedBaseLearnerPredictions, [61](#)
- getTaskClassLevels, [61](#), [61](#), [62–70](#), [167](#)
- getTaskCosts, [62](#), [62](#), [63–70](#), [167](#)
- getTaskData, [62](#), [63](#), [64–70](#), [83](#), [167](#), [171](#)
- getTaskDescription, [62](#), [63](#), [64](#), [65–70](#), [167](#)
- getTaskFeatureNames, [62–64](#), [65](#), [66–70](#), [83](#),
[167](#)
- getTaskFormula, [62–65](#), [65](#), [66–70](#), [83](#), [167](#)
- getTaskId, [62–66](#), [66](#), [67–70](#), [167](#)
- getTaskNFeats, [62–66](#), [67](#), [68–70](#), [167](#)
- getTaskSize, [62–67](#), [67](#), [68–70](#), [167](#)
- getTaskTargetNames, [61–68](#), [68](#), [69](#), [70](#), [167](#)
- getTaskTargets, [62–68](#), [68](#), [70](#), [83](#), [167](#)
- getTaskType, [62–69](#), [69](#), [167](#)
- getTuneResult, [57](#), [70](#), [101](#), [103](#), [115](#), [116](#),
[174](#), [178](#), [181](#)
- ggplot, [137](#)
- gmean (measures), [120](#)
- gpr (measures), [120](#)
- grad, [35](#)
- hamloss (measures), [120](#)
- hasLearnerProperties
(LearnerProperties), [76](#)
- hasProperties, [70](#)
- hist, [29](#), [72](#)
- holdout (resample), [154](#)
- hout (makeResampleDesc), [108](#)
- imputations, [71](#), [73](#), [74](#), [95](#), [97](#), [152](#)
- impute, [72](#), [72](#), [95–97](#), [151](#), [152](#)
- imputeConstant (imputations), [71](#)
- imputeHist (imputations), [71](#)
- imputeLearner (imputations), [71](#)
- imputeMax (imputations), [71](#)
- imputeMean (imputations), [71](#)
- imputeMedian (imputations), [71](#)
- imputeMin (imputations), [71](#)
- imputeMode (imputations), [71](#)
- imputeNormal (imputations), [71](#)
- imputeUniform (imputations), [71](#)
- integer, [111](#)
- irace, [172](#), [174](#), [176](#)
- iris, [17](#), [74](#)
- iris.task, [74](#)
- isFailureModel, [75](#)
- jacobian, [35](#)
- joinClassLevels, [75](#)
- Learner, [11](#), [12](#), [20](#), [21](#), [25](#), [28](#), [33](#), [49](#), [51](#), [54](#),
[58](#), [77](#), [79](#), [82](#), [86–88](#), [90–94](#), [96](#), [98](#),
[101](#), [104–108](#), [112](#), [113](#), [115–120](#),
[137](#), [148](#), [153–155](#), [159](#), [161–164](#),
[170](#), [176](#), [177](#), [179](#), [180](#)
- Learner (makeLearner), [97](#)
- learnerArgsToControl, [76](#)
- LearnerParam, [49](#), [54](#), [107](#)
- LearnerProperties, [50](#), [54](#), [58](#), [76](#), [97](#), [98](#),
[153](#), [157](#), [161–164](#)

- learners, [77](#), [97](#), [158](#)
- LearningCurveData
 - (generateLearningCurveData), [33](#)
- listFilterMethods, [26](#), [32](#), [53](#), [77](#), [92](#), [93](#)
- listLearners, [77](#), [78](#)
- listMeasures, [79](#), [120](#)
- lung, [80](#)
- lung.task, [80](#)
- mae (measures), [120](#)
- makeAggregation, [6](#), [81](#)
- makeBaggingWrapper, [82](#), [87](#), [91](#), [92](#), [94](#), [97](#), [104–108](#), [113](#), [116](#), [117](#), [119](#)
- makeClassifTask, [83](#), [87](#), [88](#)
- makeClusterTask (makeClassifTask), [83](#)
- makeCostMeasure, [21](#), [85](#), [90](#), [100](#), [120](#), [124](#), [129](#)
- makeCostSensClassifWrapper, [83](#), [85](#), [86](#), [87](#), [88](#), [91](#), [92](#), [94](#), [97](#), [104–108](#), [113](#), [116](#), [117](#), [119](#)
- makeCostSensRegrWrapper, [83](#), [85](#), [87](#), [87](#), [88](#), [91](#), [92](#), [94](#), [97](#), [104–108](#), [113](#), [116](#), [117](#), [119](#)
- makeCostSensTask (makeClassifTask), [83](#)
- makeCostSensWeightedPairsWrapper, [85](#), [87](#), [88](#)
- makeCustomResampledMeasure, [21](#), [86](#), [88](#), [100](#), [124](#), [129](#)
- makeDiscreteParam, [172](#), [175](#)
- makeDownsampleWrapper, [19](#), [33](#), [83](#), [87](#), [90](#), [92](#), [94](#), [97](#), [104–108](#), [113](#), [116](#), [117](#), [119](#), [173](#), [175](#)
- makeFeatSelControlExhaustive (FeatSelControl), [22](#)
- makeFeatSelControlGA (FeatSelControl), [22](#)
- makeFeatSelControlRandom (FeatSelControl), [22](#)
- makeFeatSelControlSequential (FeatSelControl), [22](#)
- makeFeatSelWrapper, [9](#), [25](#), [52](#), [83](#), [87](#), [91](#), [91](#), [94](#), [97](#), [104–108](#), [113](#), [116](#), [117](#), [119](#), [160](#)
- makeFilter, [92](#)
- makeFilterWrapper, [26](#), [32](#), [52](#), [53](#), [83](#), [87](#), [91](#), [92](#), [93](#), [97](#), [104–108](#), [113](#), [116](#), [117](#), [119](#), [136](#), [137](#)
- makeFixedHoldoutInstance, [94](#), [110](#)
- makeImputeMethod, [72–74](#), [95](#), [97](#), [152](#)
- makeImputeWrapper, [72](#), [74](#), [83](#), [87](#), [91](#), [92](#), [94](#), [95](#), [96](#), [104–108](#), [113](#), [116](#), [117](#), [119](#), [152](#)
- makeLearner, [20](#), [49](#), [50](#), [54](#), [58](#), [77](#), [82](#), [86–88](#), [90](#), [91](#), [93](#), [96](#), [97](#), [104–108](#), [112](#), [115](#), [117](#), [118](#), [120](#), [137](#), [153](#), [155](#), [159](#), [161–164](#), [170](#), [177](#), [179](#)
- makeMeasure, [21](#), [86](#), [90](#), [99](#), [120](#), [124](#), [129](#)
- makeModelMultiplexer, [57](#), [70](#), [101](#), [103](#), [116](#), [174](#), [178](#), [181](#)
- makeModelMultiplexerParamSet, [57](#), [70](#), [101](#), [102](#), [116](#), [174](#), [178](#), [181](#)
- makeMulticlassWrapper, [83](#), [87](#), [91](#), [92](#), [94](#), [97](#), [103](#), [105–108](#), [113](#), [116](#), [117](#), [119](#)
- makeMultilabelBinaryRelevanceWrapper, [56](#), [78](#), [83](#), [87](#), [91](#), [92](#), [94](#), [97](#), [104](#), [104](#), [106–108](#), [113](#), [116](#), [117](#), [119](#)
- makeMultilabelTask (makeClassifTask), [83](#)
- makeOverBaggingWrapper, [83](#), [87](#), [91](#), [92](#), [94](#), [97](#), [104](#), [105](#), [105](#), [107](#), [108](#), [113](#), [116](#), [117](#), [119](#), [128](#), [166](#)
- makeOversampleWrapper (makeUndersampleWrapper), [116](#)
- makePreprocWrapper, [83](#), [87](#), [91](#), [92](#), [94](#), [97](#), [104–106](#), [107](#), [108](#), [113](#), [116](#), [117](#), [119](#)
- makePreprocWrapperCaret, [83](#), [87](#), [91](#), [92](#), [94](#), [97](#), [104–107](#), [108](#), [113](#), [116](#), [117](#), [119](#)
- makeRegrTask (makeClassifTask), [83](#)
- makeResampleDesc, [61](#), [108](#), [111](#), [156](#), [157](#)
- makeResampleInstance, [19](#), [61](#), [110](#), [111](#), [156](#), [157](#)
- makeRLearner (RLearner), [157](#)
- makeRLearnerClassif (RLearner), [157](#)
- makeRLearnerCluster (RLearner), [157](#)
- makeRLearnerMultilabel (RLearner), [157](#)
- makeRLearnerRegr (RLearner), [157](#)
- makeRLearnerSurv (RLearner), [157](#)
- makeSMOTERWrapper, [83](#), [87](#), [91](#), [92](#), [94](#), [97](#), [104–108](#), [112](#), [116](#), [117](#), [119](#)
- makeStackedLearner, [113](#)
- makeSurvTask (makeClassifTask), [83](#)
- makeTuneControlCMAES (TuneControl), [172](#)
- makeTuneControlDesign (TuneControl), [172](#)
- makeTuneControlGenSA (TuneControl), [172](#)
- makeTuneControlGrid (TuneControl), [172](#)
- makeTuneControlIrace, [101](#)

- makeTuneControlIrace (TuneControl), 172
- makeTuneControlRandom (TuneControl), 172
- makeTuneMultiCritControlGrid
 - (TuneMultiCritControl), 174
- makeTuneMultiCritControlNSGA2
 - (TuneMultiCritControl), 174
- makeTuneMultiCritControlRandom
 - (TuneMultiCritControl), 174
- makeTuneWrapper, 11, 56, 57, 70, 83, 87, 91, 92, 94, 97, 101, 103–108, 113, 115, 117, 119, 174, 178, 181
- makeUndersampleWrapper, 83, 87, 91, 92, 94, 97, 104–108, 113, 116, 116, 119, 128, 166
- makeWeightedClassesWrapper, 83, 87, 91, 92, 94, 97, 104–108, 113, 116, 117, 118
- makeWrappedModel, 119
- mcc (measures), 120
- mcp (measures), 120
- mean, 86
- meancosts (measures), 120
- Measure, 11, 12, 16, 20, 27, 28, 31, 33, 34, 38, 51, 56, 80, 81, 86, 88, 90, 91, 100, 115, 128, 130–132, 138–140, 142, 144, 145, 155, 159, 160, 177, 179, 181
- Measure (makeMeasure), 99
- measureACC (measures), 120
- measureAUC (measures), 120
- measureBAC (measures), 120
- measureBrier (measures), 120
- measureEXPVAR (measures), 120
- measureFDR (measures), 120
- measureFN (measures), 120
- measureFNR (measures), 120
- measureFP (measures), 120
- measureFPR (measures), 120
- measureGMEAN (measures), 120
- measureGPR (measures), 120
- measureHAMLOSS (measures), 120
- measureMAE (measures), 120
- measureMCC (measures), 120
- measureMEDAE (measures), 120
- measureMEDSE (measures), 120
- measureMMCE (measures), 120
- measureMSE (measures), 120
- measureNPV (measures), 120
- measurePPV (measures), 120
- measureRMSE (measures), 120
- measureRSQ (measures), 120
- measures, 21, 86, 90, 99, 100, 120, 129, 154
- measureSAE (measures), 120
- measureSSE (measures), 120
- measureTN (measures), 120
- measureTNR (measures), 120
- measureTP (measures), 120
- measureTPR (measures), 120
- medae (measures), 120
- medse (measures), 120
- mergeBenchmarkResultLearner, 124
- mergeBenchmarkResultTask, 125
- mergeSmallFactorLevels, 14, 18, 20, 125, 127, 153, 168
- mmce (measures), 120
- model.matrix, 17
- ModelMultiplexer, 103
- ModelMultiplexer
 - (makeModelMultiplexer), 101
- mse (measures), 120
- mtcars, 126
- mtcars.task, 126
- multiclass.auc (measures), 120
- MultilabelTask (makeClassifTask), 83
- normalize, 126, 127
- normalizeFeatures, 14, 18, 20, 126, 126, 153, 168
- npv (measures), 120
- nsga2, 175
- optimizeSubInts, 180, 181
- options, 14
- OptPath, 25, 177, 180
- oversample, 106, 117, 127, 166
- Param, 103
- ParamSet, 48, 49, 58, 103, 107, 115, 158, 177, 179
- PartialPredictionData
 - (generatePartialPredictionData), 34
- performance, 21, 81, 86, 90, 100, 124, 128
- pid.task, 129
- PimaIndiansDiabetes, 129
- plotBMRBoxplots, 12, 16, 27, 28, 31, 39–48, 130, 132–137, 139–145

- plotBMRRanksAsBarChart, [12](#), [16](#), [27](#), [28](#), [31](#), [39–48](#), [130](#), [131](#), [133–137](#), [139–145](#)
- plotBMRSummary, [12](#), [16](#), [27](#), [28](#), [31](#), [39–48](#), [130](#), [132](#), [132](#), [134–137](#), [139–145](#)
- plotCalibration, [30](#), [130](#), [132](#), [133](#), [133](#), [135–137](#), [139–145](#)
- plotCritDifferences, [12](#), [16](#), [27](#), [28](#), [31](#), [39–48](#), [130](#), [132–134](#), [134](#), [136](#), [137](#), [139–145](#)
- plotFilterValues, [26](#), [32](#), [53](#), [94](#), [130](#), [132–135](#), [135](#), [137](#), [139–145](#)
- plotFilterValuesGGVIS, [26](#), [32](#), [53](#), [94](#), [130](#), [132–136](#), [136](#), [139–145](#)
- plotLearnerPrediction, [137](#)
- plotLearningCurve, [34](#), [130](#), [132–137](#), [138](#), [140–145](#)
- plotLearningCurveGGVIS, [34](#), [130](#), [132–137](#), [139](#), [139](#), [141–145](#)
- plotPartialPrediction, [35](#), [37](#), [130](#), [132–137](#), [139](#), [140](#), [140](#), [142–145](#)
- plotPartialPredictionGGVIS, [35](#), [37](#), [130](#), [132–137](#), [139–141](#), [141](#), [143–145](#)
- plotROCCurves, [38](#), [130](#), [132–137](#), [139–142](#), [142](#), [144](#), [145](#)
- plotThreshVsPerf, [38](#), [130](#), [132–137](#), [139–143](#), [143](#), [145](#)
- plotThreshVsPerfGGVIS, [38](#), [130](#), [132–137](#), [139–144](#), [144](#)
- plotTuneMultiCritResult, [145](#), [146](#), [176](#), [179](#)
- plotTuneMultiCritResultGGVIS, [146](#), [146](#), [176](#), [179](#)
- plotViperCharts, [10](#), [59](#), [60](#), [147](#), [149](#), [163](#), [164](#)
- posthoc.friedman.nemenyi.test, [27](#), [31](#)
- ppv (measures), [120](#)
- predict, [36](#)
- predict.WrappedModel, [10](#), [50](#), [59](#), [60](#), [98](#), [148](#), [148](#), [149](#), [163–165](#), [170](#)
- Prediction, [10](#), [29](#), [38](#), [50](#), [56](#), [58](#), [60](#), [81](#), [89](#), [100](#), [128](#), [147–149](#), [149](#), [156](#), [164](#), [165](#), [181](#)
- predictLearner, [150](#)
- preProcess, [108](#)
- randomForest, [151](#)
- rank, [16](#), [131](#)
- regr.randomForest, [151](#)
- RegrTask, [21](#)
- RegrTask (makeClassifTask), [83](#)
- reimpute, [72–74](#), [95–97](#), [151](#)
- removeConstantFeatures, [14](#), [18](#), [20](#), [126](#), [127](#), [152](#), [168](#)
- removeHyperPars, [50](#), [54](#), [58](#), [77](#), [98](#), [153](#), [161–164](#)
- repcv (resample), [154](#)
- resample, [38](#), [45](#), [46](#), [60](#), [61](#), [98](#), [110](#), [111](#), [154](#), [156](#), [157](#)
- ResampleDesc, [11](#), [20](#), [33](#), [91](#), [110](#), [111](#), [114](#), [115](#), [154](#), [155](#), [159](#), [172](#), [177](#), [179](#)
- ResampleDesc (makeResampleDesc), [108](#)
- ResampleInstance, [11](#), [19](#), [33](#), [91](#), [95](#), [108](#), [111](#), [115](#), [154](#), [155](#), [159](#), [172](#), [177](#), [179](#)
- ResampleInstance
(makeResampleInstance), [111](#)
- ResamplePrediction, [38](#), [46](#), [61](#), [89](#), [110](#), [111](#), [156](#), [156](#), [157](#)
- ResampleResult, [11](#), [12](#), [29](#), [38](#), [57](#), [61](#), [110](#), [111](#), [147](#), [155](#), [156](#), [156](#)
- RLearner, [150](#), [157](#), [158](#), [171](#)
- RLearnerClassif, [158](#)
- RLearnerClassif (RLearner), [157](#)
- RLearnerCluster, [158](#)
- RLearnerCluster (RLearner), [157](#)
- RLearnerMultilabel, [158](#)
- RLearnerMultilabel (RLearner), [157](#)
- RLearnerRegr, [158](#)
- RLearnerRegr (RLearner), [157](#)
- RLearnerSurv, [158](#)
- RLearnerSurv (RLearner), [157](#)
- rmse (measures), [120](#)
- rpart, [55](#)
- rsq (measures), [120](#)
- sae (measures), [120](#)
- scale_x_log10, [132](#)
- sd, [168](#)
- selectFeatures, [9](#), [22](#), [25](#), [52](#), [91](#), [92](#), [159](#)
- setAggregation, [81](#), [86](#), [109](#), [160](#)
- setHyperPars, [50](#), [54](#), [58](#), [77](#), [98](#), [153](#), [161](#), [162–164](#)
- setHyperPars2, [162](#)
- setId, [50](#), [54](#), [58](#), [77](#), [98](#), [153](#), [161](#), [162](#), [163](#), [164](#)
- setPredictThreshold, [10](#), [50](#), [54](#), [58–60](#), [77](#), [98](#), [148](#), [149](#), [153](#), [161](#), [162](#), [163](#), [164](#)

- setPredictType, [10](#), [50](#), [54](#), [58–60](#), [77](#), [82](#),
[98](#), [148–150](#), [153](#), [161–163](#), [163](#)
- setThreshold, [97](#), [98](#), [104](#), [148](#), [163](#), [164](#)
- silhouette (measures), [120](#)
- smote, [106](#), [112](#), [117](#), [128](#), [165](#)
- Sonar, [166](#)
- sonar.task, [166](#)
- sse (measures), [120](#)
- subsample (resample), [154](#)
- subsetTask, [62–70](#), [83](#), [167](#)
- summarizeColumns, [14](#), [18](#), [20](#), [126](#), [127](#), [153](#),
[167](#)
- summarizeLevels, [168](#)
- summary, [167](#)
- Surv, [63](#), [85](#)
- SurvTask (makeClassifTask), [83](#)
- Task, [11](#), [13](#), [18–20](#), [26](#), [28](#), [32–34](#), [49](#), [51](#), [53](#),
[62–69](#), [75](#), [79–81](#), [85](#), [89](#), [100](#), [111](#),
[125–128](#), [137](#), [148](#), [152–155](#), [159](#),
[165–171](#), [177](#), [179](#), [181](#)
- Task (makeClassifTask), [83](#)
- TaskDesc, [30](#), [32](#), [36](#), [51](#), [62](#), [64](#), [66–69](#), [83](#),
[120](#), [150](#), [169](#)
- test.join (aggregations), [7](#)
- test.max (aggregations), [7](#)
- test.mean, [100](#)
- test.mean (aggregations), [7](#)
- test.median (aggregations), [7](#)
- test.min (aggregations), [7](#)
- test.range (aggregations), [7](#)
- test.rmse (aggregations), [7](#)
- test.sd (aggregations), [7](#)
- test.sum (aggregations), [7](#)
- testgroup.mean (aggregations), [7](#)
- ThreshVsPerfData
(generateThreshVsPerfData), [37](#)
- timeboth (measures), [120](#)
- timepredict (measures), [120](#)
- timetrain (measures), [120](#)
- tn (measures), [120](#)
- tnr (measures), [120](#)
- tp (measures), [120](#)
- tpr (measures), [120](#)
- train, [21](#), [34](#), [55](#), [119](#), [148](#), [155](#), [170](#)
- train.max (aggregations), [7](#)
- train.mean (aggregations), [7](#)
- train.median (aggregations), [7](#)
- train.min (aggregations), [7](#)
- train.range (aggregations), [7](#)
- train.rmse (aggregations), [7](#)
- train.sd (aggregations), [7](#)
- train.sum (aggregations), [7](#)
- trainLearner, [63](#), [171](#)
- TuneControl, [57](#), [70](#), [101](#), [103](#), [115](#), [116](#), [172](#),
[174](#), [176–178](#), [180](#), [181](#)
- TuneControlCMAES, [174](#)
- TuneControlCMAES (TuneControl), [172](#)
- TuneControlGenSA, [174](#)
- TuneControlGenSA (TuneControl), [172](#)
- TuneControlGrid, [174](#)
- TuneControlGrid (TuneControl), [172](#)
- TuneControlIrace, [174](#)
- TuneControlIrace (TuneControl), [172](#)
- TuneControlRandom, [174](#)
- TuneControlRandom (TuneControl), [172](#)
- TuneMultiCritControl, [146](#), [174](#), [176](#), [179](#)
- TuneMultiCritControlGrid, [176](#)
- TuneMultiCritControlGrid
(TuneMultiCritControl), [174](#)
- TuneMultiCritControlNSGA2, [176](#)
- TuneMultiCritControlNSGA2
(TuneMultiCritControl), [174](#)
- TuneMultiCritControlRandom, [176](#)
- TuneMultiCritControlRandom
(TuneMultiCritControl), [174](#)
- TuneMultiCritResult, [145](#), [146](#), [176](#), [179](#)
- tuneParams, [48](#), [57](#), [70](#), [101](#), [103](#), [115](#), [116](#),
[172](#), [174](#), [175](#), [177](#), [181](#)
- tuneParamsMultiCrit, [145](#), [146](#), [176](#), [177](#),
[179](#)
- TuneResult, [47](#), [70](#), [178](#), [180](#)
- tuneThreshold, [24](#), [57](#), [70](#), [101](#), [103](#), [104](#),
[116](#), [173](#), [174](#), [178](#), [180](#)
- undersample, [117](#)
- undersample (oversample), [127](#)
- which, [149](#), [170](#)
- wdbc, [181](#)
- wdbc.task, [181](#)
- WrappedModel, [21](#), [22](#), [34](#), [44](#), [51](#), [52](#), [55](#), [70](#),
[75](#), [100](#), [120](#), [128](#), [148](#), [150](#), [155](#),
[157](#), [170](#), [181](#)
- WrappedModel (makeWrappedModel), [119](#)
- yeast.task, [182](#)